

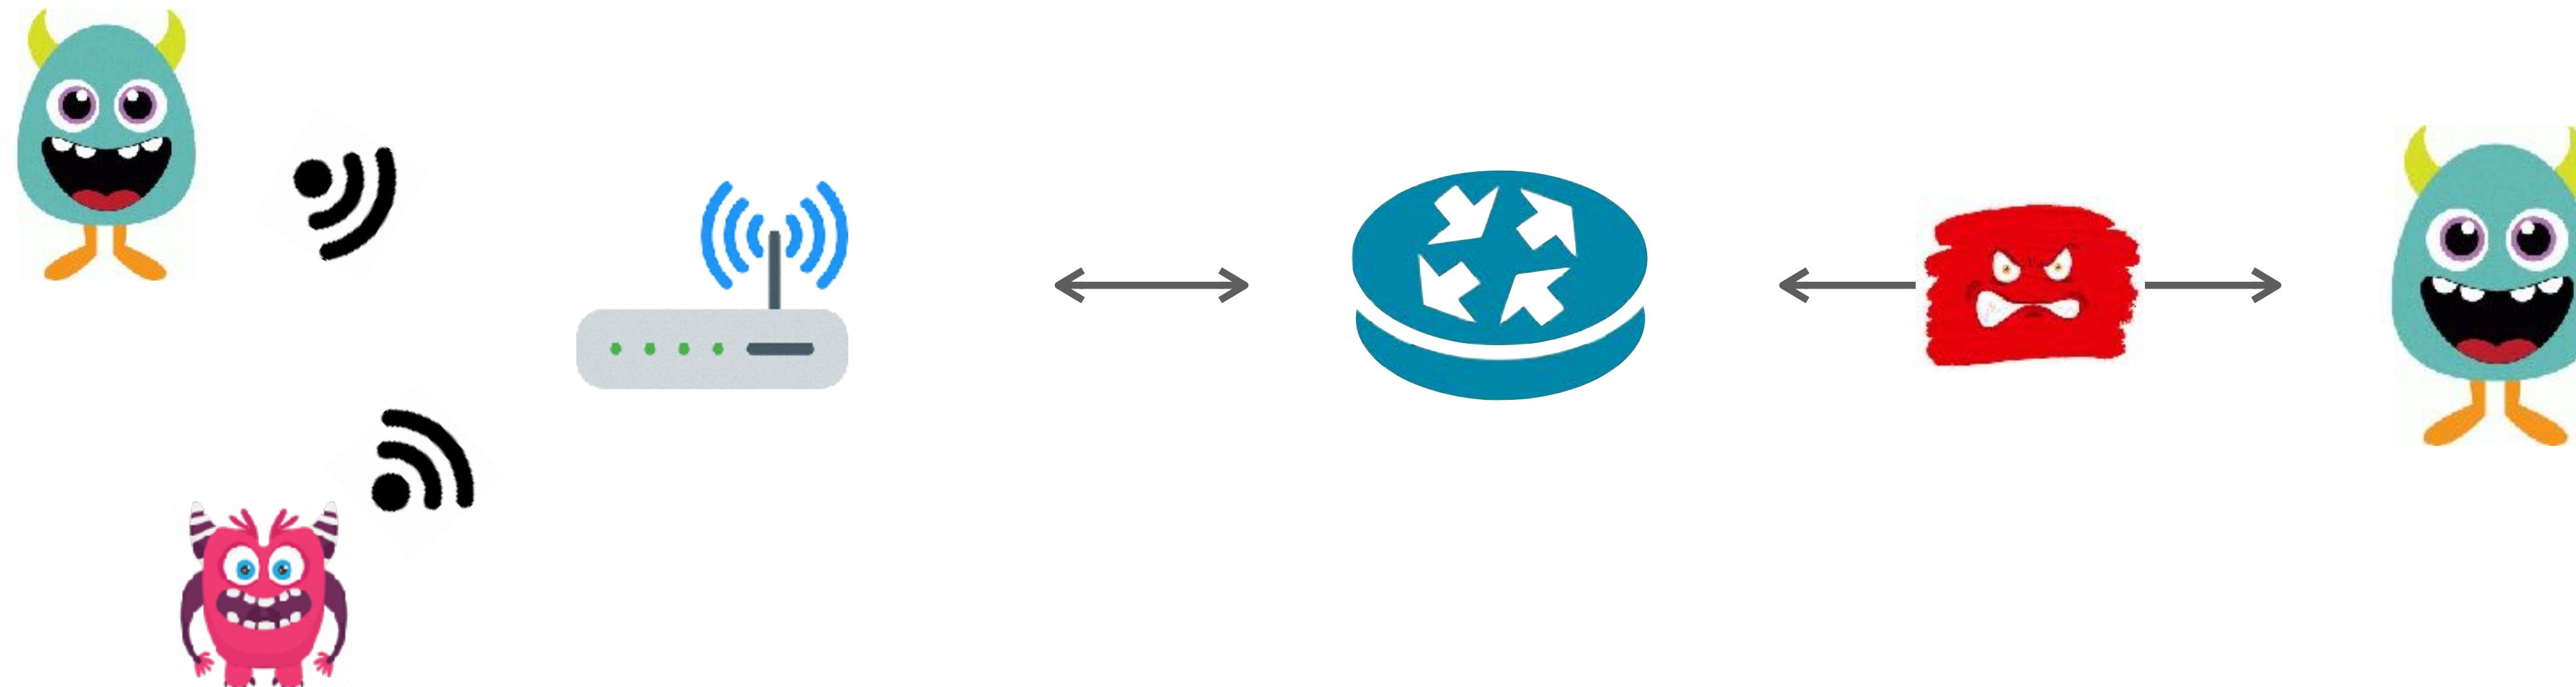
DoS Attacks and Network Defenses

CS155 Computer and Network Security

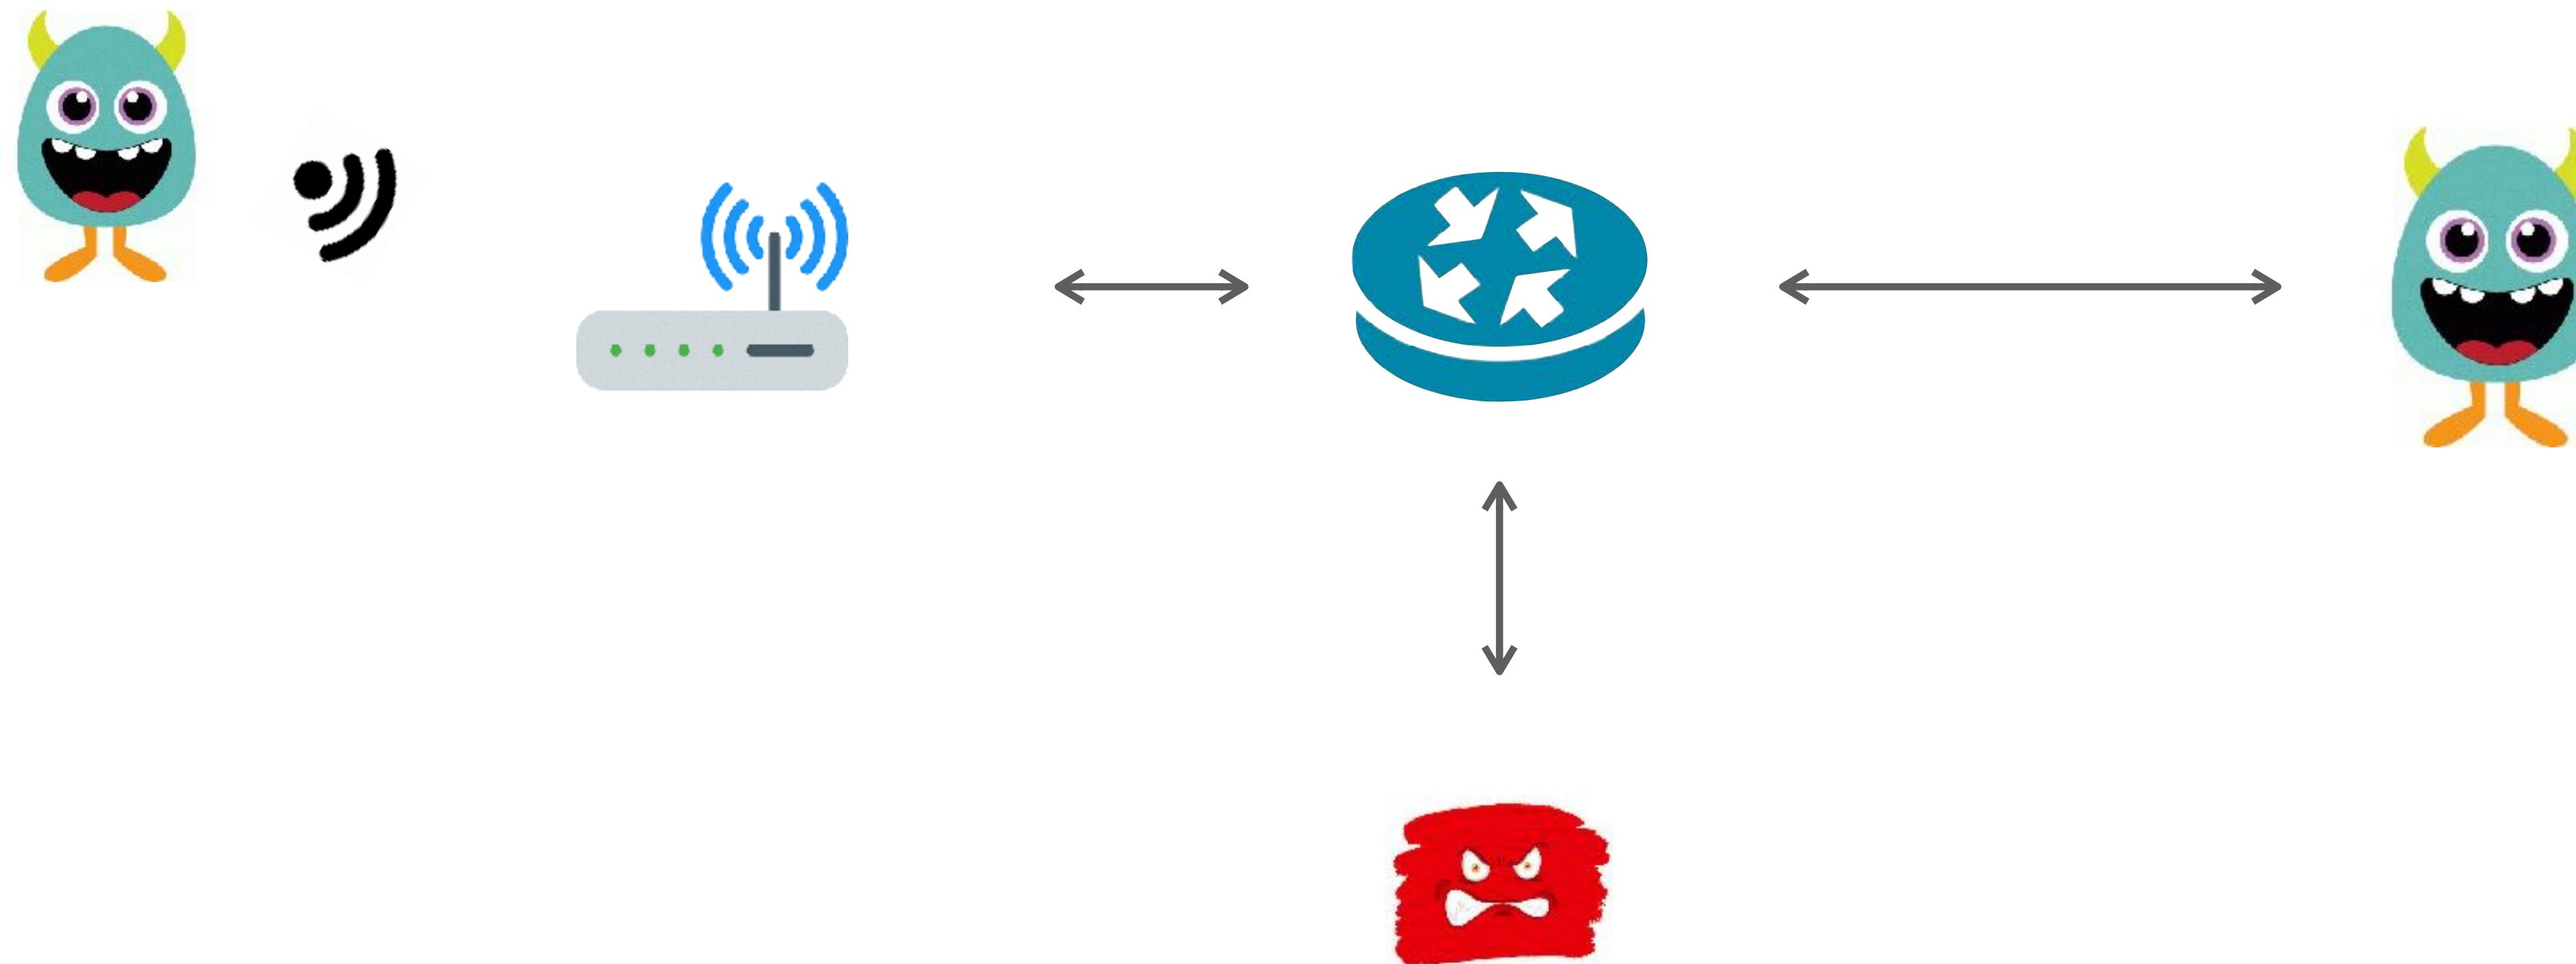
Acknowledgments: Lecture slides are from the Computer Security course taught by Dan Boneh and Zakir Durumeric at Stanford University. When slides are obtained from other sources, a reference will be noted on the bottom of that slide. A full list of references is provided on the last slide.

Stanford University

Review: On Path Attacker



Review: Off Path Attacker



No security guarantees

Confidentiality – Ethernet, IP, UDP, and TCP do not provide any confidentiality. All traffic is in cleartext.

On-path attacker can do anything. Some attacks allow an off-path attacker to become on-path and MITM connections.

Integrity – No guarantees that attacker hasn't modified traffic. Ethernet, IP and UDP have no protection against spoofed packets. TCP provides weak guarantee of source authentication.

Availability – Attackers can attempt to inject RST packets. More today.

Assume network is malicious

Always Assume: The network is out to get you.

Solution: Always use TLS if you want any protection against large-scale eavesdropping (e.g., intelligence agencies), or guarantee that data hasn't been modified or corrupted by an on-path attacker

Note! HTTPS and TLS aren't just for sensitive material! There have been attacks where malicious Javascript or malware is injected into websites.

Denial of Service Attacks

Goal: take large service/network/org offline by overwhelming it with network traffic such that they can't process real requests

How: find mechanism where attacker doesn't spend a lot of effort, but requests are difficult/expensive for victim to process

Types of Attacks

DoS Bug: design flaw that allows one machine to disrupt a service. Generally a protocol asymmetry, e.g., easy to send request, difficult to create response. Or requires server state.

DoS Flood: control a large number of requests from a botnet or other machines you control

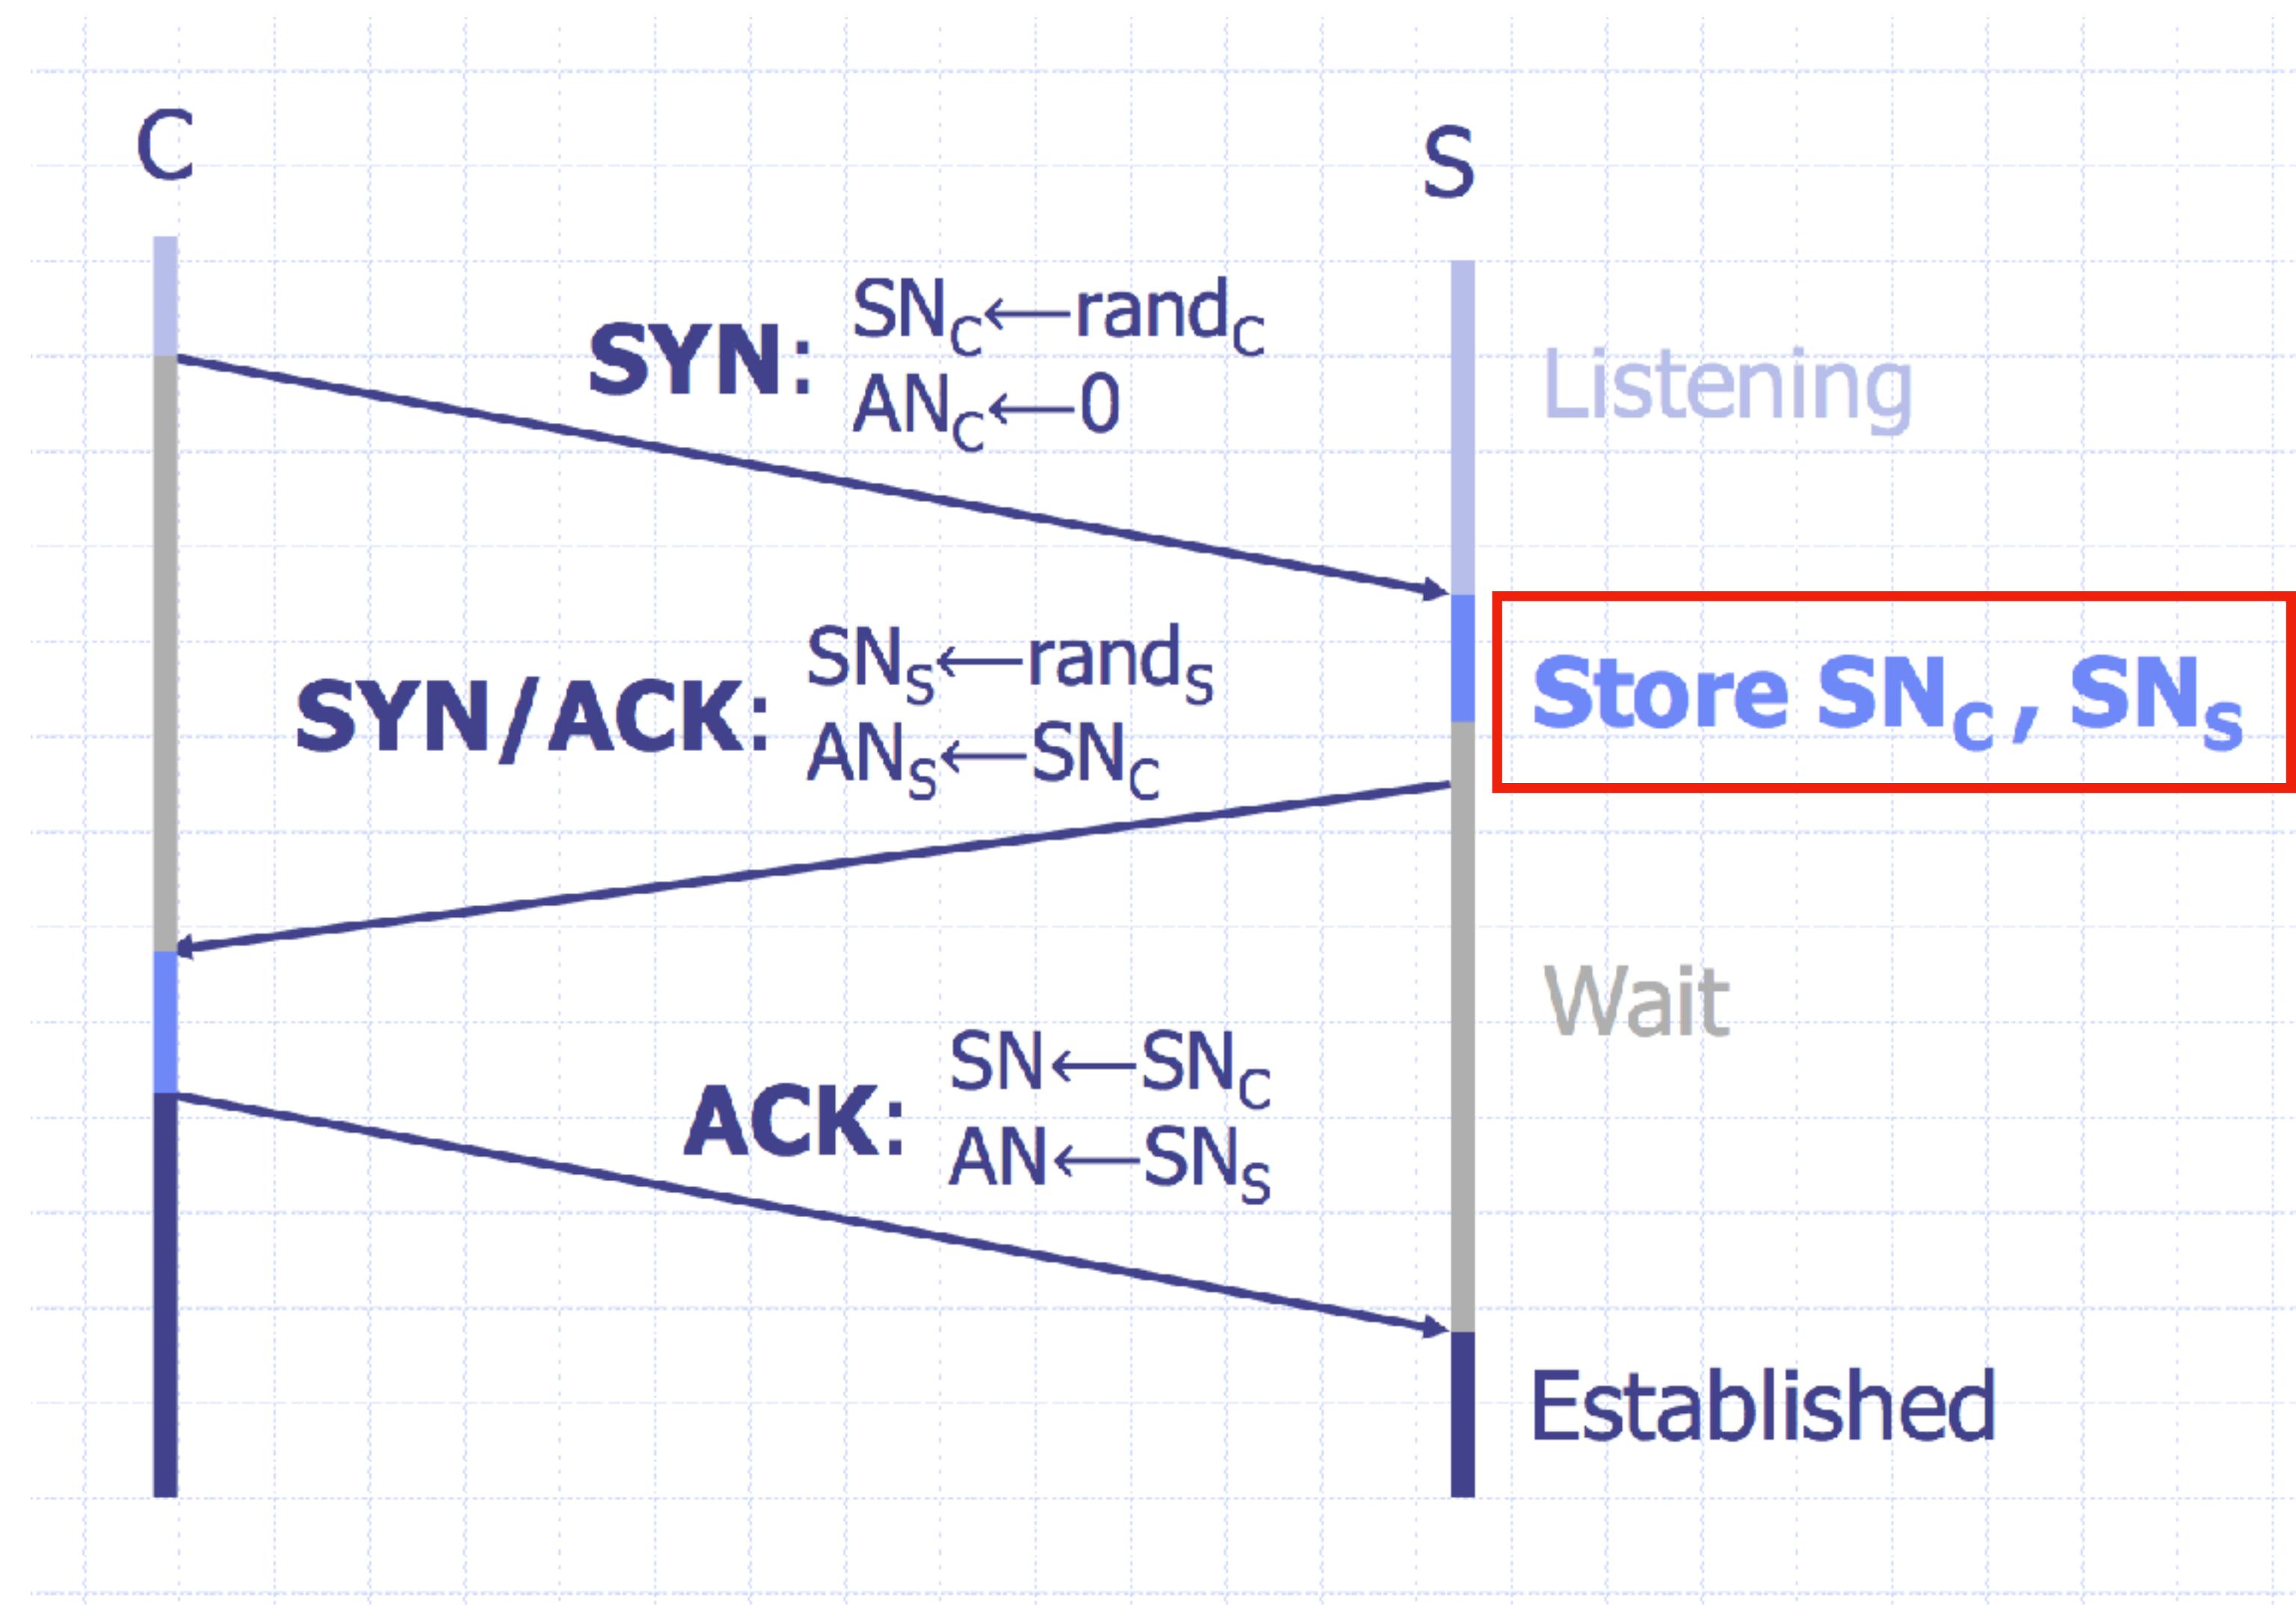
DoS Opportunities at Every Layer

Link Layer: send too much traffic for switches/routers to handle

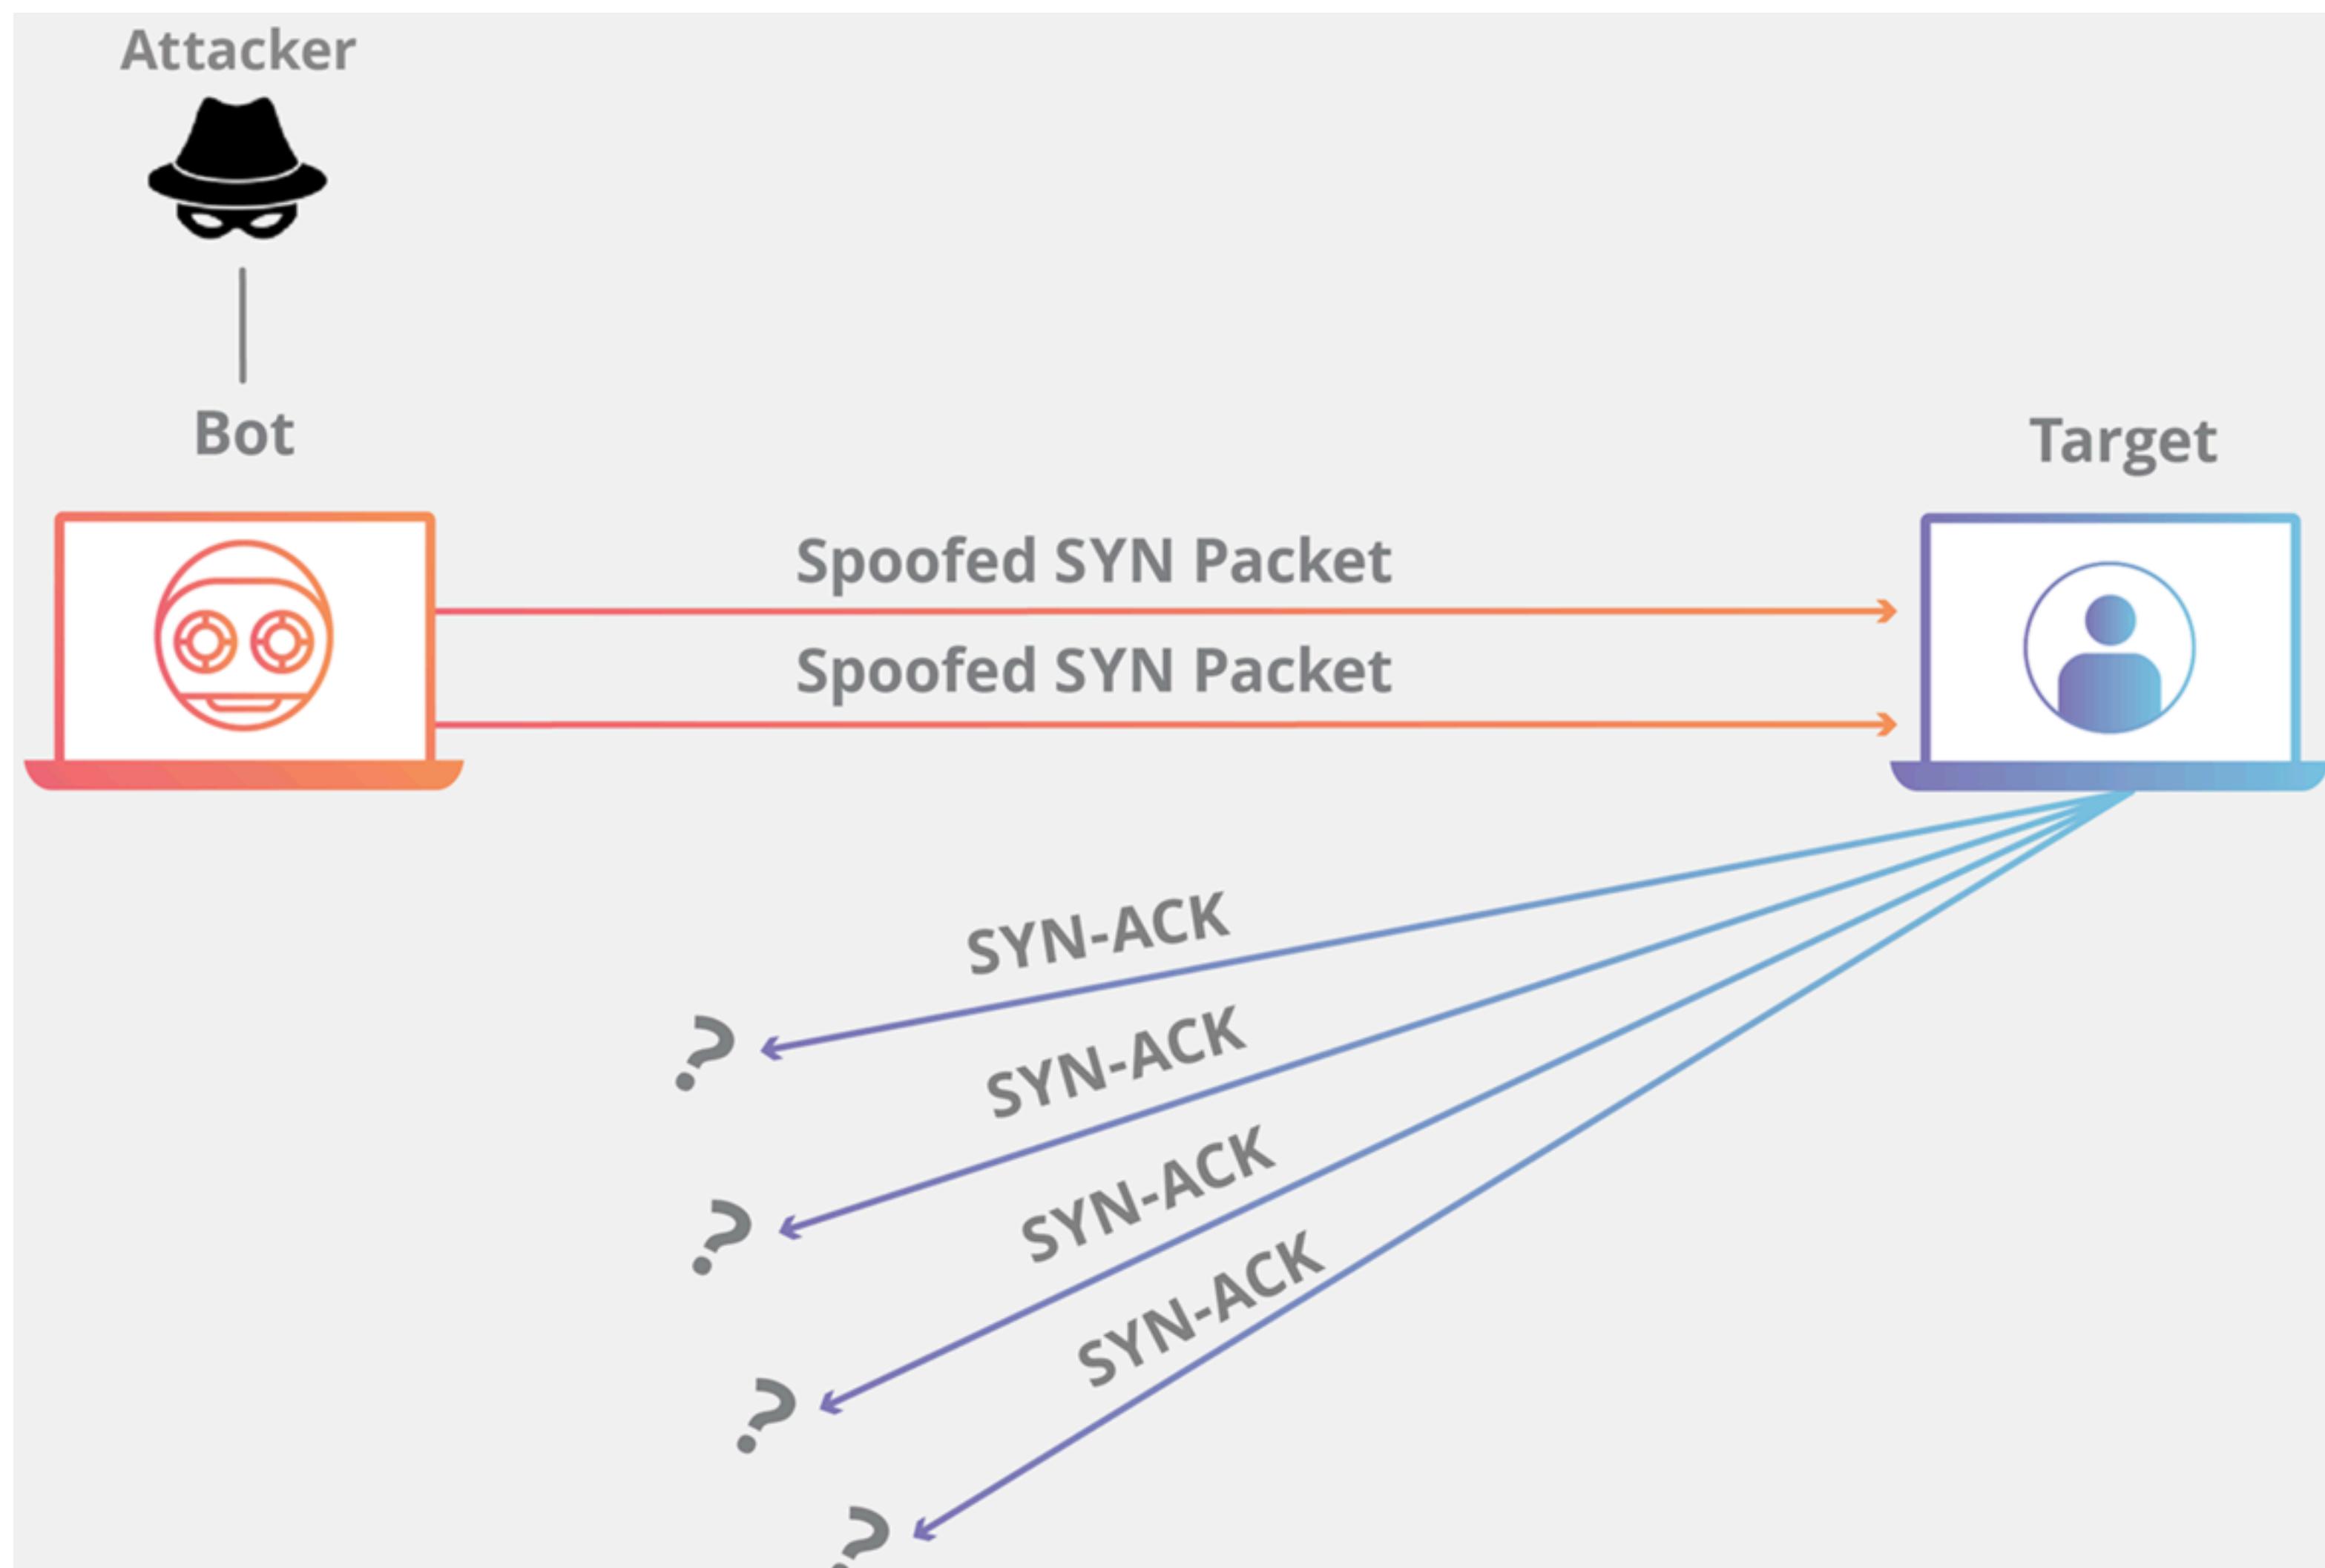
TCP/UDP: require servers to maintain large number of concurrent connections or state

Application Layer: require servers to perform expensive queries or cryptographic operations

TCP Handshake



SYN Floods

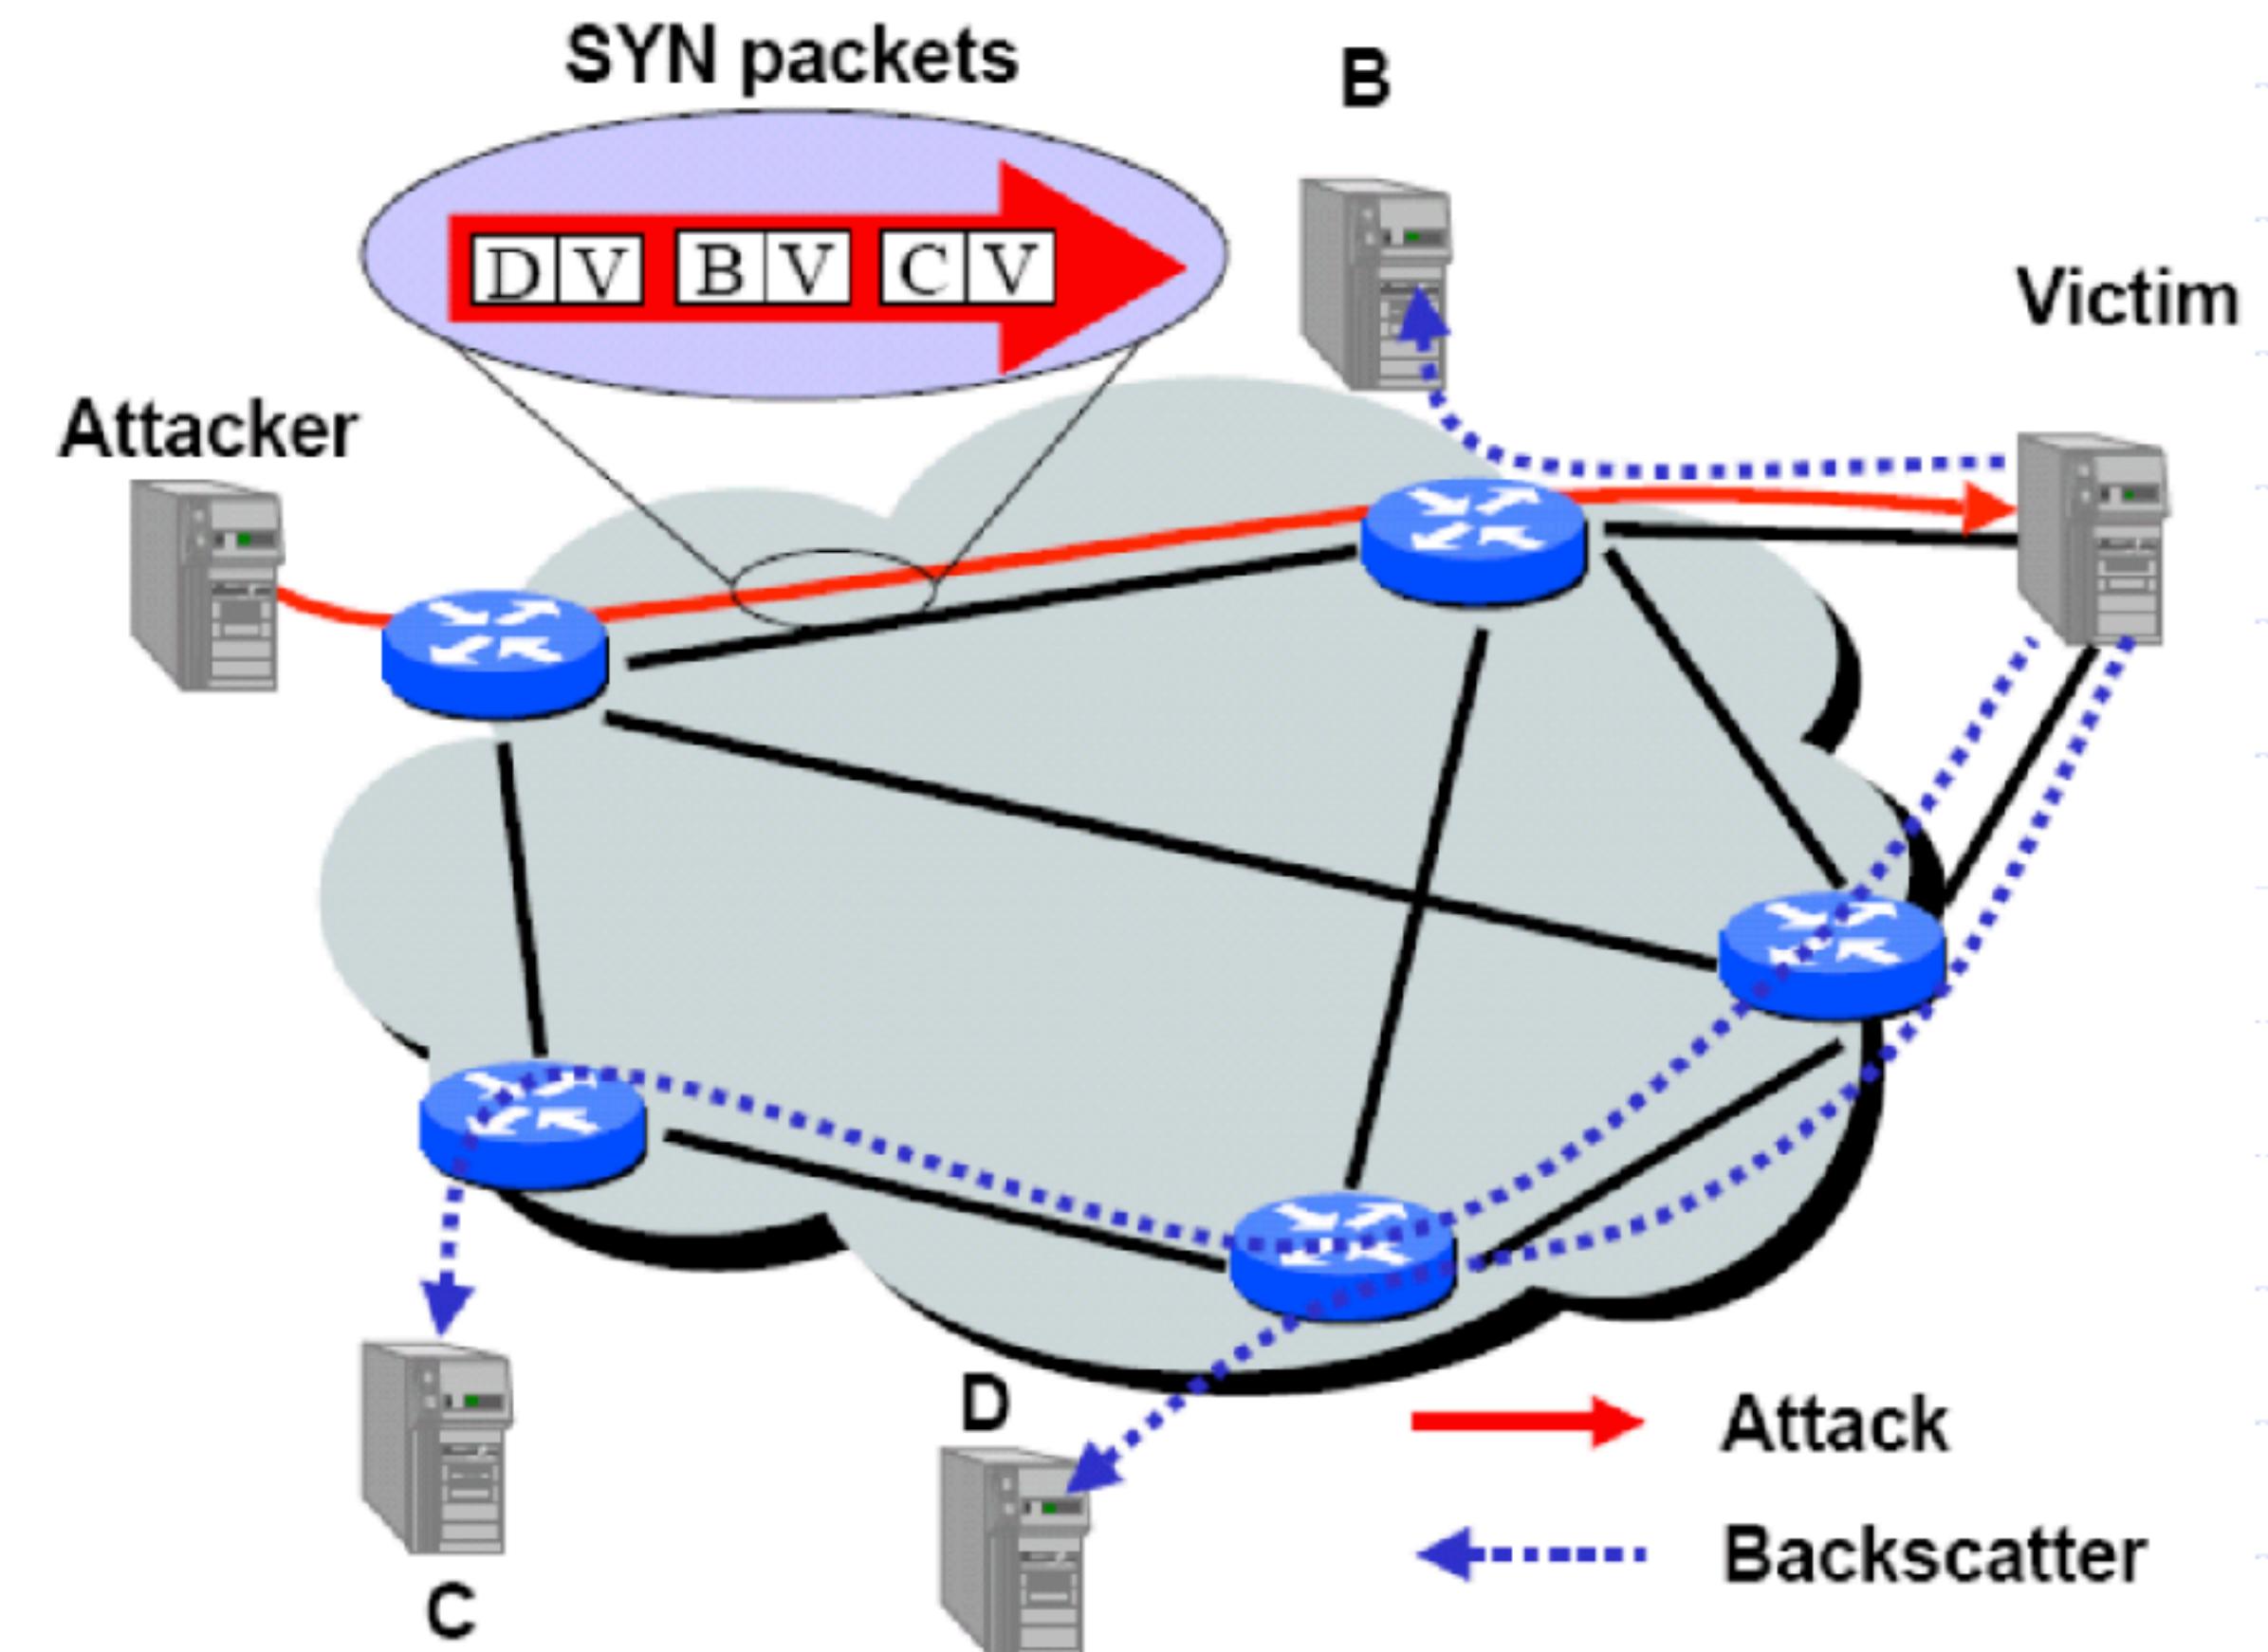


Backscatter

SYN with forged source IP ->
SYN/ACK to random host

Listen to unused IP addresss
space (darknet)

Lonely SYN/ACK packet likely to
be result of SYN attack



Core Problem

Problem: server commits resources (memory) before confirming identify of the client (when client responds)

Bad Solution:

- Increase backlog queue size
- Decrease timeout

Real Solution: Avoid state until 3-way handshake completes

SYN Cookies

Idea: Instead of storing SN_c and SN_s ...
send a cookie back to the client.

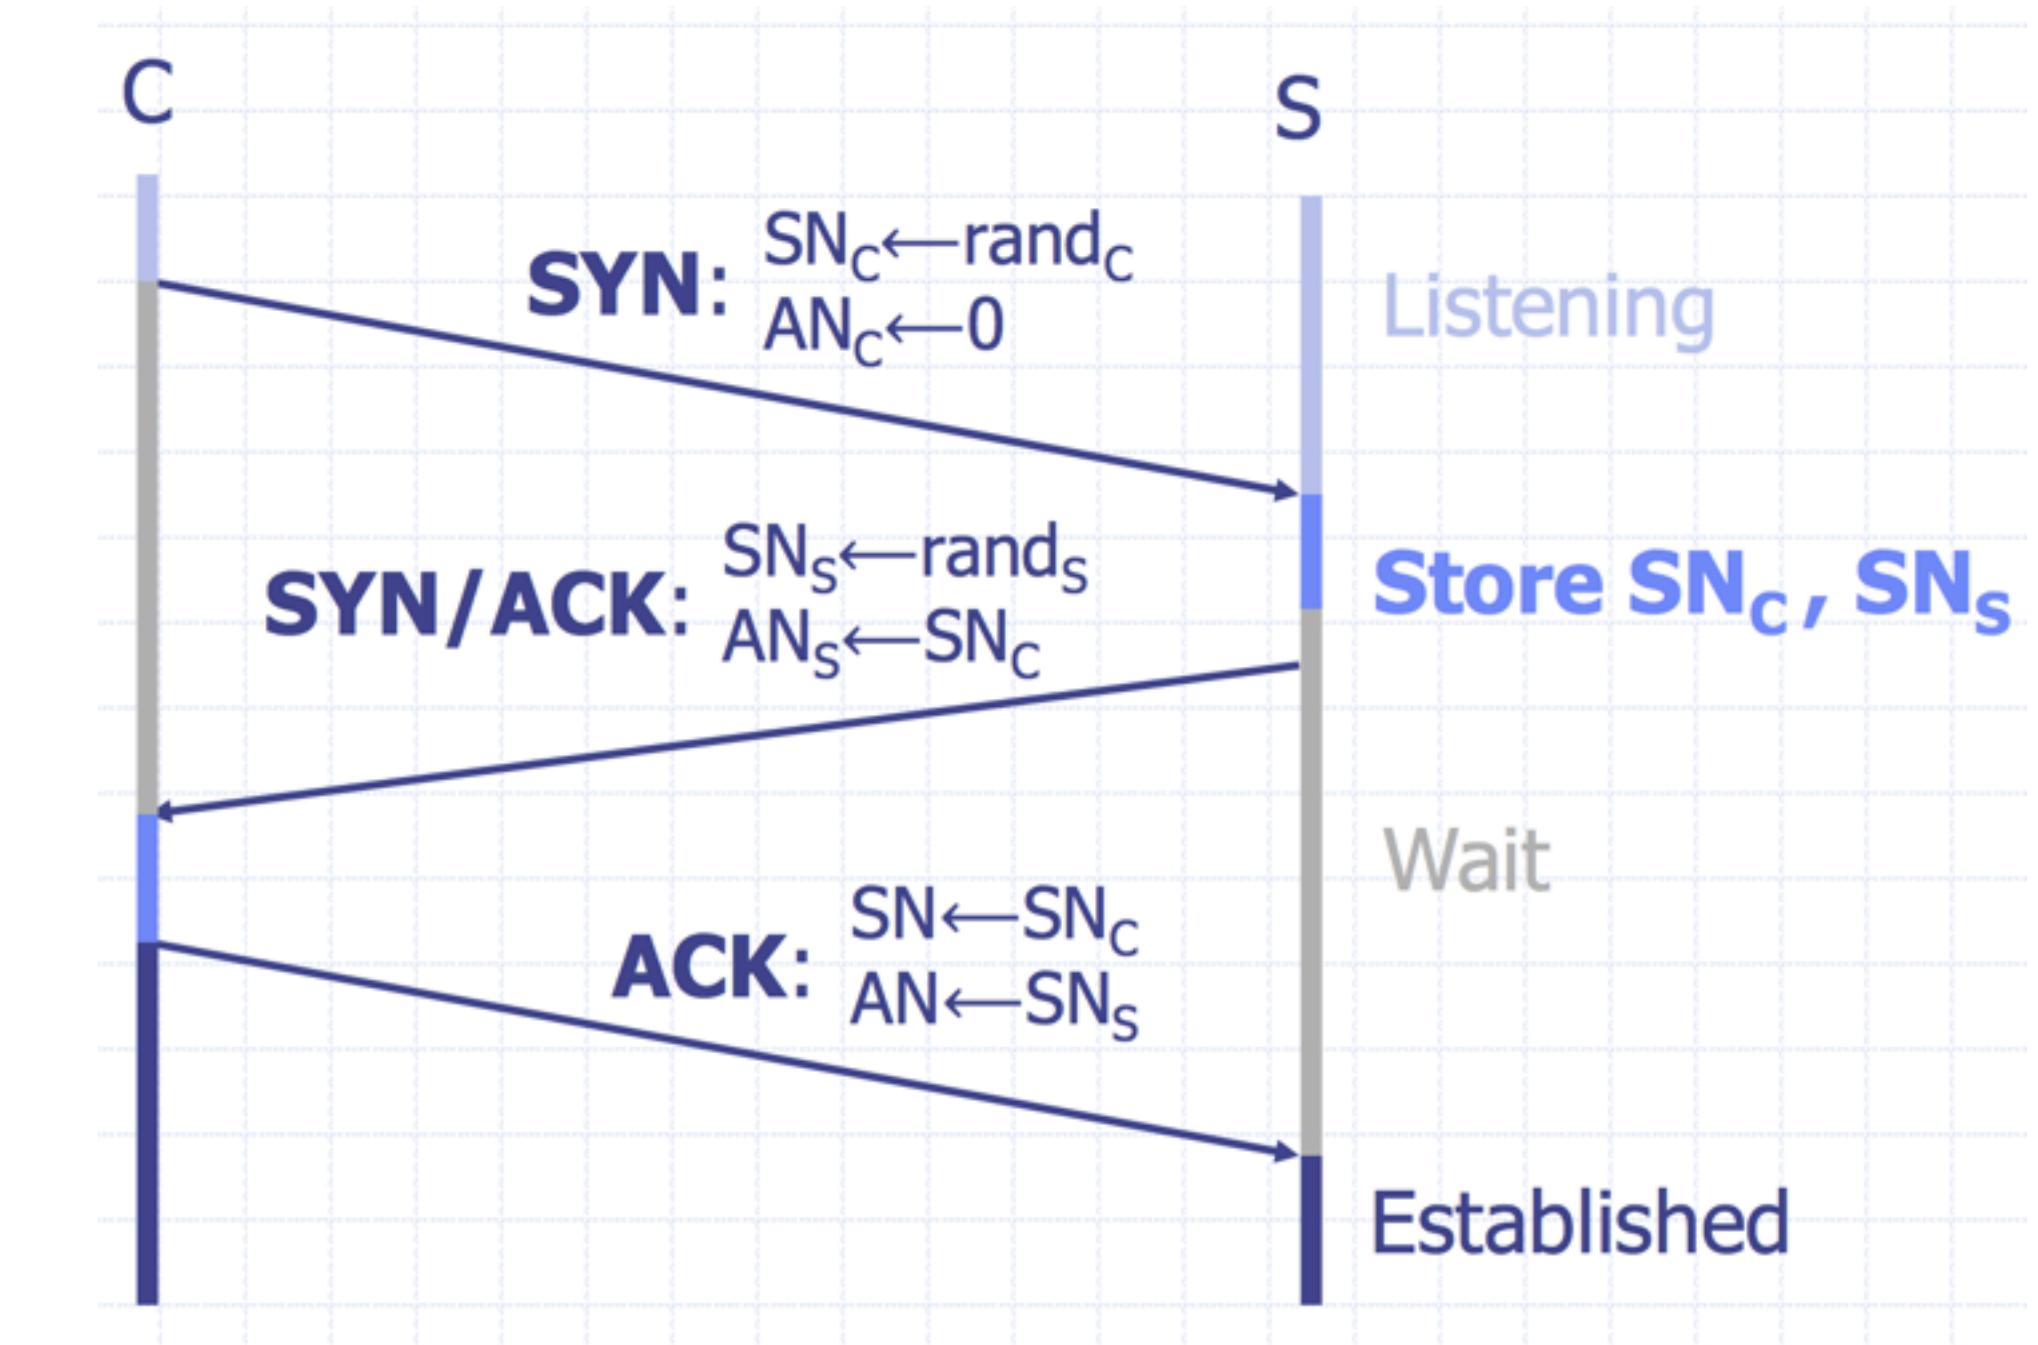
$L = MAC_{key} (SAddr, SPort, DAddr, DPort, SN_c, T)$
key: picked at random during boot

$T = 5$ -bit counter incremented every 64 secs.

$SN_s = (T \parallel mss \parallel L)$

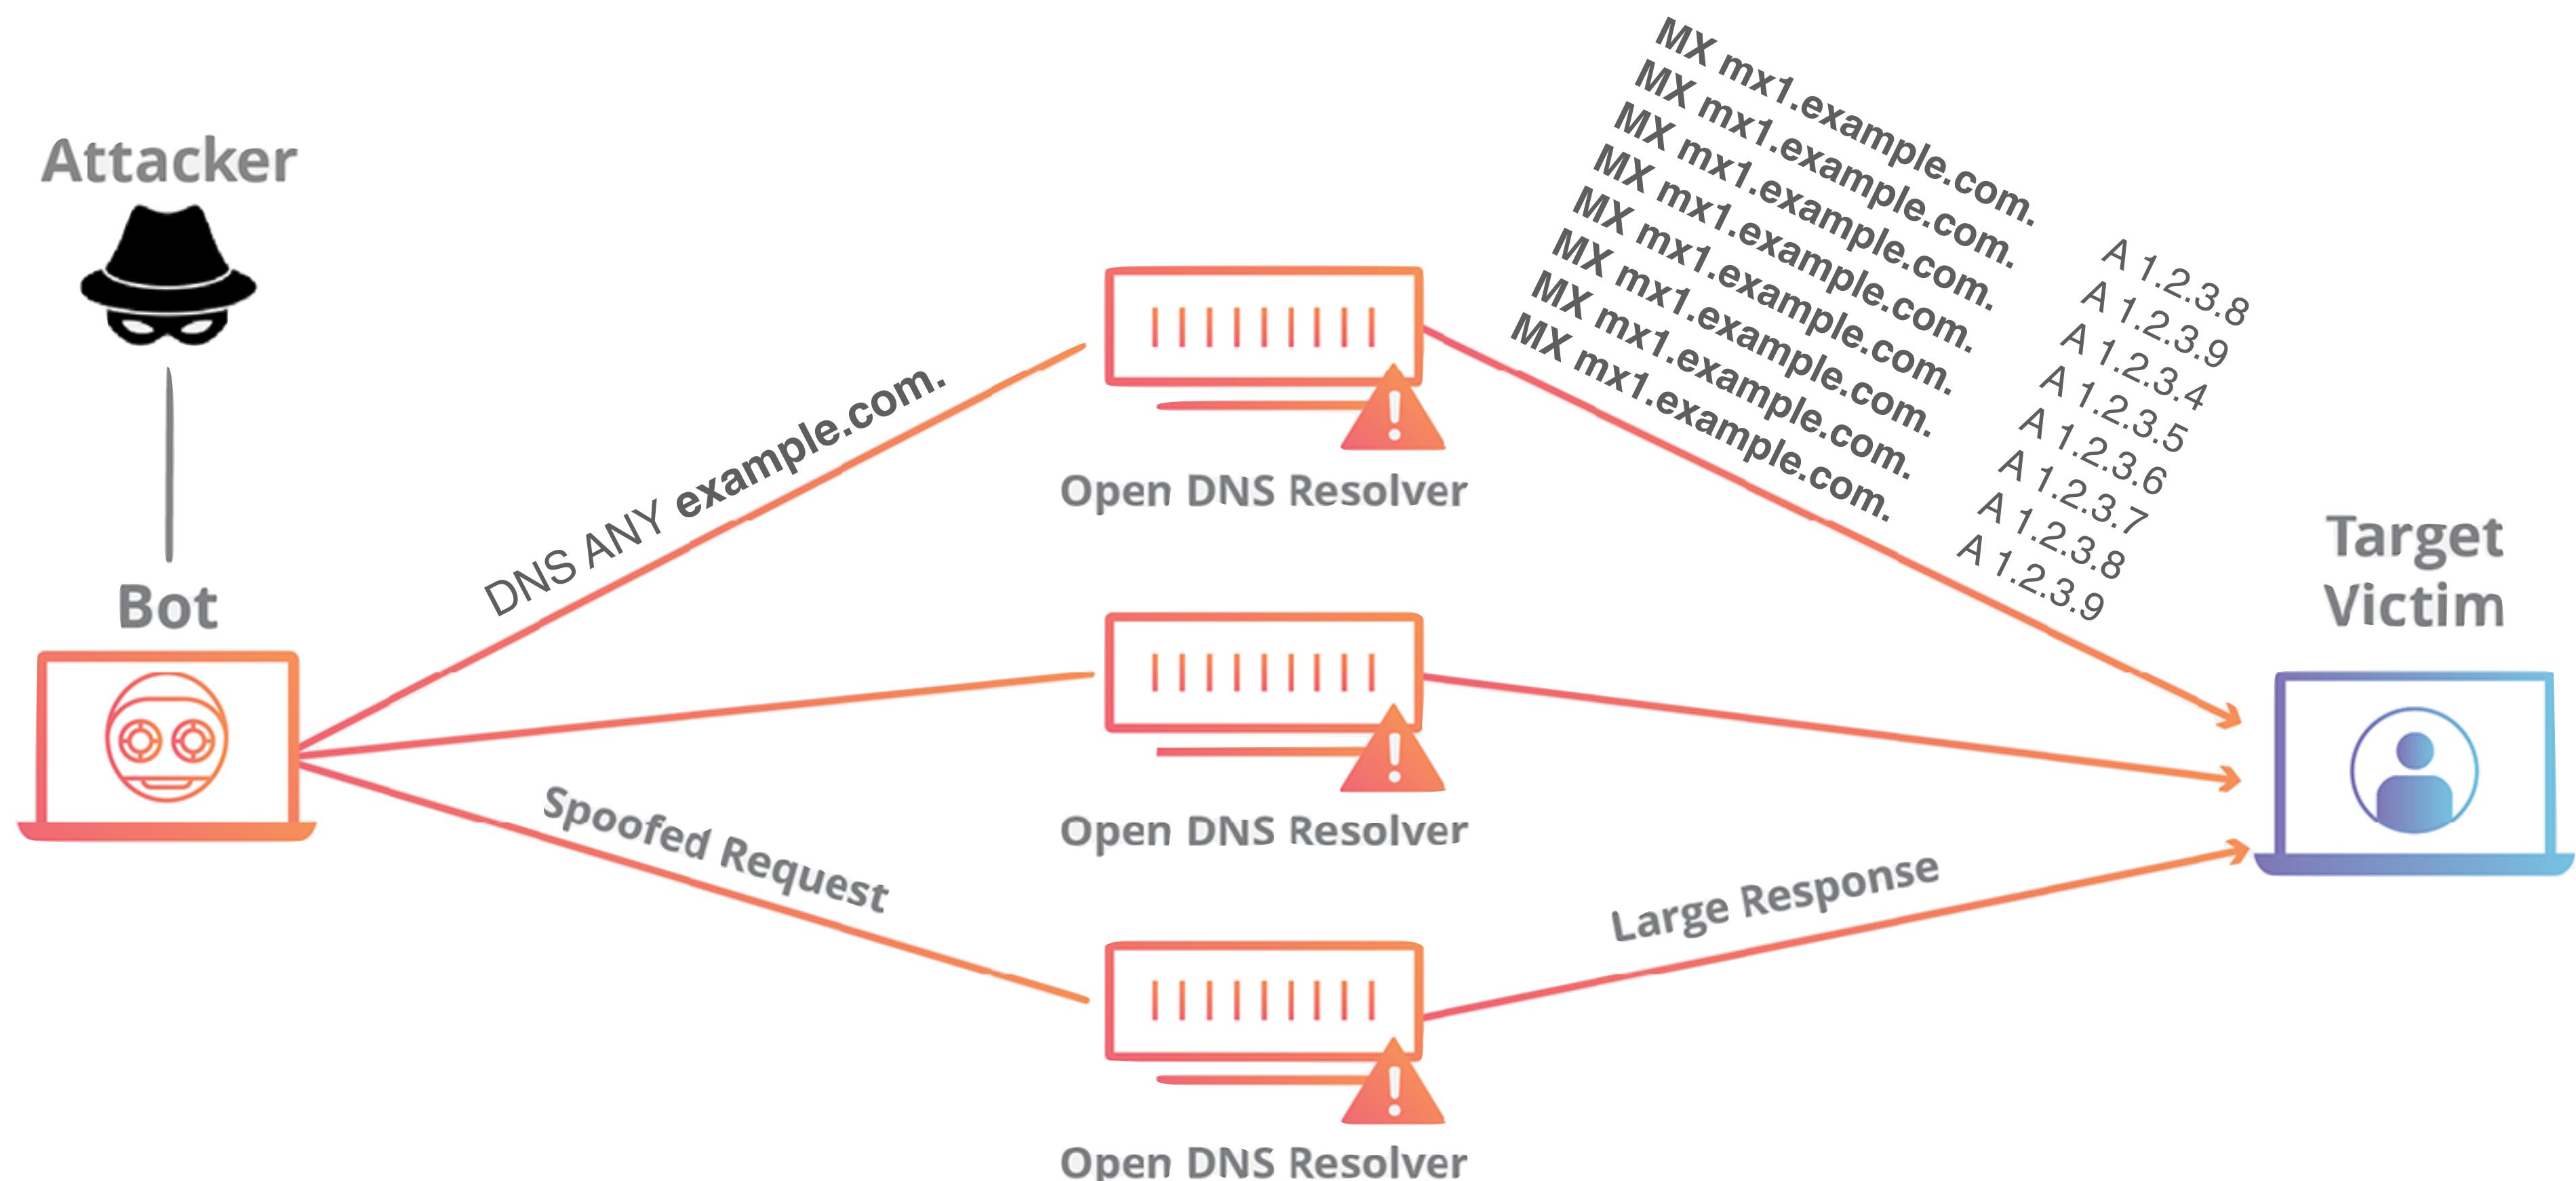
Honest client sends ACK ($AN=SN_s$, $SN=SN_c+1$)

Server allocates space for socket only if valid SNs



Server does not save state
(loses TCP options)

Amplification Attacks



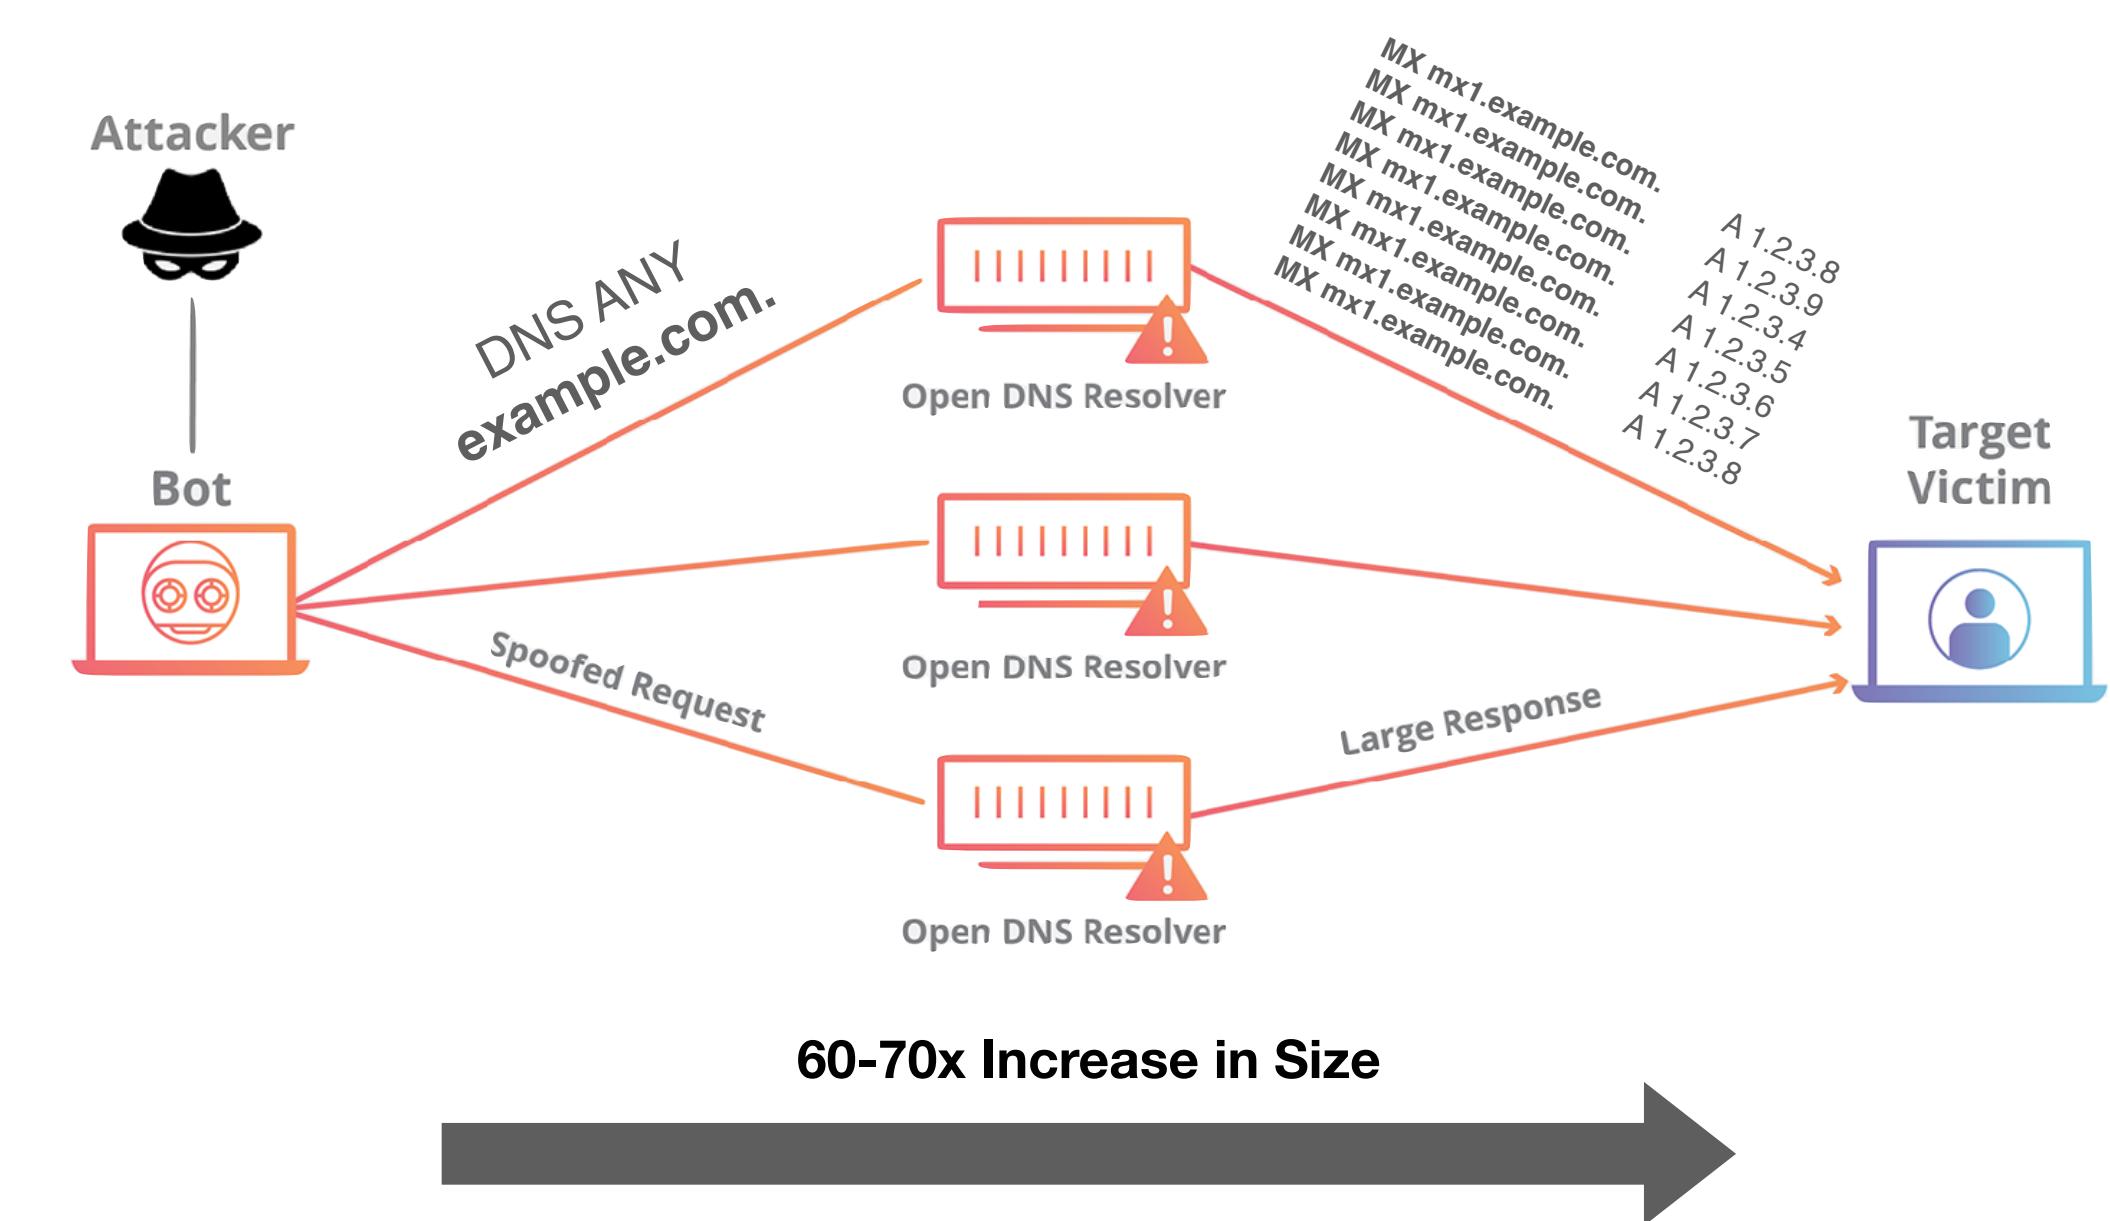
60-70x Increase in Size

Amplification Attacks

Services that respond to a single (small) UDP packet with a large UDP packet can be used to amplify DOS attacks

Attacker forges packet and sets source IP to victim's IP address. When service responds, it sends large amount of data to the spoofed victim

The attacker needs a large number of these services to amplify packets. Otherwise, the victim could just drop the packets from the small number of hosts



Common UDP Amplifiers

DNS: ANY query returns *all* records server has about a domain

NTP: MONLIST returns list of last 600 clients who asked for the time recently

DNS: Do not have recursive resolvers on the public Internet.

NTP: Do not respond to commands like MONLIST

Both are considered misconfigurations today, but often 100Ks of misconfigured hosts on the public Internet

Amplification Attacks

2013: DDoS attack generated 300 Gbps (DNS)

- 31,000 misconfigured open DNS resolvers, each at 10 Mbps
- Source: **3 networks that allowed IP spoofing**

2014: 400 Gbps DDoS attacked used 4,500 NTP servers

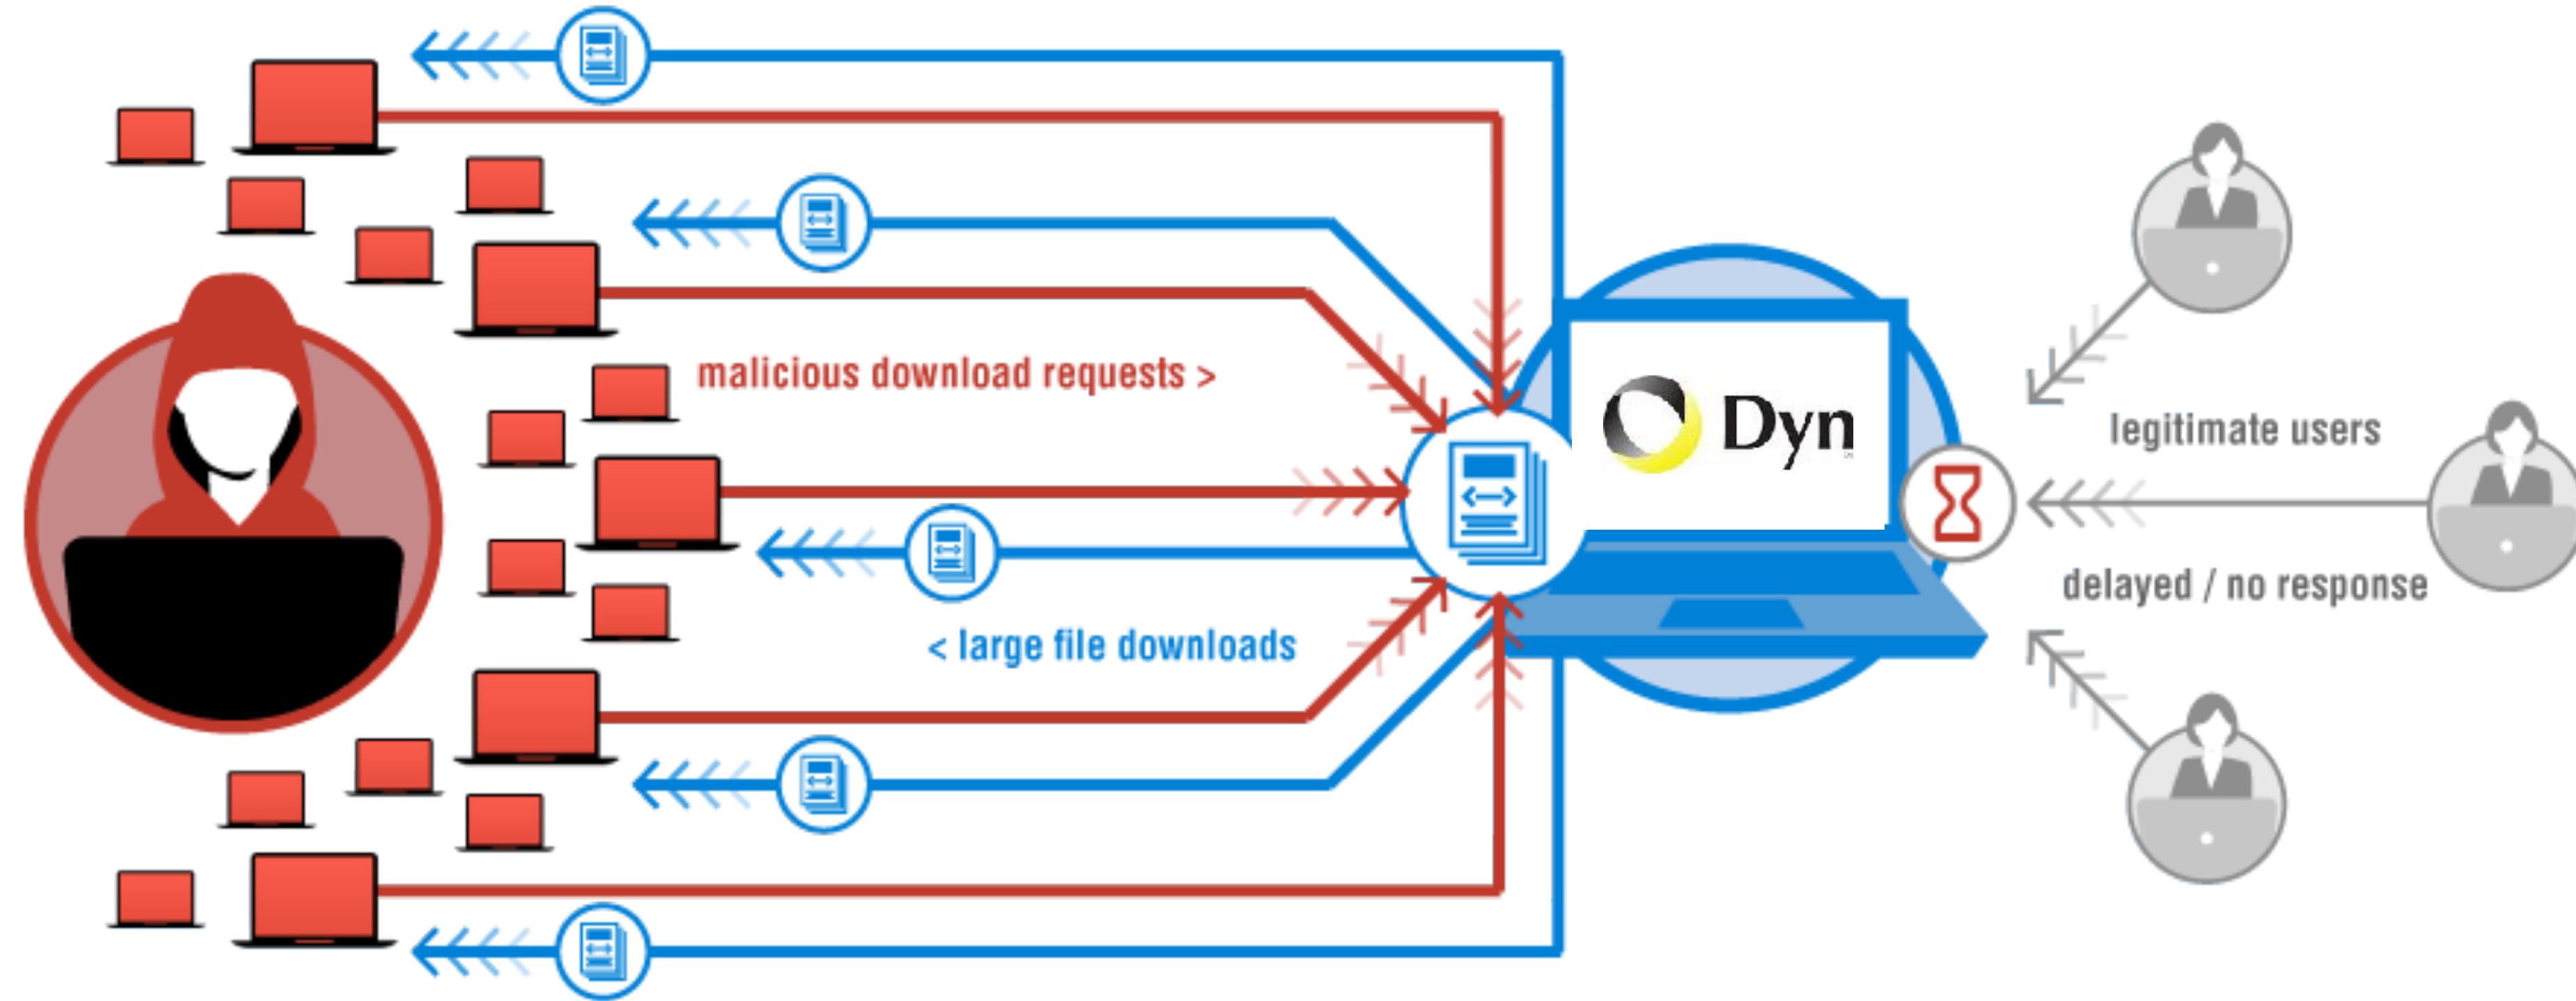
THE WALL STREET JOURNAL.

October 21, 2016

Cyberattack Knocks Out Access to Websites

Popular sites such as Twitter, Netflix and PayPal were unreachable for part of the day





“We are still working on analyzing the data but the estimate at the time of this report is up to 100,000 malicious endpoints. [...] There have been some reports of a magnitude in the 1.2 Tbps range; at this time we are unable to verify that claim.”

A Botnet of IoT Devices



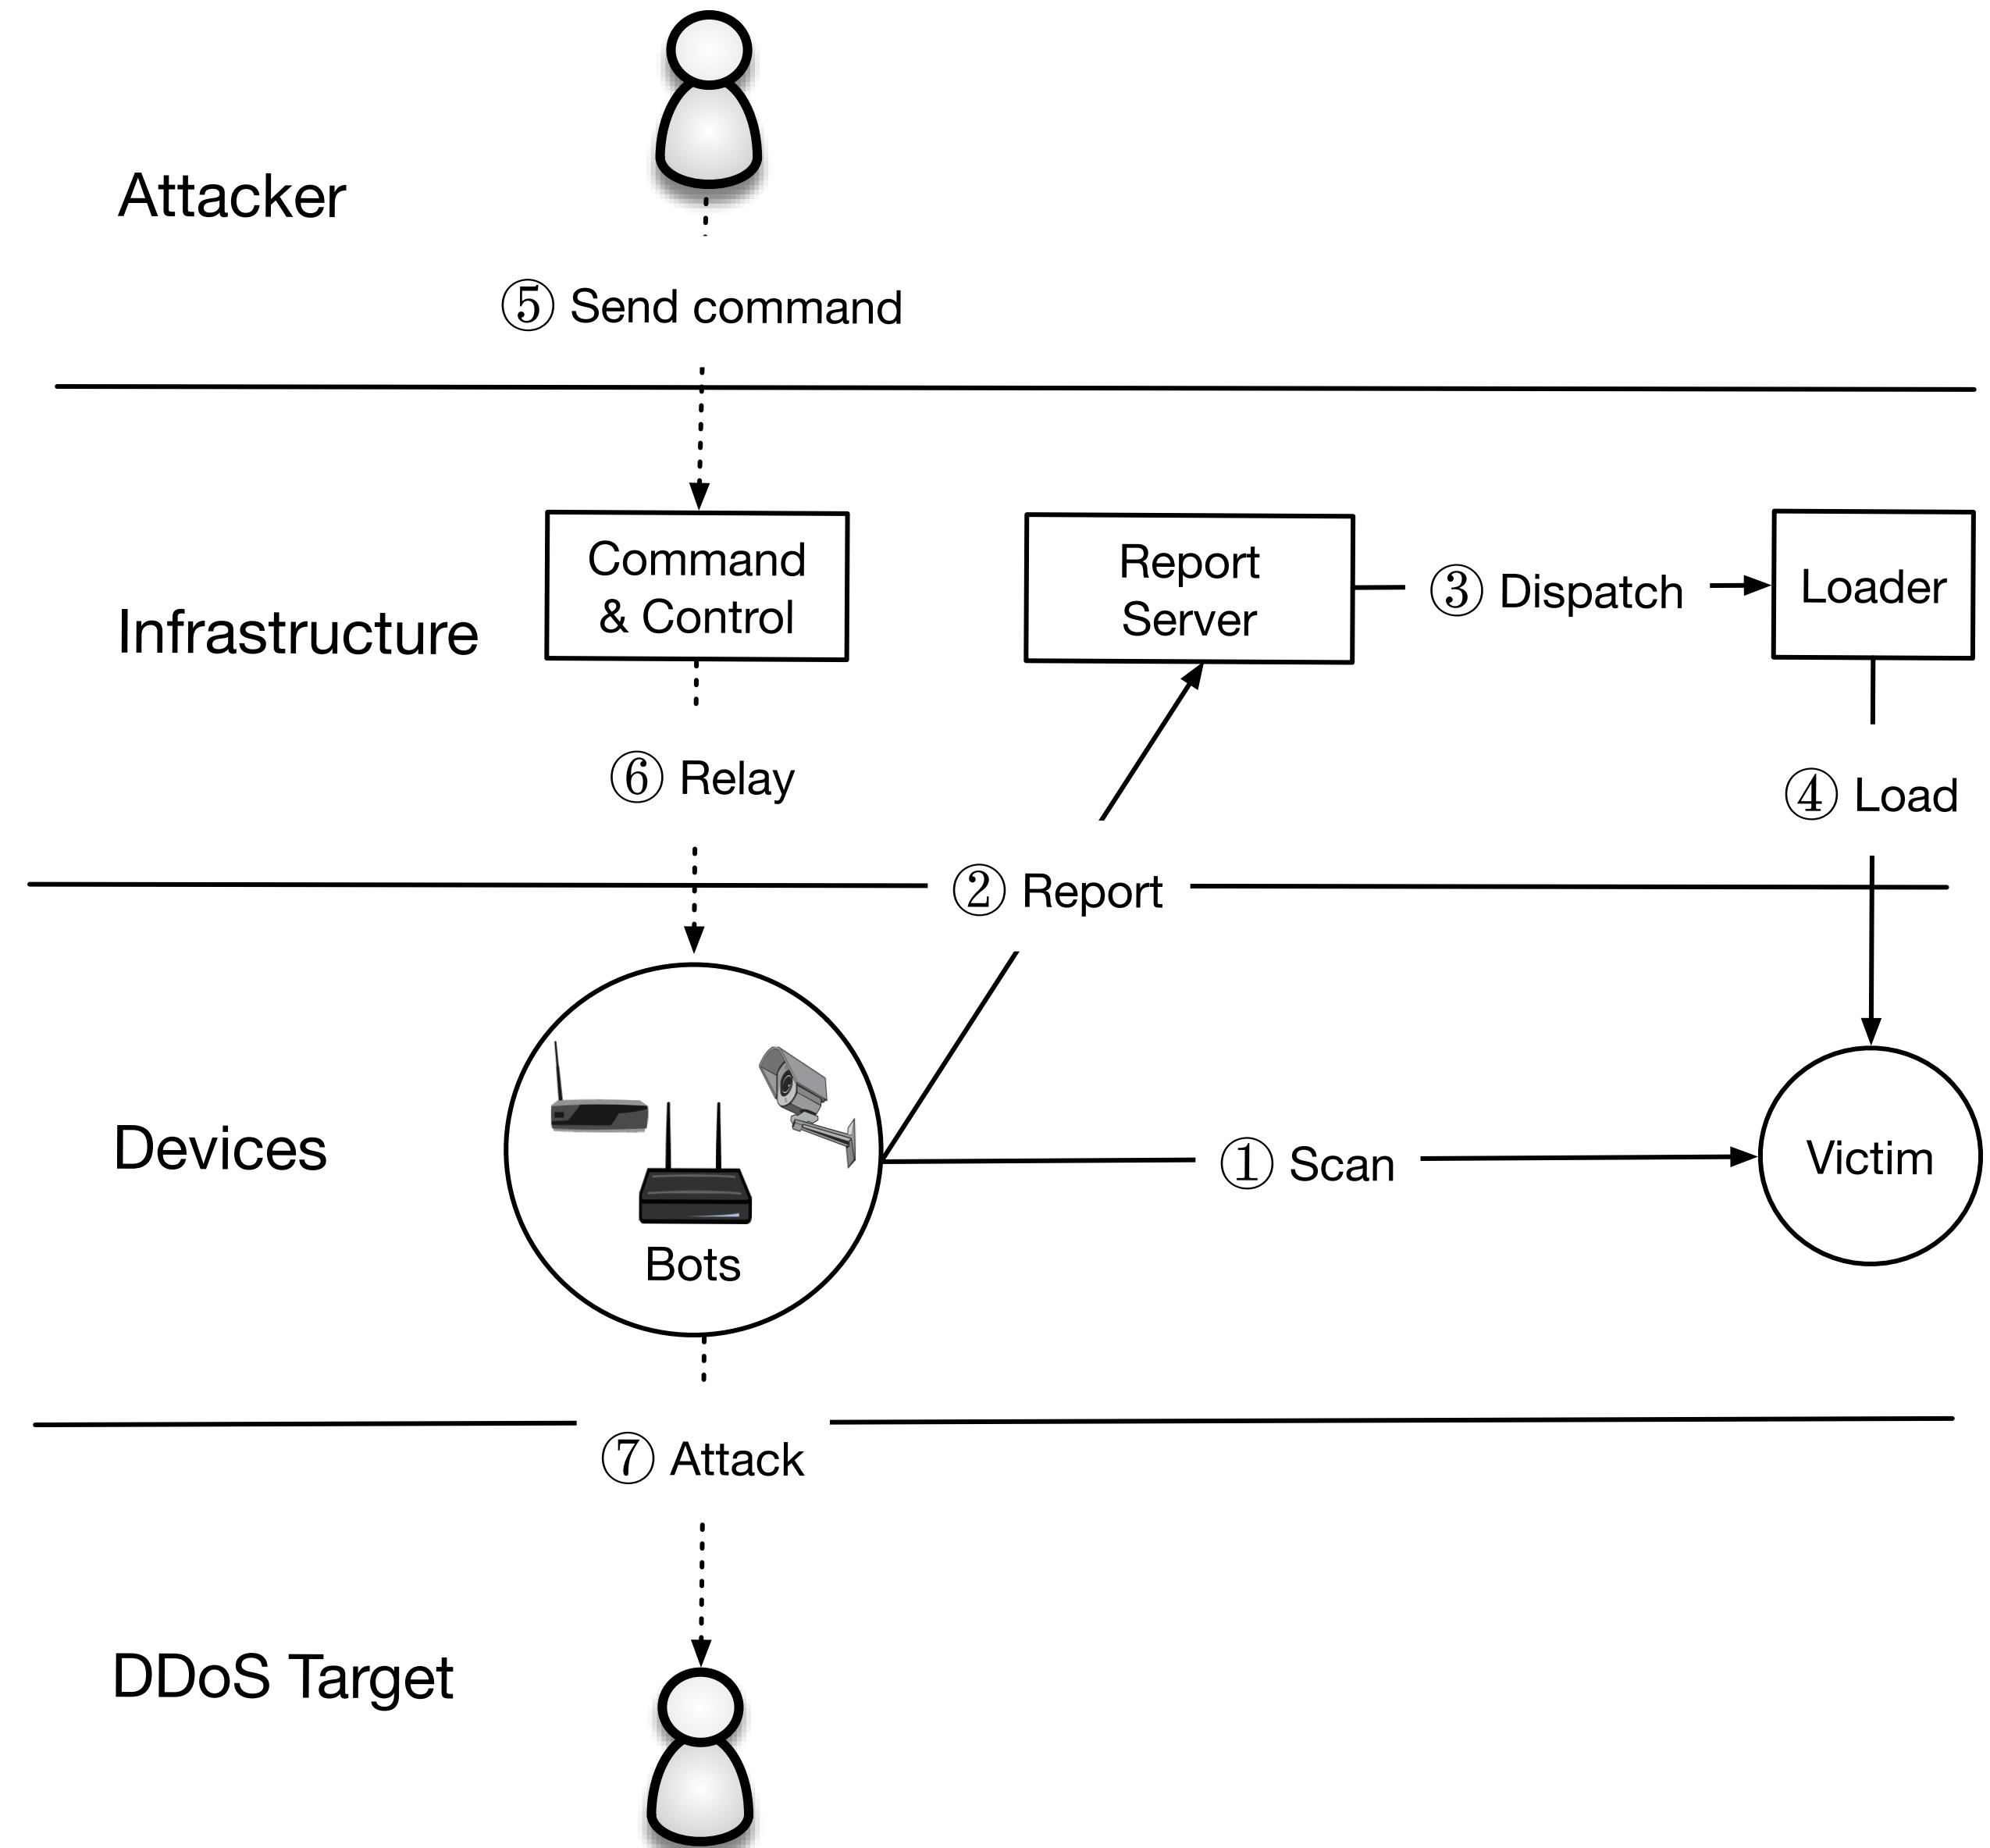
Not Amplification.
Flood with SYN, ACK, UDP, and GRE packets

The Mirai Malware

Bot master will issue commands to scan or start an attack

Attack Command:

- Action (e.g., START, STOP)
- Target IP(s)
- Attack Type (e.g., GRE, DNS, TCP)
- Attack Duration

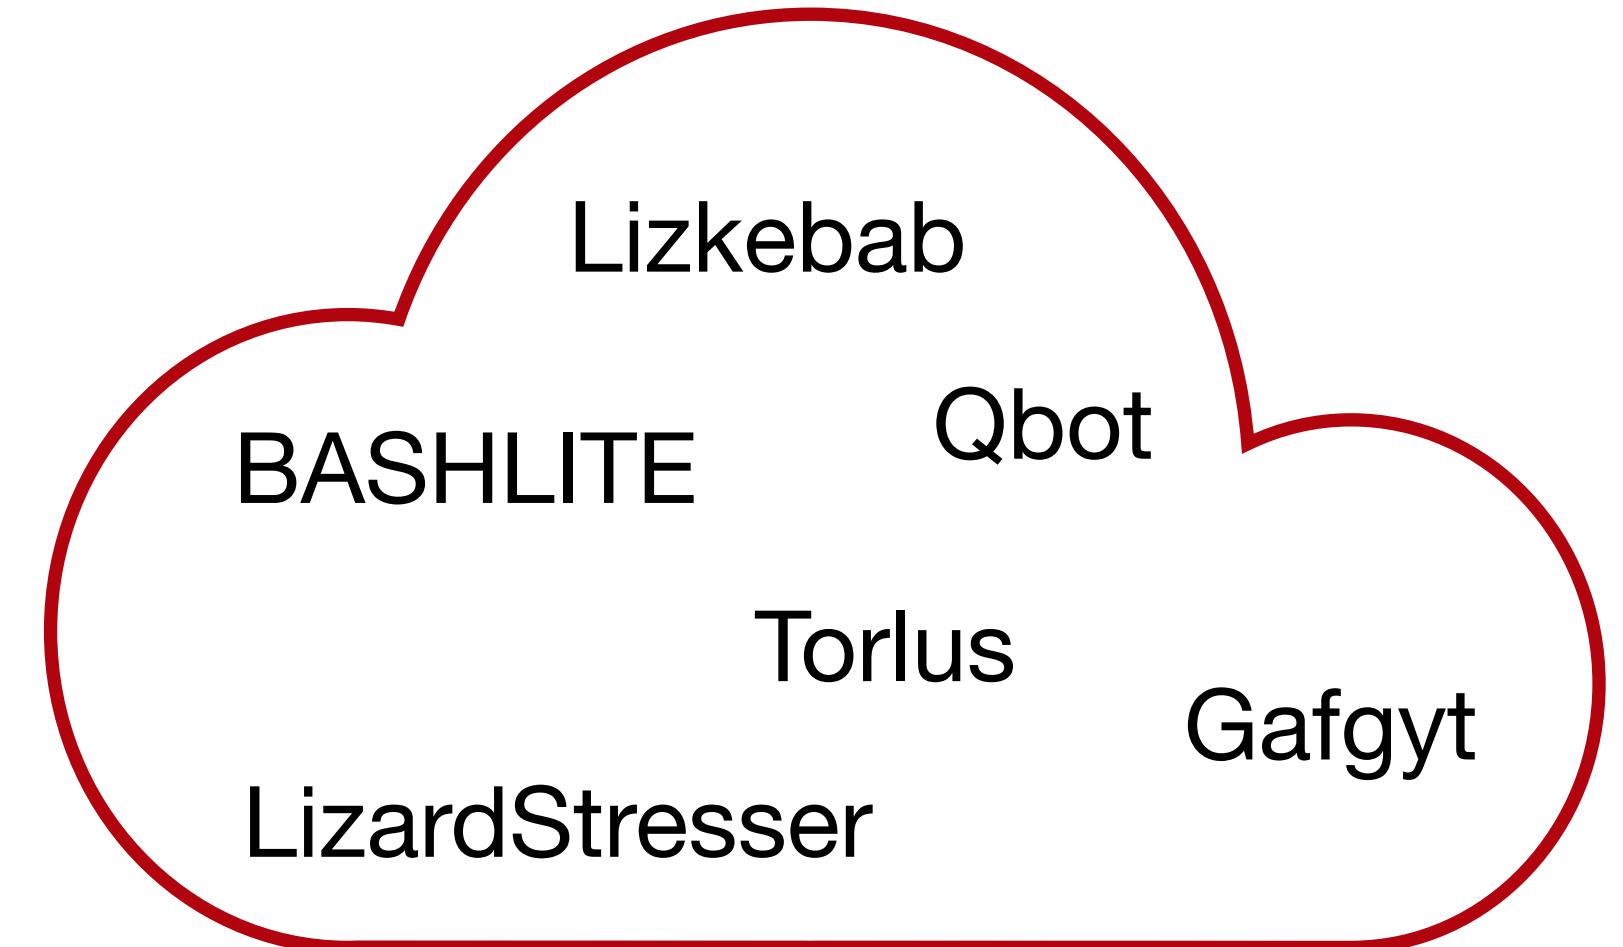


What made Mirai Successful?

The Mirai malware is (astoundingly) badly written. It uses no new or complex techniques.

Mirai was successful because:

1. IoT security bar is very low
2. Attack simplicity enabled the malware to compromise heterogeneous hardware
3. Stateless scanning was an improvement over prior versions

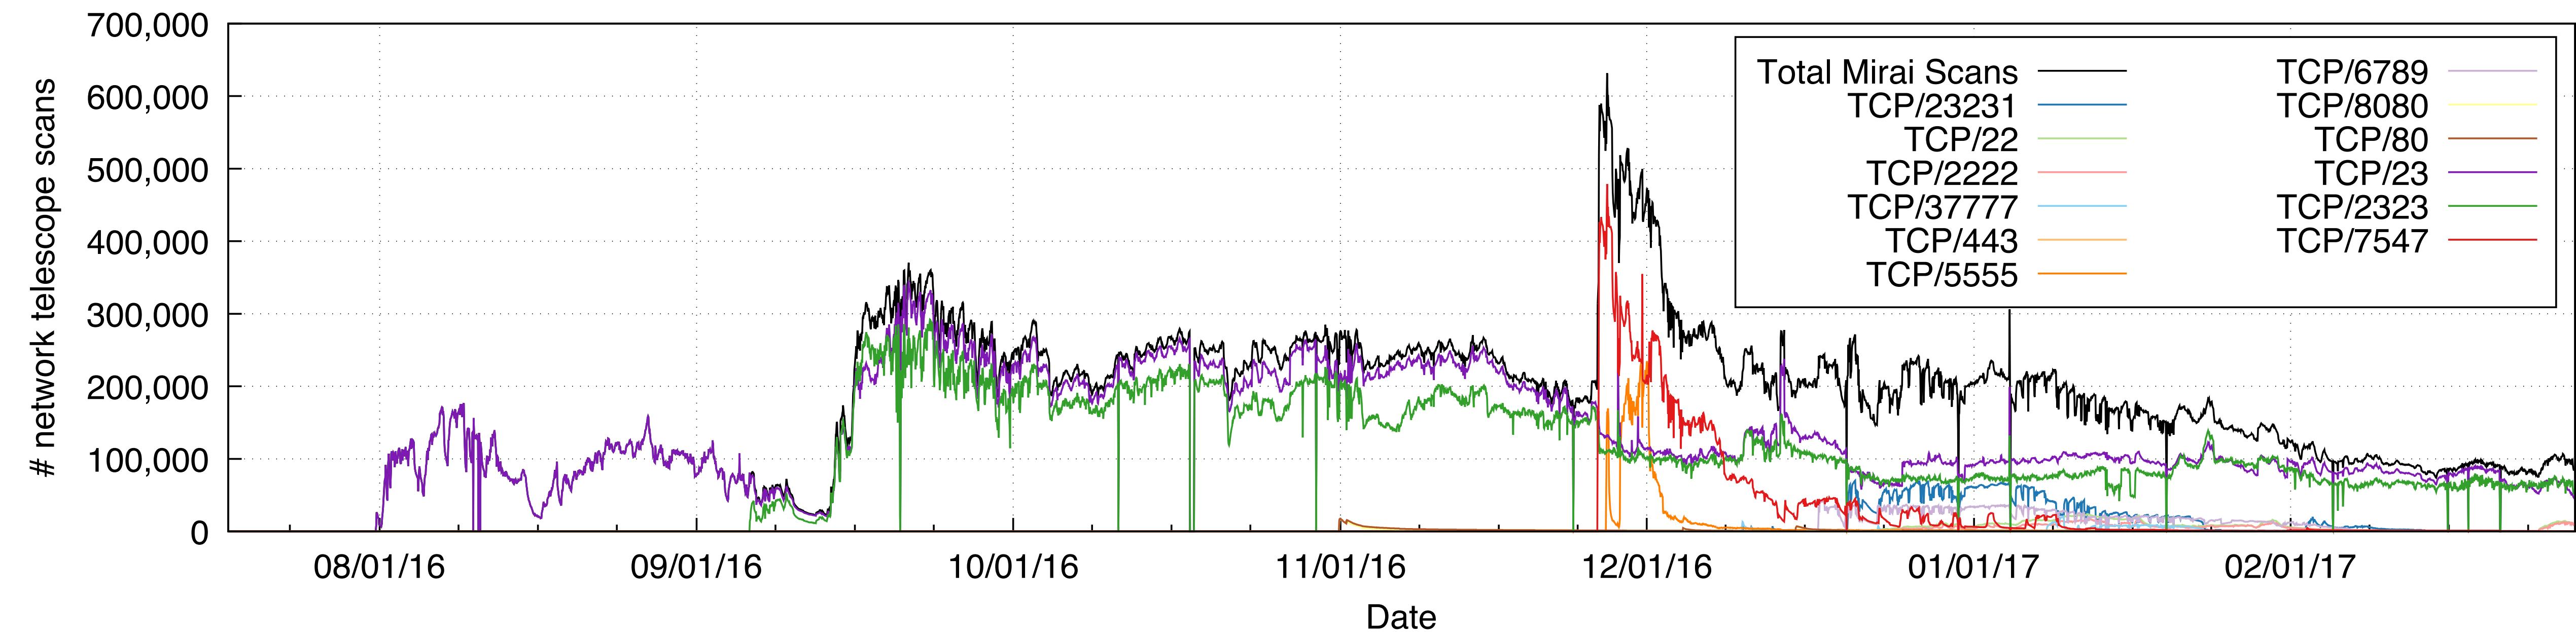


Mirai

Password Guessing

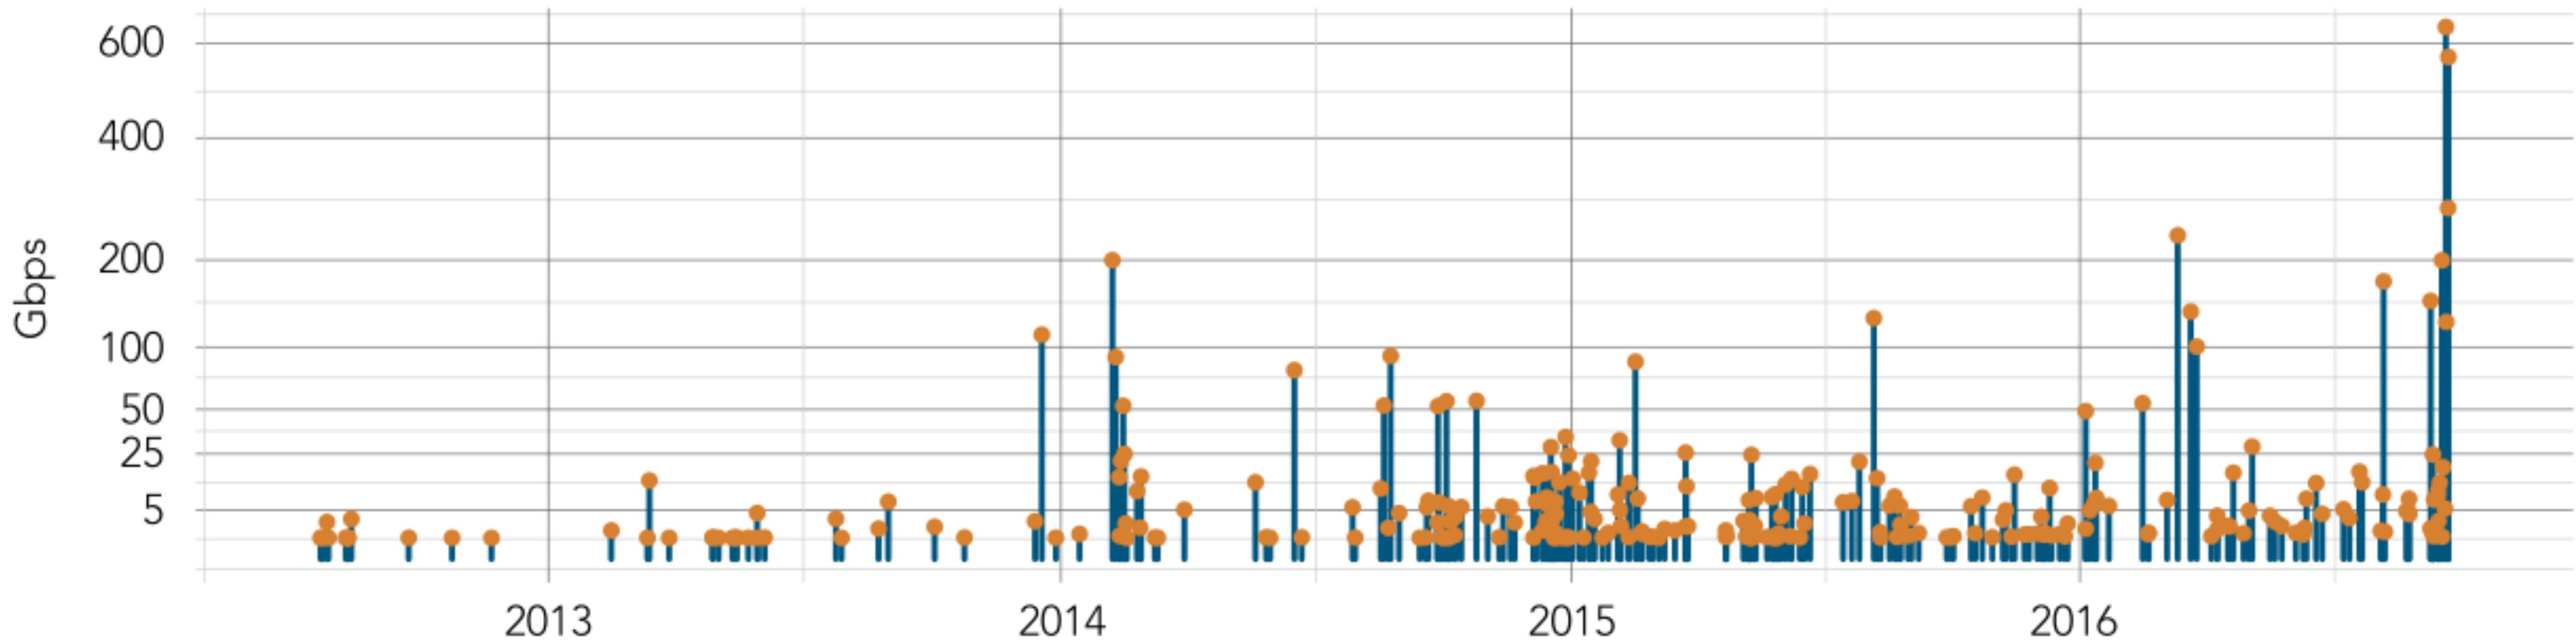
Password	Device Type	Password	Device Type	Password	Device Type
123456	ACTi IP Camera	klv1234	HiSilicon IP Camera	1111	Xerox Printer
anko	ANKO Products DVR	jvbzd	HiSilicon IP Camera	Zte521	ZTE Router
pass	Axis IP Camera	admin	IPX-DDK Network Camera	1234	Unknown
888888	Dahua DVR	system	IQinVision Cameras	12345	Unknown
666666	Dahua DVR	meinsm	Mobotix Network Camera	admin1234	Unknown
vizxv	Dahua IP Camera	54321	Packet8 VOIP Phone	default	Unknown
7ujMko0vizxv	Dahua IP Camera	00000000	Panasonic Printer	fucker	Unknown
7ujMko0admin	Dahua IP Camera	realtek	RealTek Routers	guest	Unknown
666666	Dahua IP Camera	1111111	Samsung IP Camera	password	Unknown
dreambox	Dreambox TV Receiver	xmhdipc	Shenzhen Anran Camera	root	Unknown
juantech	Guangzhou Juan Optical	smcadmin	SMC Routers	service	Unknown
xc3511	H.264 Chinese DVR	ikwb	Toshiba Network Camera	support	Unknown
OxhlwSG8	HiSilicon IP Camera	ubnt	Ubiquiti AirOS Router	tech	Unknown
cat1029	HiSilicon IP Camera	supervisor	VideoIQ	user	Unknown
hi3518	HiSilicon IP Camera	<none>	Vivotek IP Camera	zlxx.	Unknown
klv123	HiSilicon IP Camera				

Mirai Population



~600K devices compromised

DDoS Attacks on Krebs on Security



“The magnitude of the attacks seen during the final week were significantly larger than the majority of attacks Akamai sees on a regular basis. [...] In fact, while the attack on September 20 was the largest attack ever mitigated by Akamai, the attack on September 22 would have qualified for the record at any other time, peaking at 555 Gbps.”

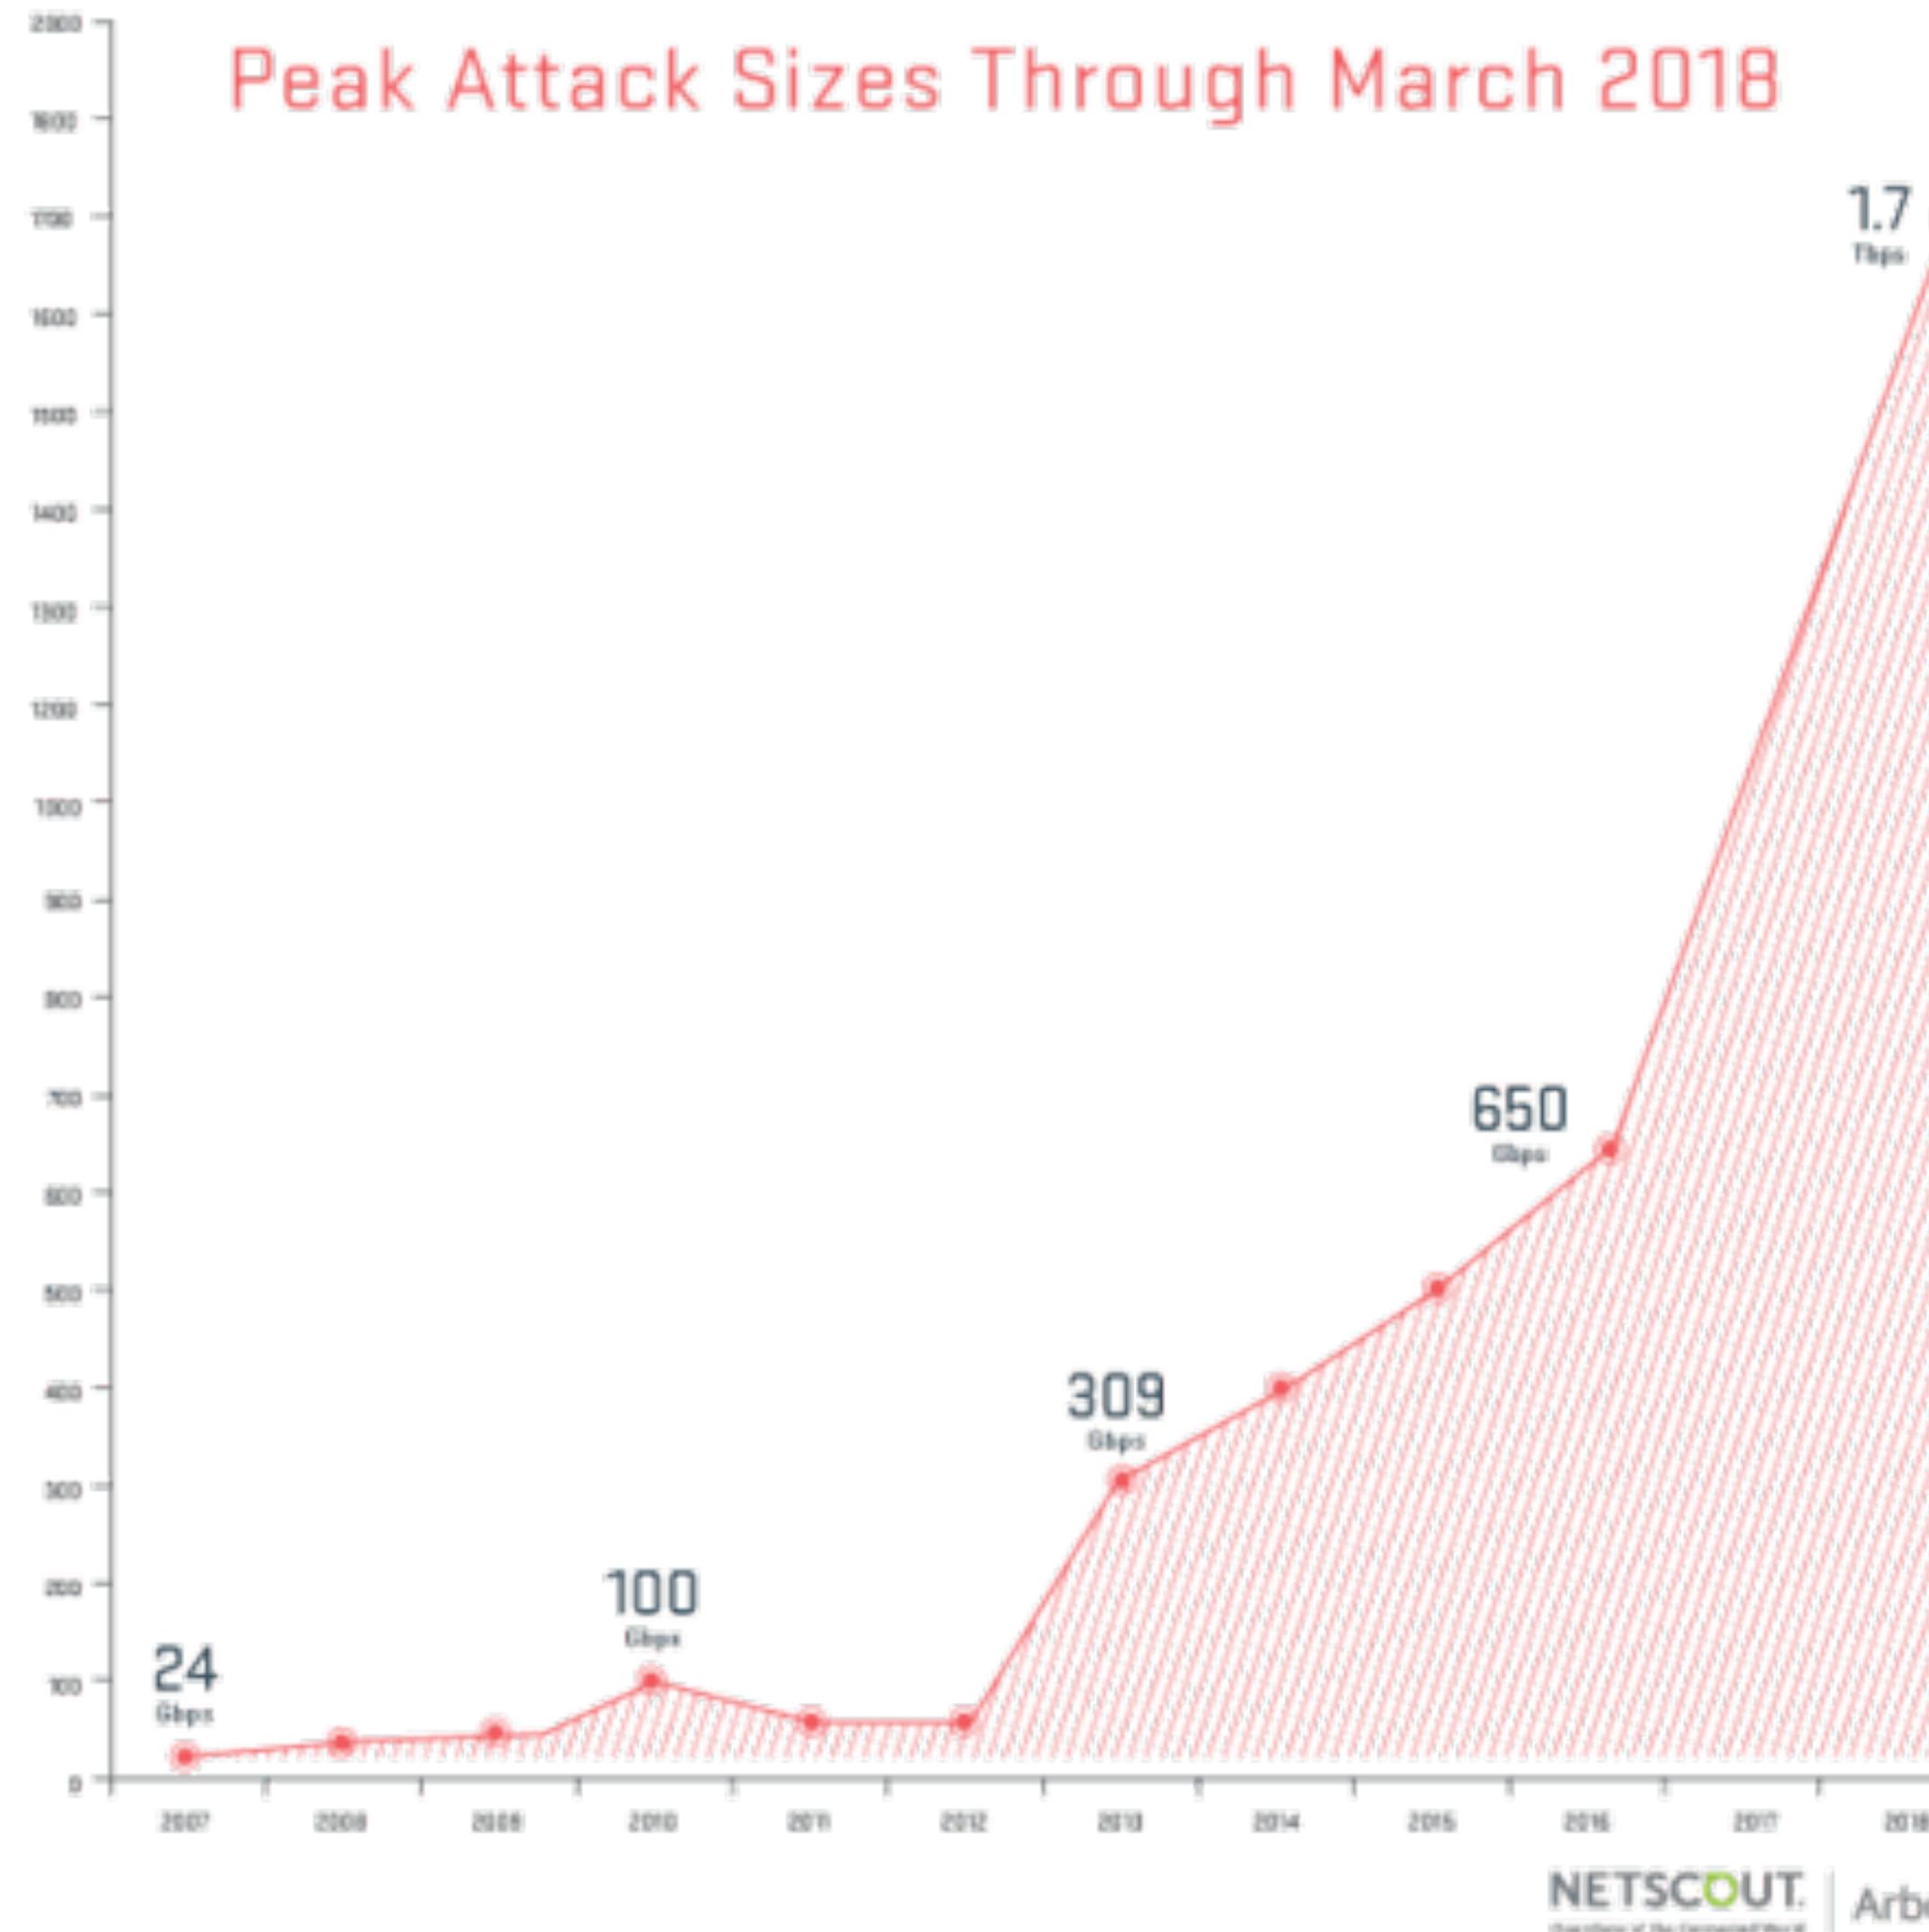
Booster Services

\$23.99	1 month
1 Month Gold	
Time per boot	2400 sec
Concurrents	1
Total network	220Gbps
Tools	Included
Support	24/7
Buy with Paypal	
	

\$34.99	1 month
1 Month Diamond	
Time per boot	3600 sec
Concurrents	2
Total network	220Gbps
Tools	Included
Support	24/7
Buy with Paypal	
	

\$44.99	10 years
Lifetime Bronze	
Time per boot	600 sec
Concurrents	2
Total network	220Gbps
Tools	Included
Support	24/7
Buy with Paypal	
	

Memcache



Memcache: retrieve large record

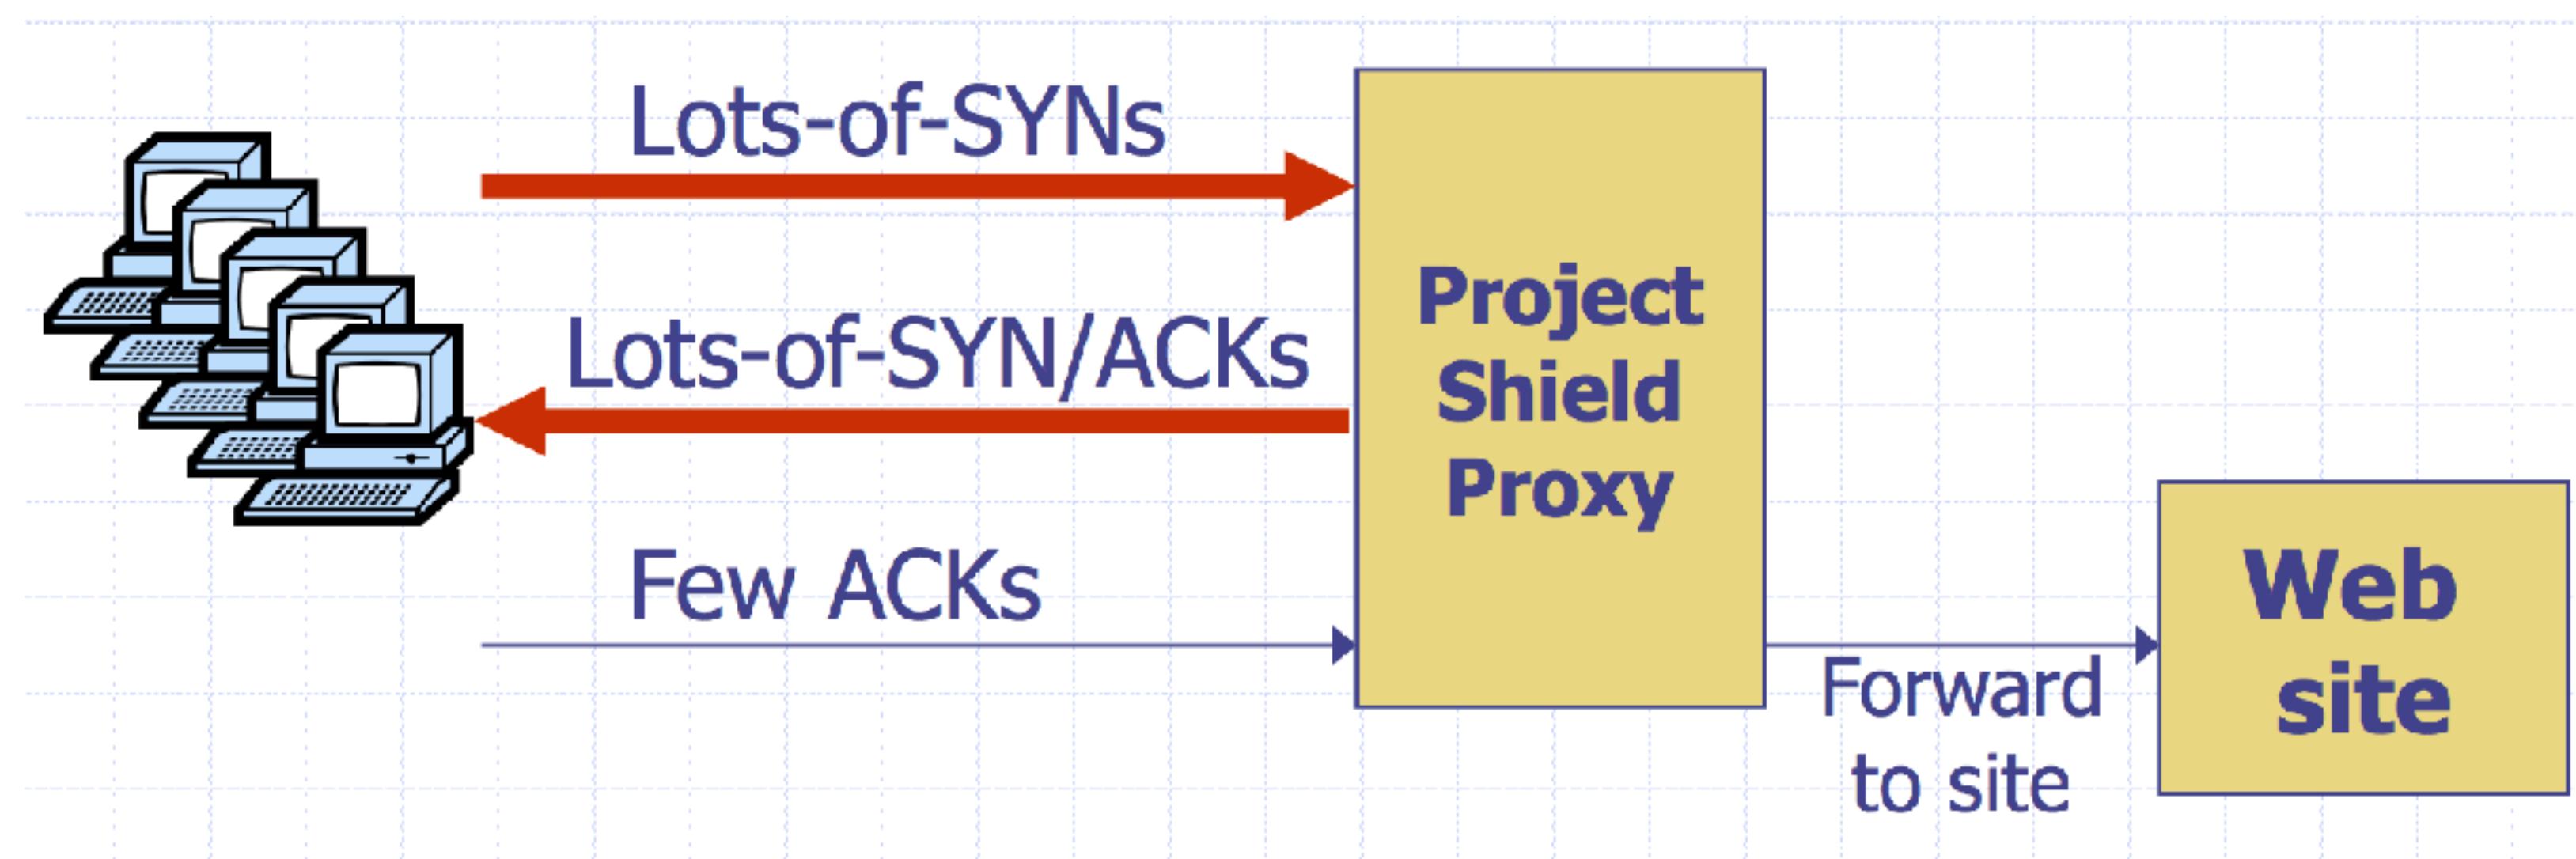
The server responds by firing back as much as 50,000 times the data it received.

Exist both a UDP and TCP version. Only works for UDP! TCP would require a three-way handshake and server would realize IP had been spoofed.

Google Project Shield

DDoS Attacks are often used to censor content. In the case of Mirai, Brian Kreb's blog was under attack.

Google Project shield uses Google bandwidth to shield vulnerable websites (e.g., news, blogs, human rights orgs)



Moving Up Stack: GET Floods

Command bot army to:

- * Complete real TCP connection
- * Complete TLS Handshake
- * GET large image or other content

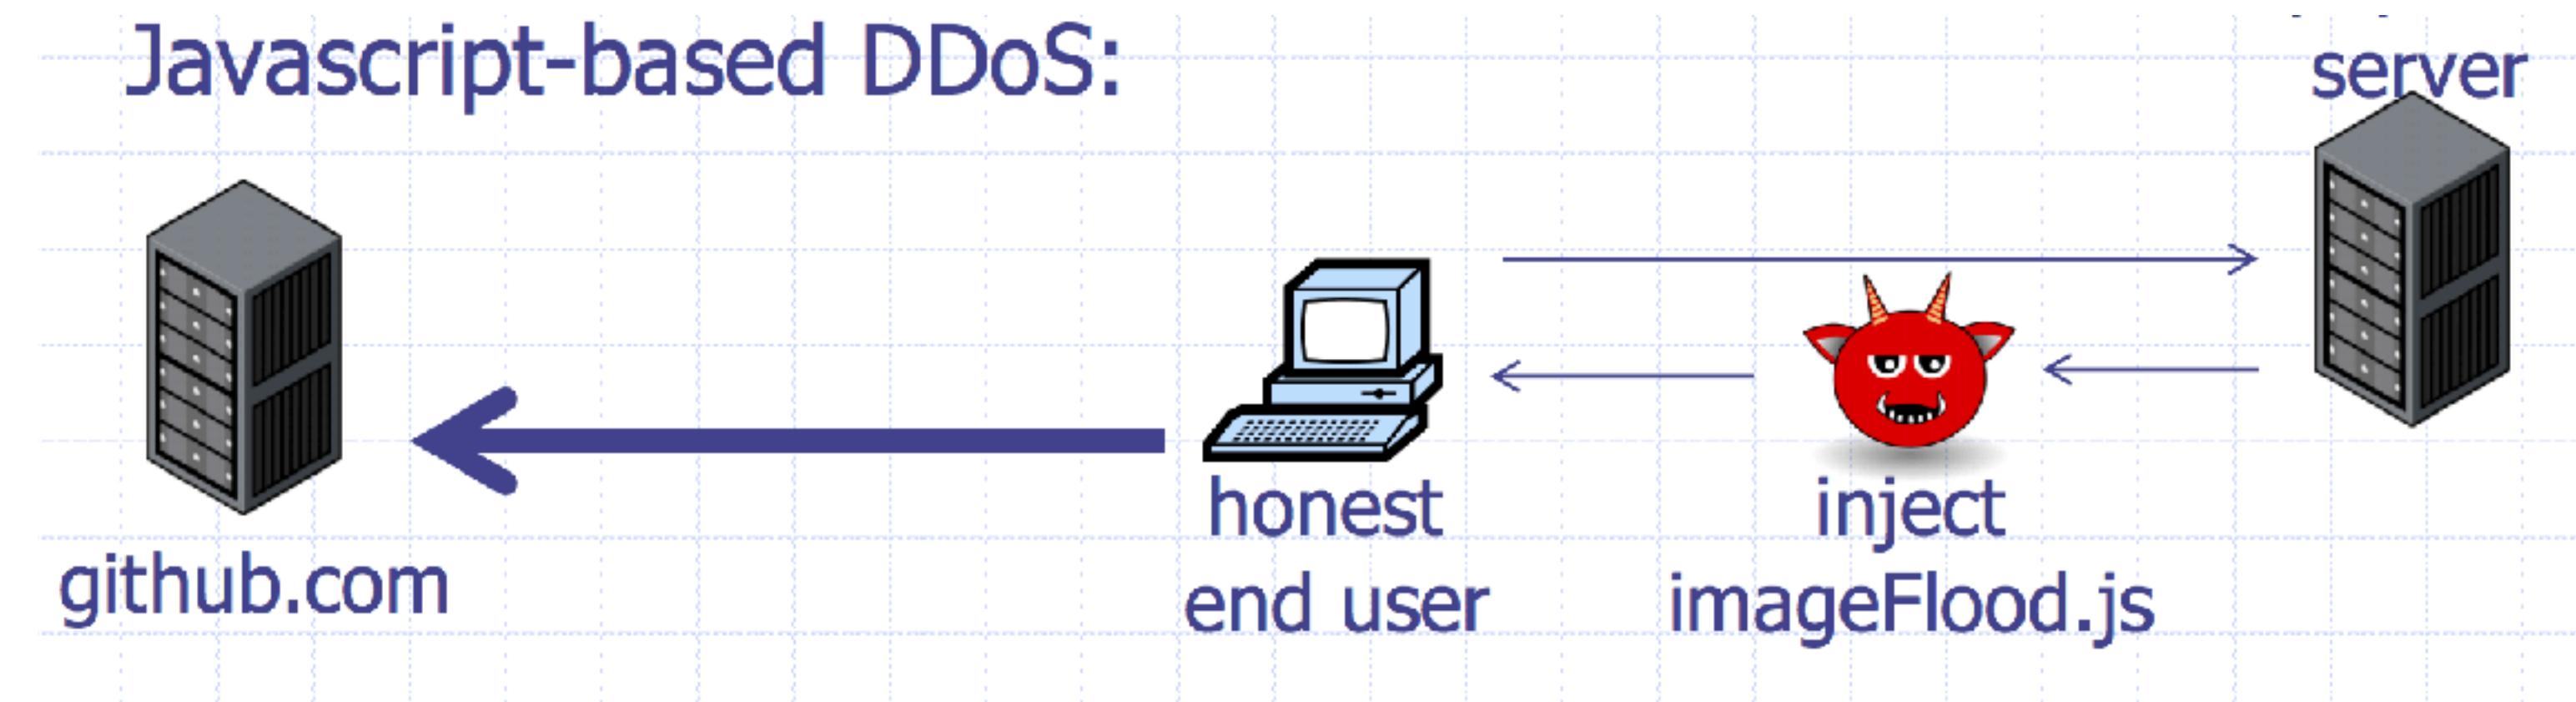
Will bypass flood protections.... but attacker can no longer use random source IPs

Victim site can block or rate limit bots

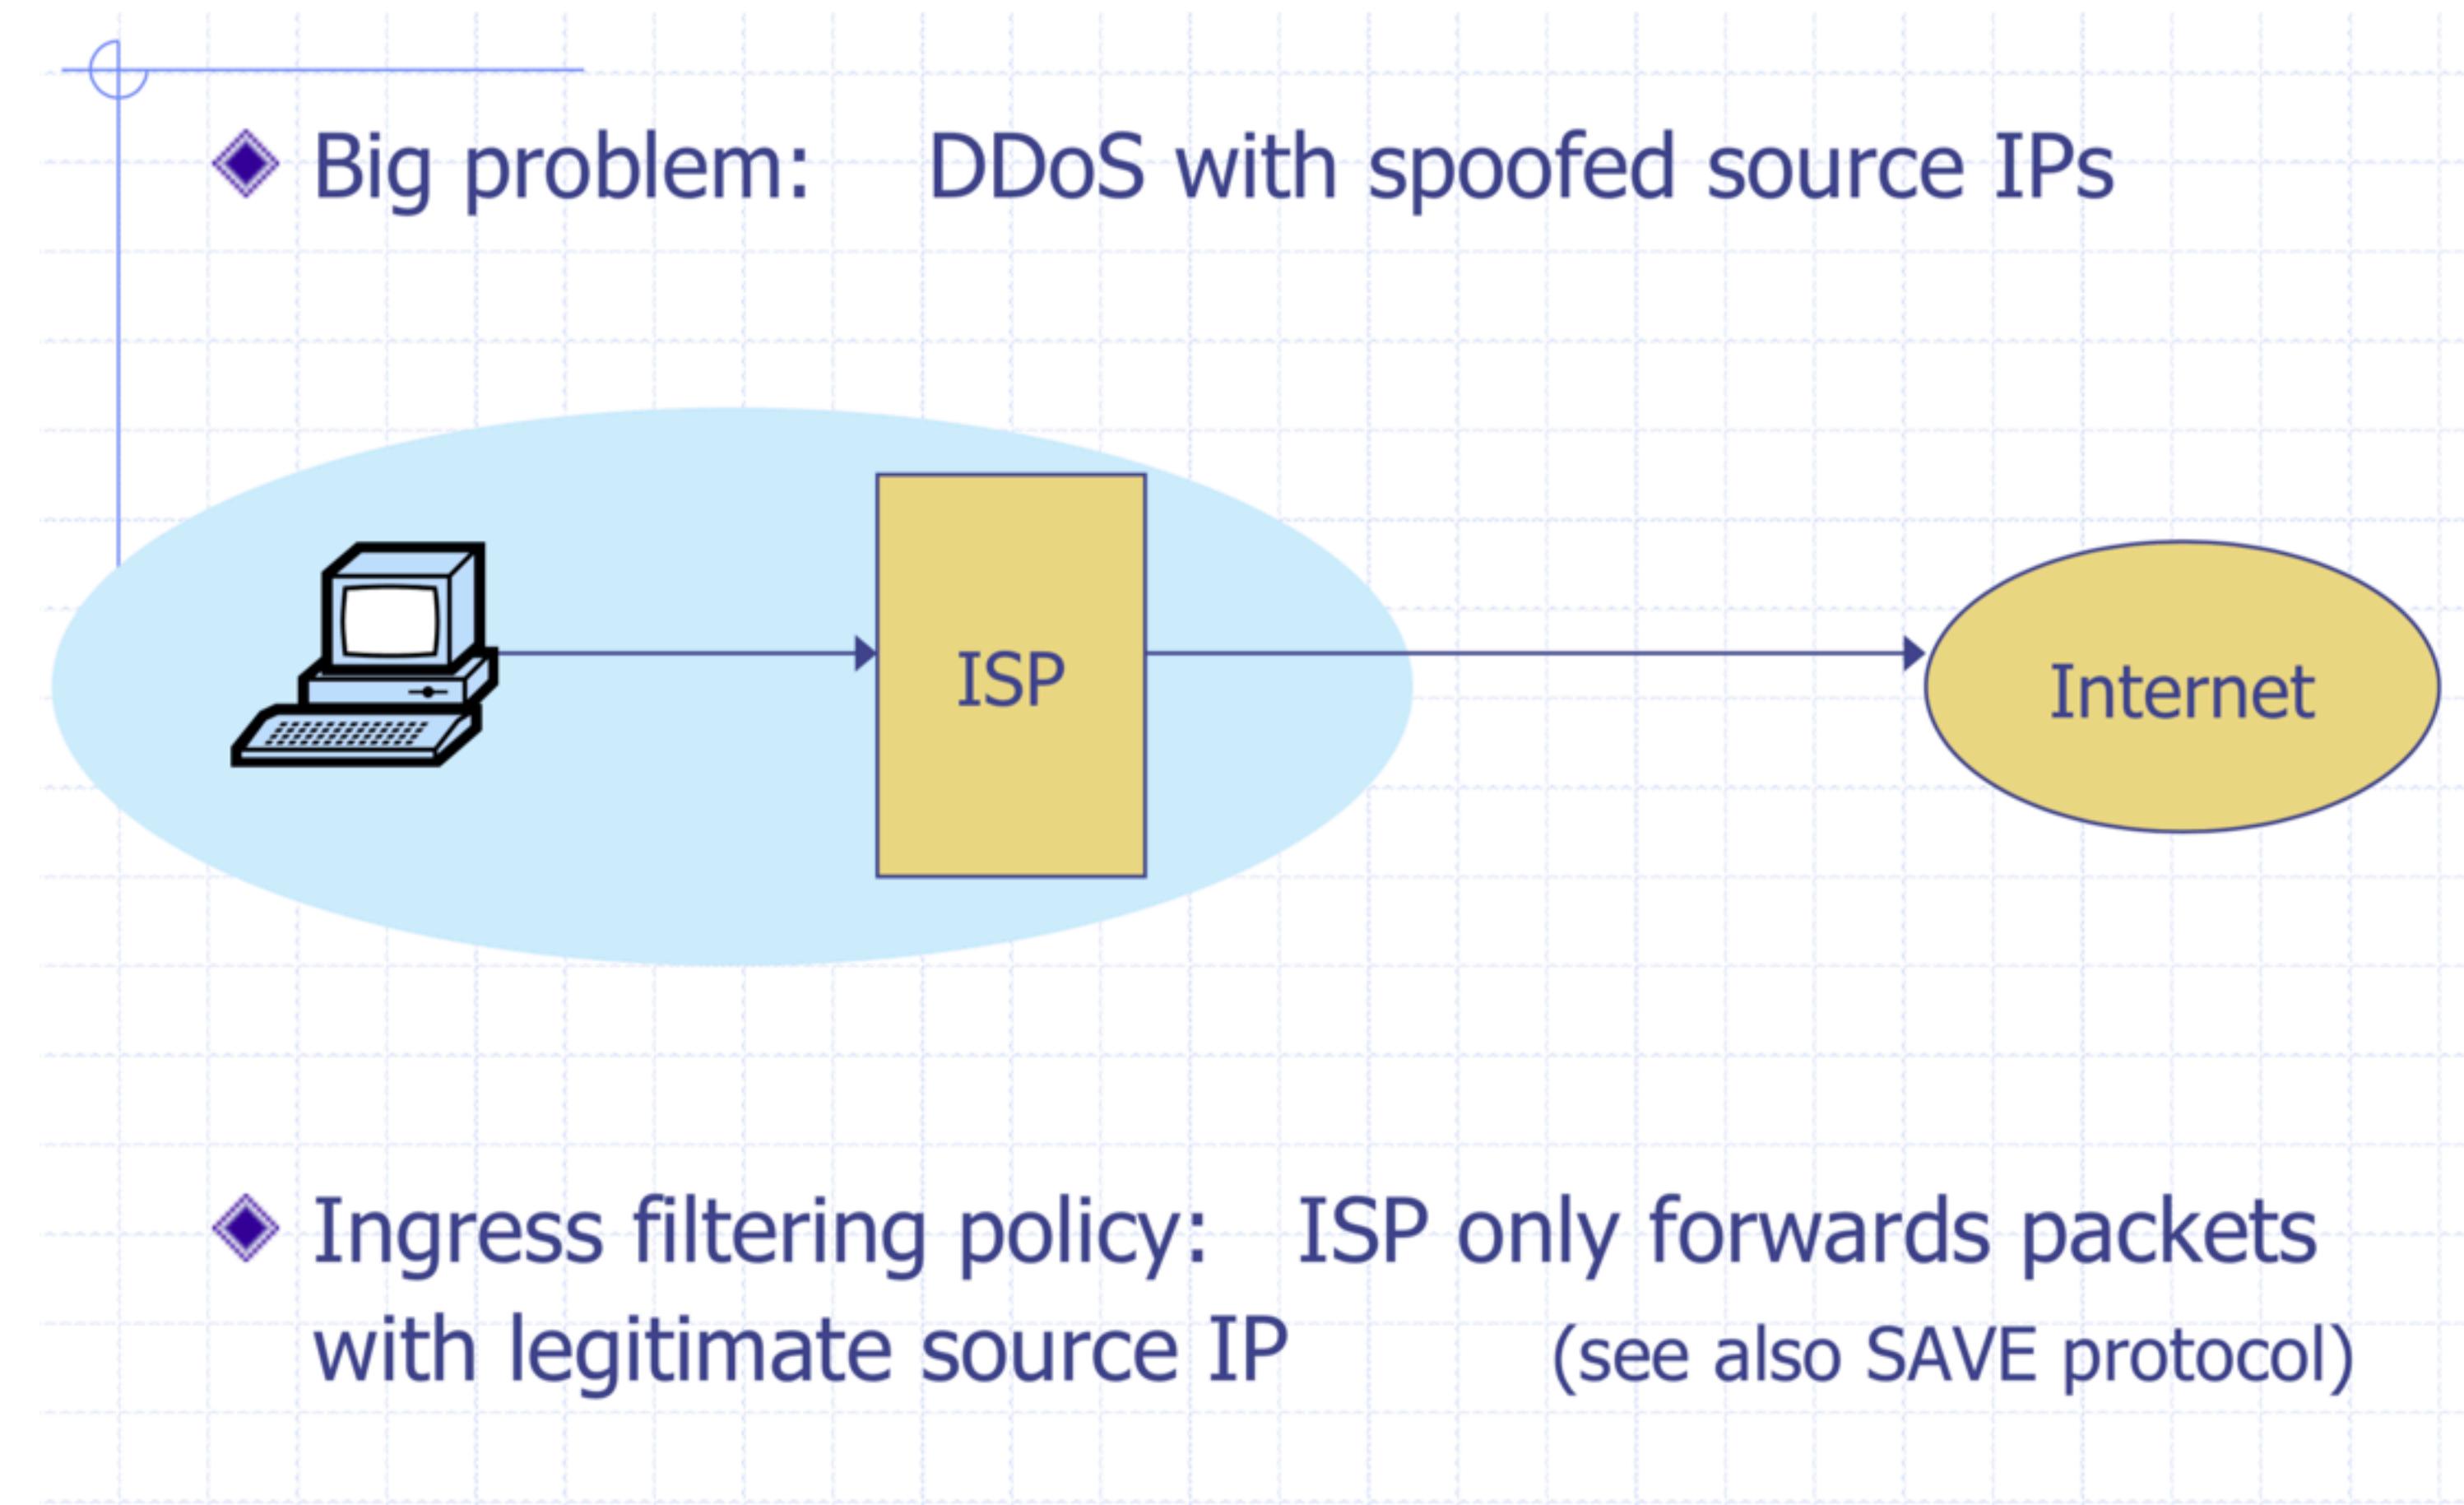
Github Attacks

1.35 Tbps attack against Github caused by javascript injected into HTTP web requests

The Chinese government was widely suspected to be behind the attack



Ingress Filtering



Ingress Filtering

All ISPs need to do this – requires global coordination

If 10% of networks don't implement, there's no defense

No incentive for an ISP to implement – doesn't affect them

As of 2017 (from CAIDA):

33% of autonomous systems allow spoofing

23% of announced IP address space allow spoofing

2013 300 Gbps attack sent attack traffic from only 3 networks

Client Puzzles

Idea: What if we force every client to do moderate amount of work for every connection they make?

Example:

- 1) Server Sends: C
- 2) Client: find $X \mid \text{LSB}_n(\text{SHA1}(C \parallel X)) = 0^n$

Assumption:

Puzzle takes 2^n for the client to compute (0.3 s on 1Ghz core)

Solution is trivial for server to check (single SHA-1 hash)

Client Puzzles

Not frequently used in the real world

Benefits:

- Can change n based on amount of attack traffic

Limitations:

- Requires changes to both protocols, clients, and servers
- Hurts low power legitimate clients during attack (e.g., phones)

Network Defenses

Local Services

Review: Popular TCP and UDP services live on standardized ports.
HTTPS servers listen on TCP/443. SSH on TCP/22.

Some services you don't want listening on the public Internet.

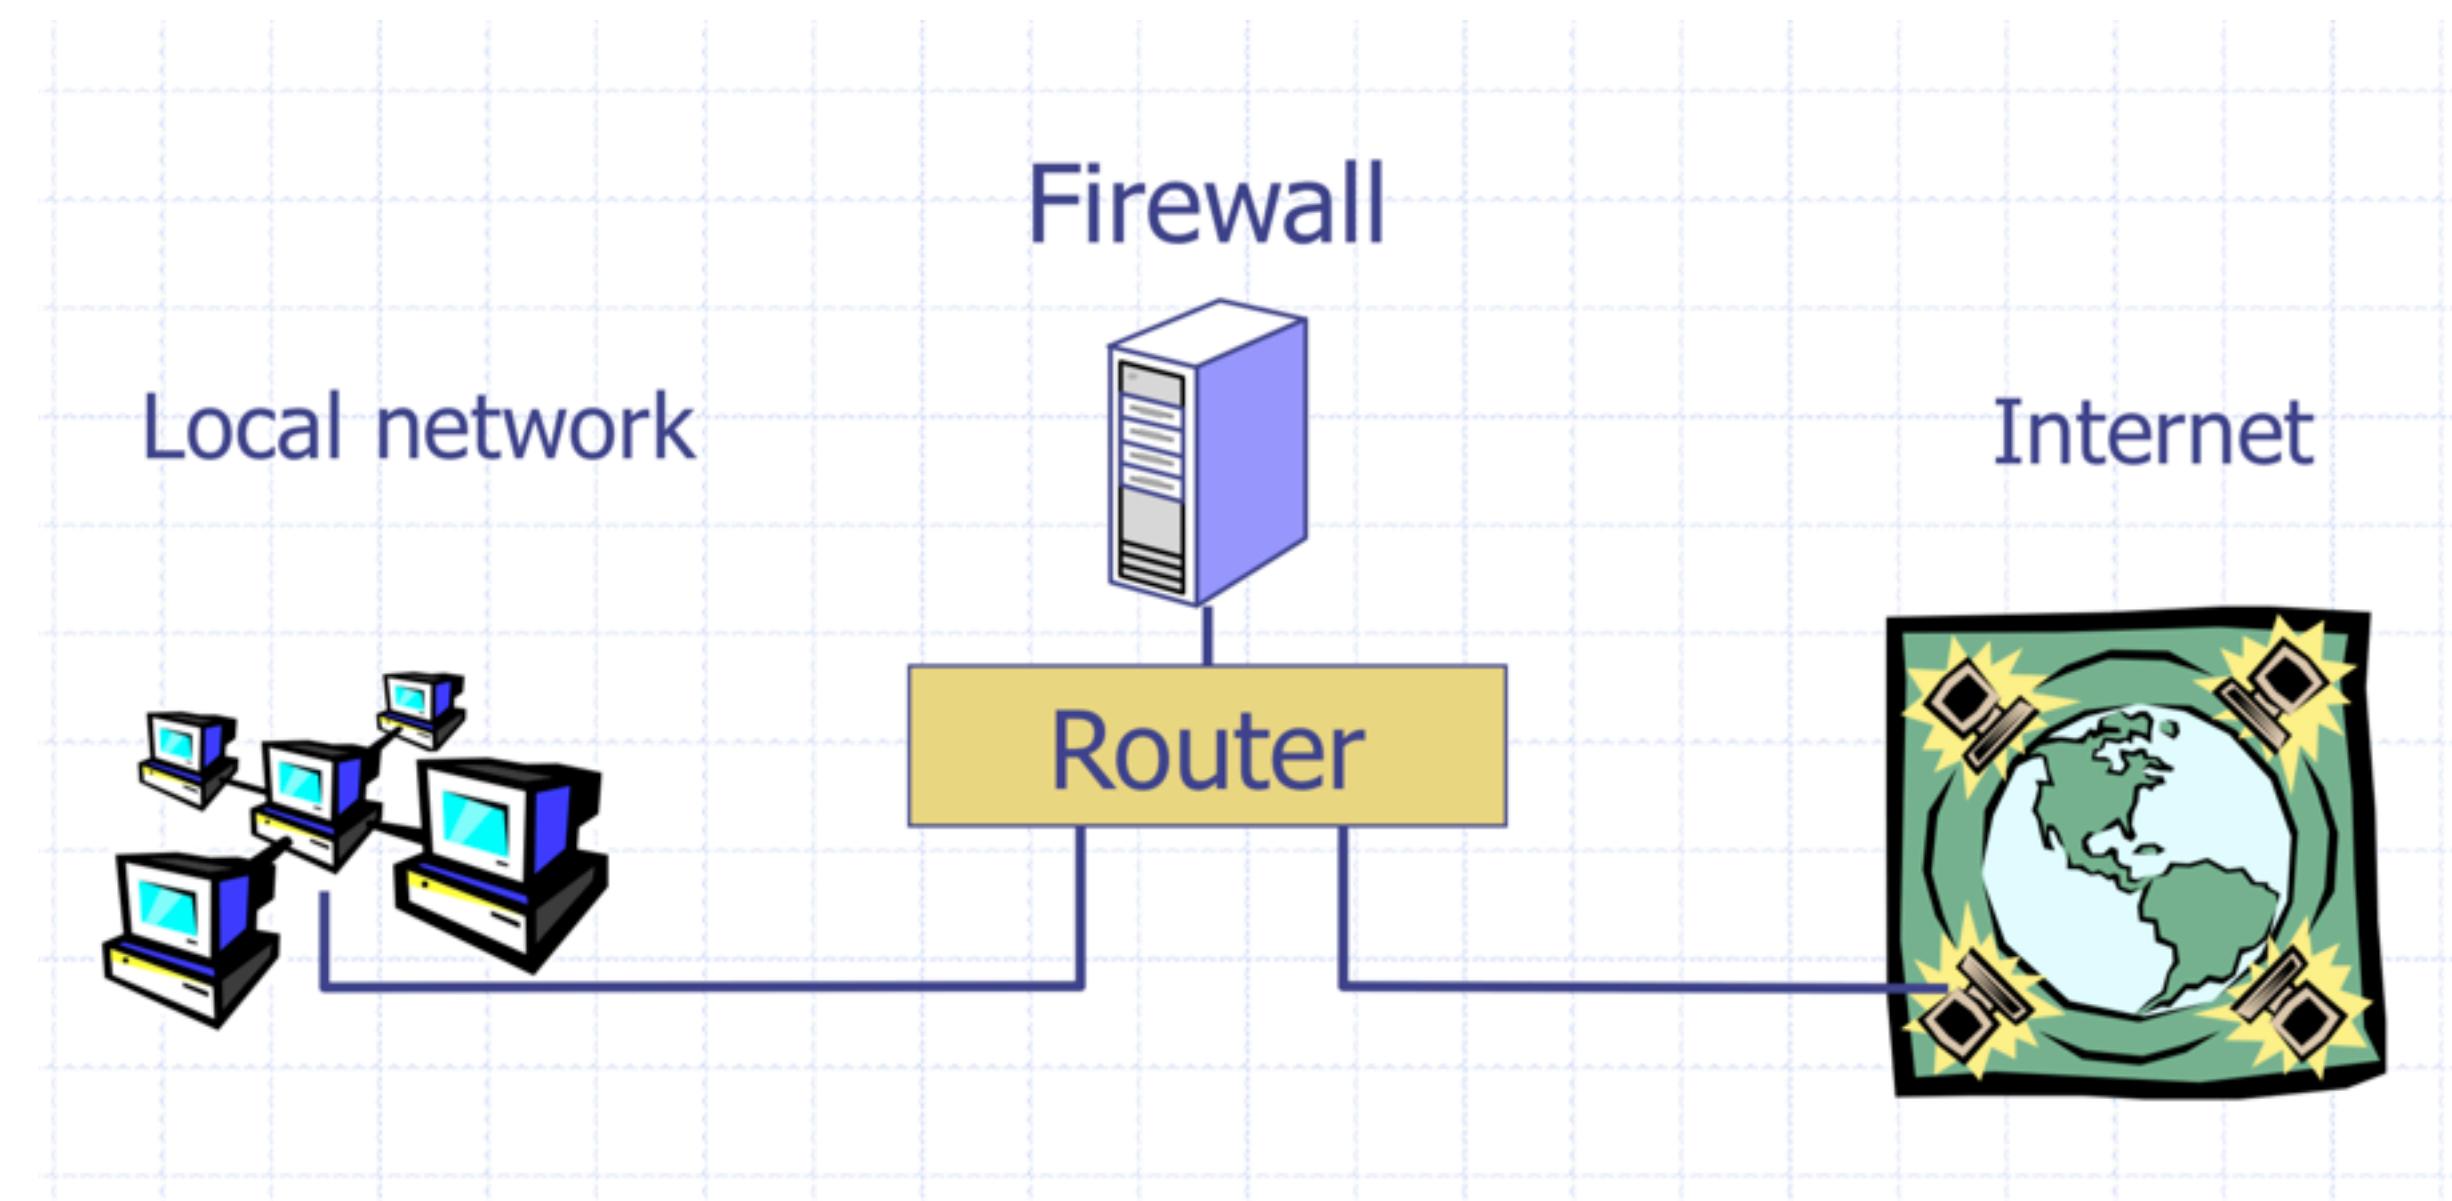
Recursive DNS Resolvers: allows attackers to mount DDoS attacks

Windows File Sharing: historically full of vulnerabilities. What if a local machine doesn't have a secure password on it?

Firewalls

Separate local area network (LAN) from the Internet. Only allow some traffic to transit.

Sometimes rules on a router. Sometimes a standalone device.



Basic Packet Filtering

Uses transport and IP layer information only

- IP Source Address, Destination Address
- Protocol (TCP, UDP, ICMP, etc.)
- TCP and UDP source and destination ports

Examples:

- “Do not allow external hosts to connect to Windows File Sharing”
-> DROP ALL INBOUND PACKETS TO TCP PORT 445

What's the rule?

What if you have a network with lots of servers but only want outsiders to be able to access a web server?

DROP ALL INBOUND PACKETS IF DEST PORT != 80

All outbound connections also have a source port! Their responses will be blocked!

IANA Port Numbering

System or Well-Known Ports [1,1023]:

Common services, e.g., HTTP -> 80, SSH -> 22

User or registered ports [1024, 49151]

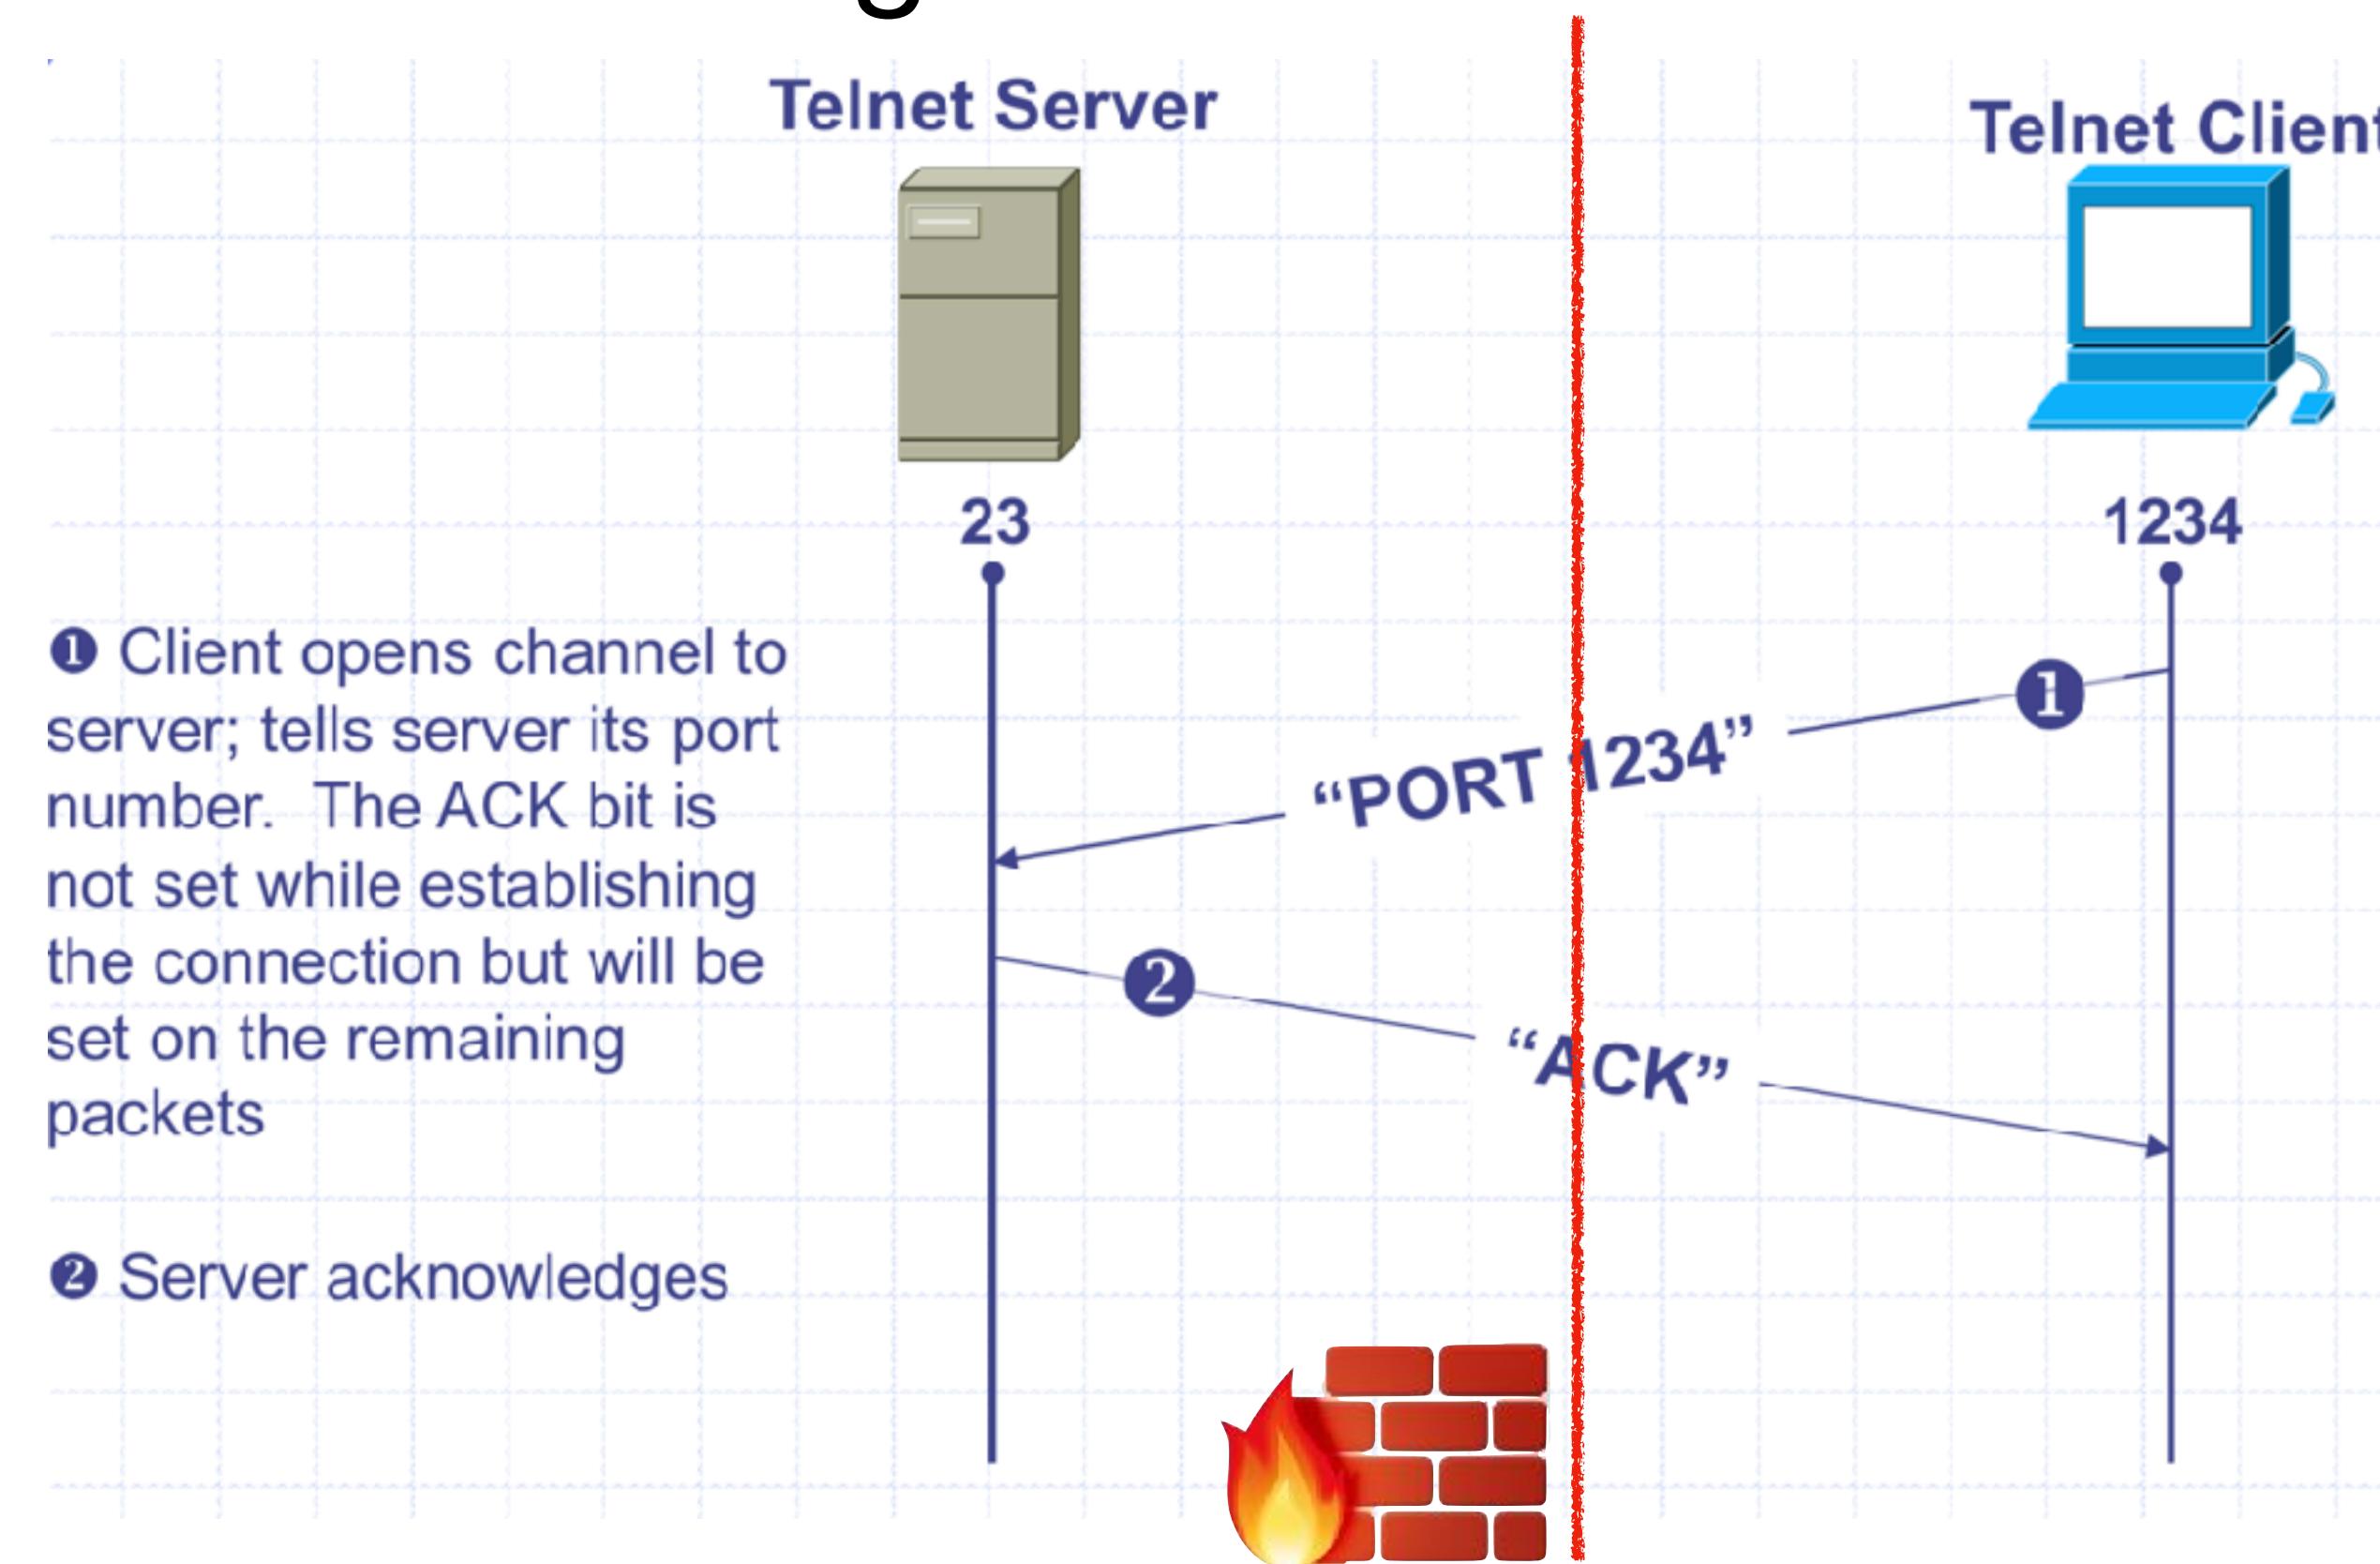
Less well-known services

Ephemeral/Dynamic/Private Ports [49152, 65535]

Short lived connections

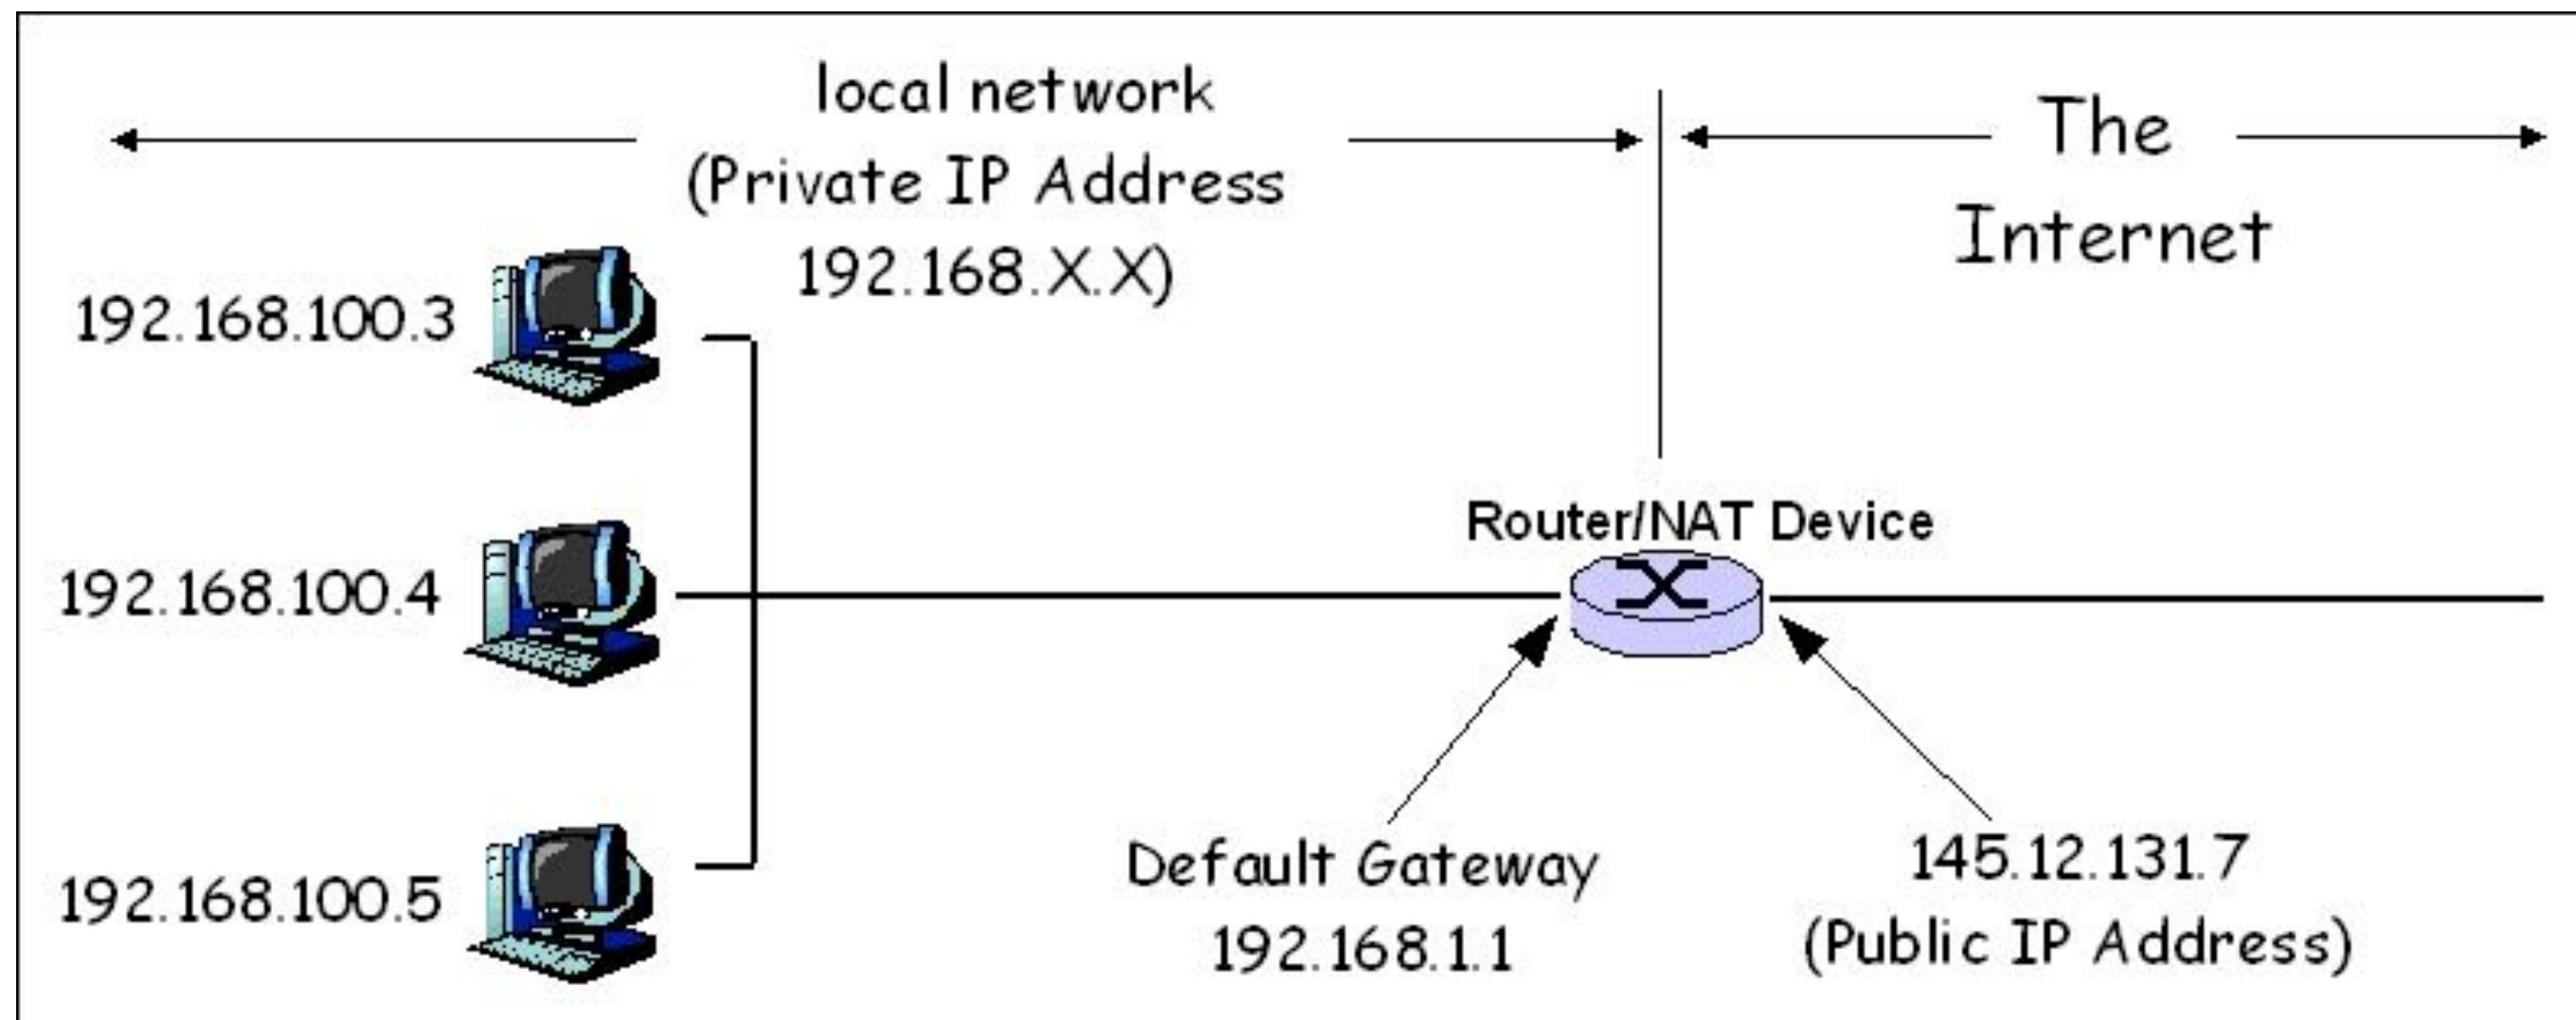
Stateful Filtering

Firewall tracks outgoing connections and allows associated inbound traffic back through



Network Address Translation (NAT)

NATs map between two different address spaces. Most home routers are NATs and firewalls.



Private Subnets

10.0.0.0 – 10.255.255.255

172.16.0.0 – 172.31.255.255

192.168.0.0 – 192.168.255.255

Local vs. Network Firewall

Firewalls we've discussed so far have all been network firewalls.
Most have lived at the edge of the organization.

Firewalls also run on individual hosts. Linux servers use **iptables**.

Typically have a combination of network and host firewalls

```
sudo iptables -A INPUT -m conntrack --ctstate ESTABLISHED,RELATED -j ACCEPT
```

```
sudo iptables -A INPUT -p tcp --dport 22 -m conntrack --ctstate NEW,ESTABLISHED -j ACCEPT
```

Complications: Normal IP Fragmentation



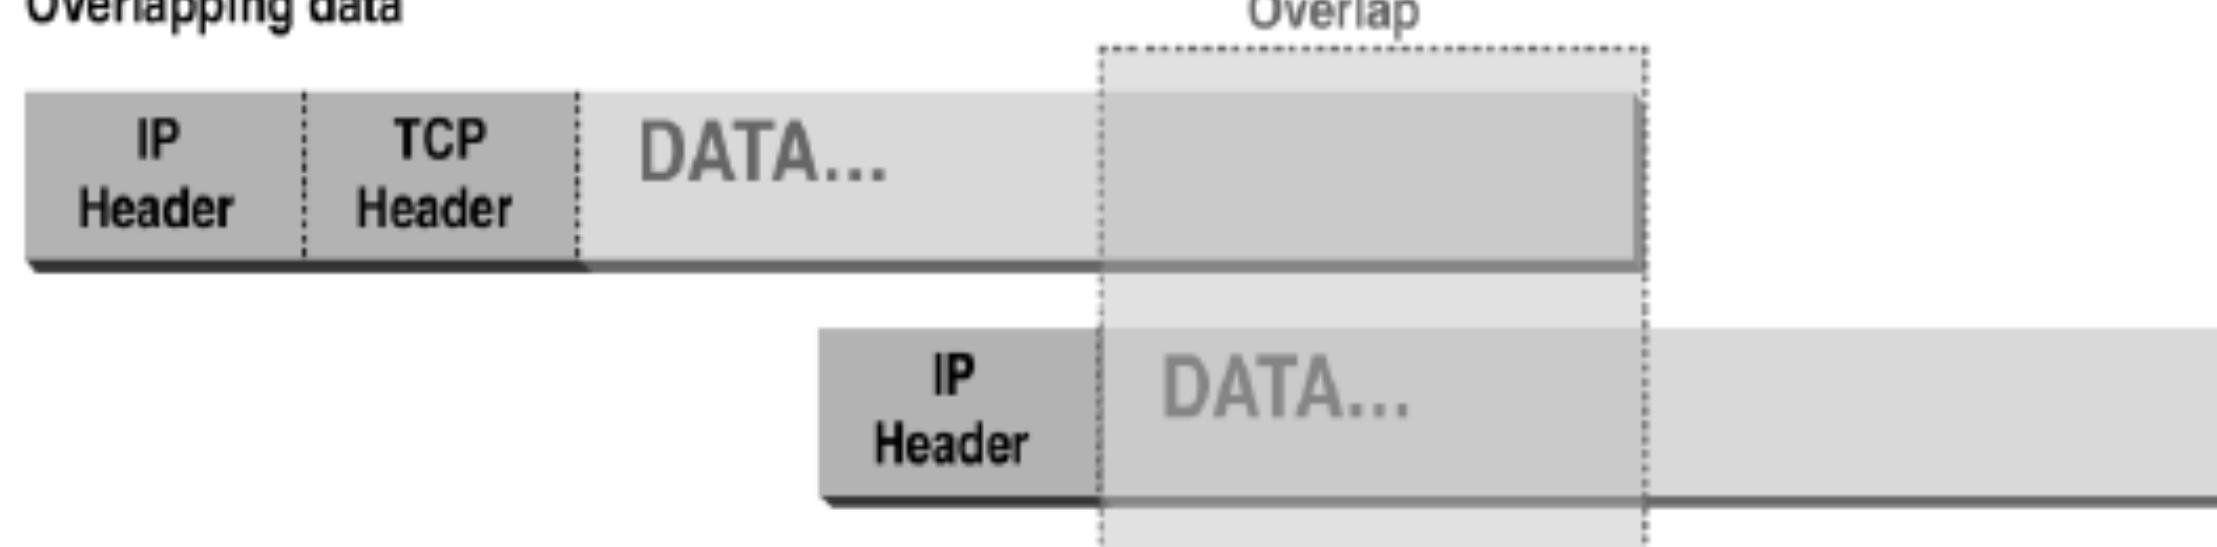
Flags and offset inside IP header indicate packet fragmentation

Abnormal Fragmentation

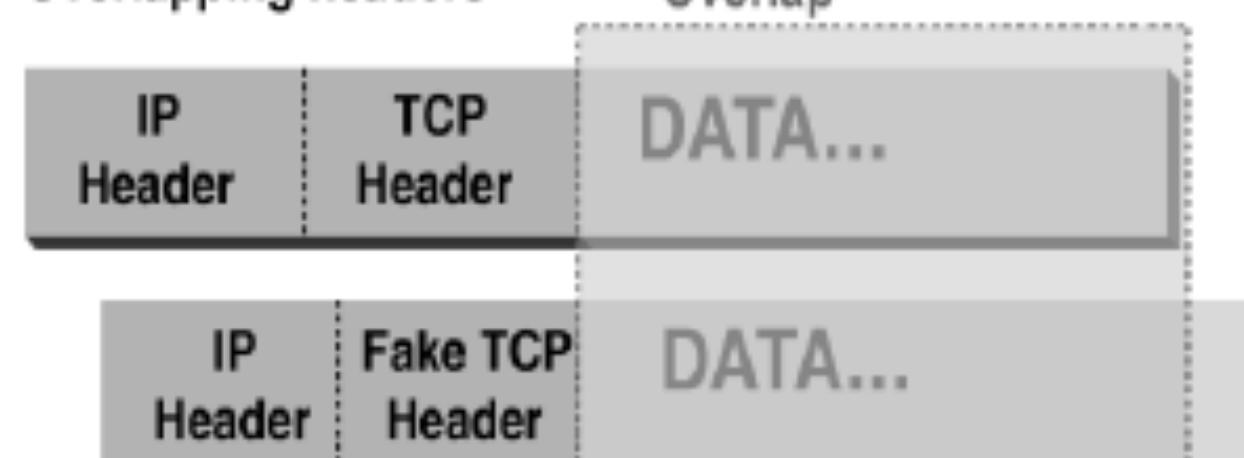
Normal



Overlapping data



Overlapping headers



Low offset allows second packet to overwrite TCP header at receiving host

Packet Fragmentation Attack

Firewall configuration

- TCP port 23 is blocked but SMTP port 25 is allowed

First packet

- Fragmentation Offset = 0, DF bit = 0 : "May Fragment", MF bit = 1 : "More Fragments"
- Destination Port = 25. TCP port 25 is allowed, so firewall allows packet

Second packet

- Fragmentation Offset = 1: second packet overwrites all but first 8 bits of the first packet
- DF bit = 0 : "May Fragment" , MF bit = 0 : "Last Fragment."
- Destination Port = 23. Normally be blocked, but sneaks by!

What happens

- Firewall ignores second packet “TCP header” because it is fragment of first
- At host, packet reassembled and received at port 23

Application Layer Filtering

Enforce protocol-specific policies:

- Virus scanning for SMTP
 - Need to understand protocol, MIME encoding, ZIP files, etc
- Look for SQL injection attacks in HTTP POSTs

Outbound Too!

Organizations will often inspect outbound traffic as well

- Block access to sites with known malicious behavior
- Prevent exfiltrating data
- Block services like bit torrent

Be careful on enterprise networks! Sometimes companies will even install their own root certificates on employee workstations to monitor TLS traffic.

Intrusion Detection Systems (IDS)

Software/device to monitor network traffic for attacks or policy violations

Violations are reported to a central security information and event management (SIEM) system where analysts can later investigate

Signature Detection: maintains long list of traffic patterns (rules) associated with attacks

Anomaly Detection: attempts to learn normal behavior and report deviations

Open Source IDS

Three Major Open Source IDS (and a *tremendous* number of commercial products)

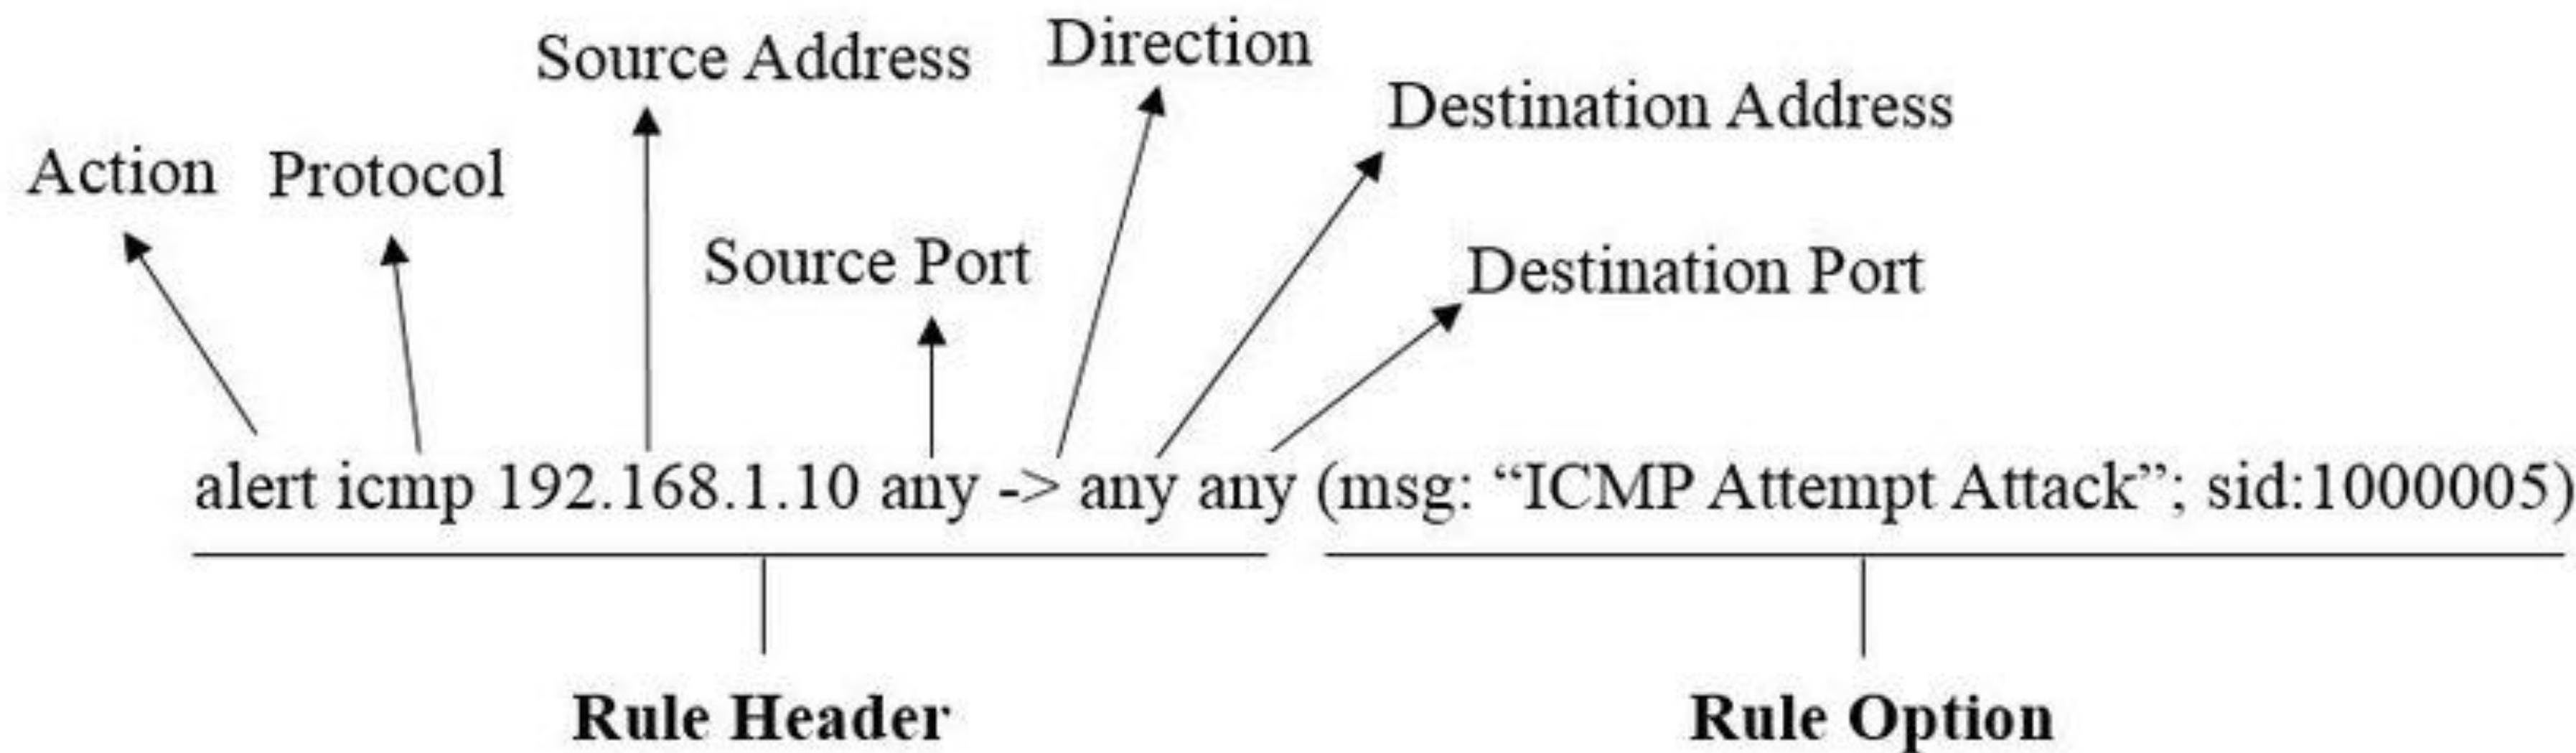
Snort

Bro Zeek

Suricata



Example Snort Rule



Snort challenges

Misuse detection – avoid known intrusions

- Database size continues to grow
 - Snort version 2.3.2 had 2,600 rules
 - Snort spends 80% of time doing string match

Anomaly detection – identify new attacks

- Probability of detection is low

Difficulties in anomaly detection

Lack of training data

- Lots of “normal” network, system call data
- Little data containing realistic attacks, anomalies

Data drift

- Statistical methods detect changes in behavior
- Attacker can attack gradually and incrementally

Main characteristics not well understood

- By many measures, attack may be within bounds of “normal” range of activities

False identifications are very costly

- Sys Admin spend many hours examining evidence

Remote Access

Virtual Private Networks (VPNs)

Problem: How do you provide secure communication for non-TLS protocols across the public Internet?

VPNs create a fake shared network on which traffic is encrypted

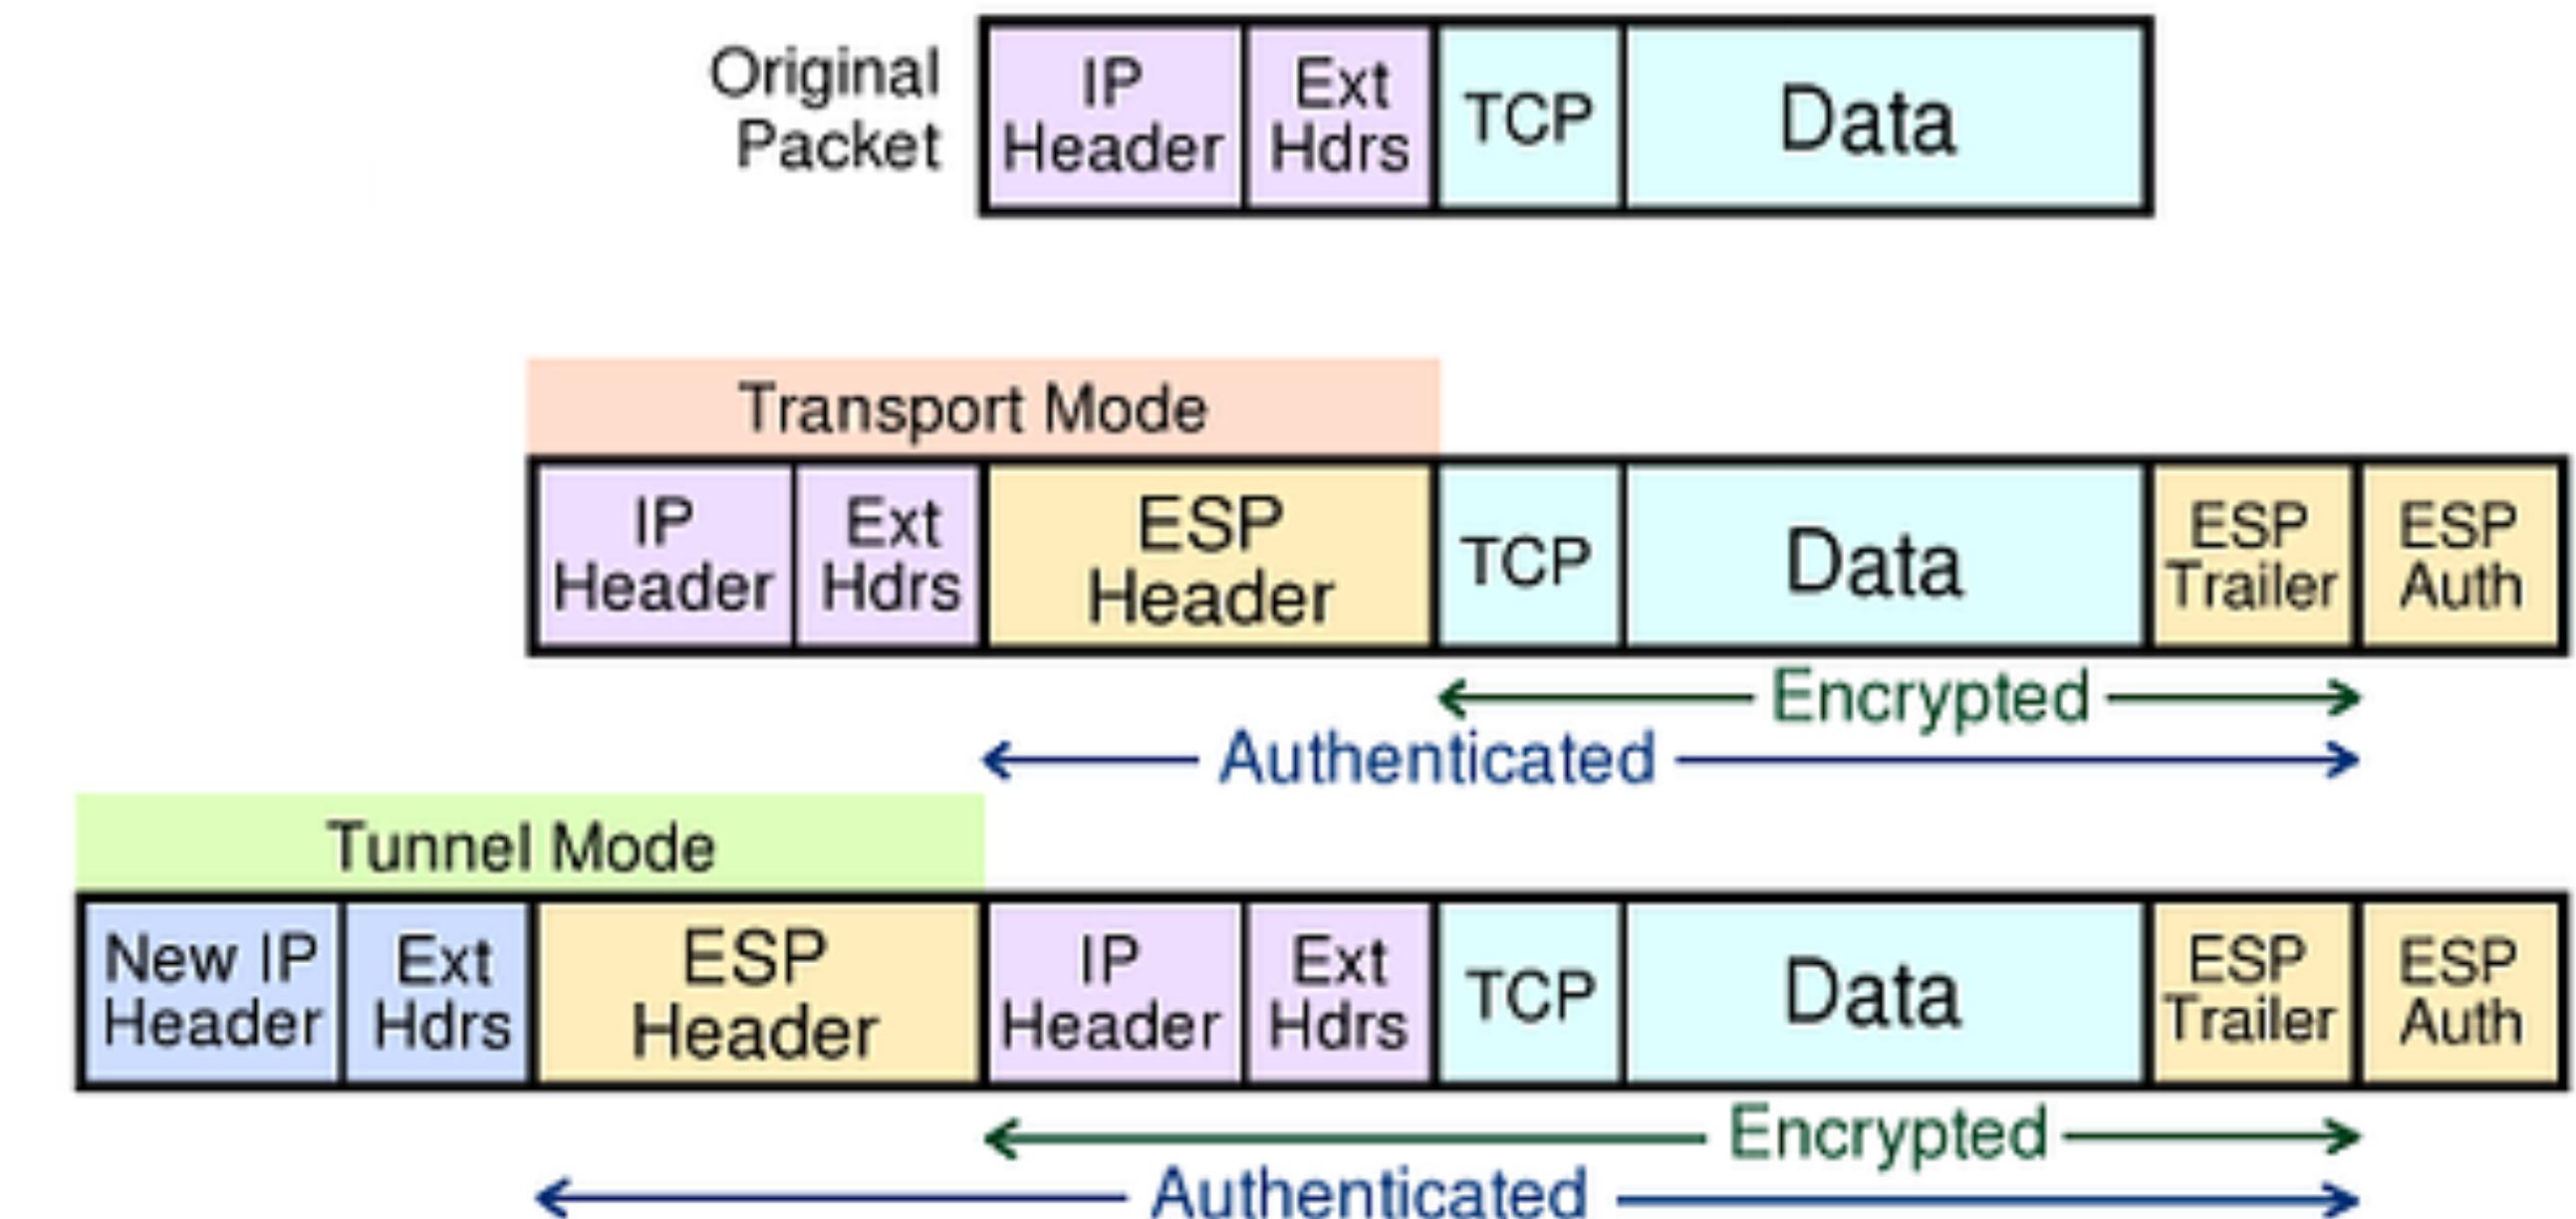
Two Broad Types:

- Remote client (e.g., traveler with laptop) to corporate network
- Connect two remote networks across Internet

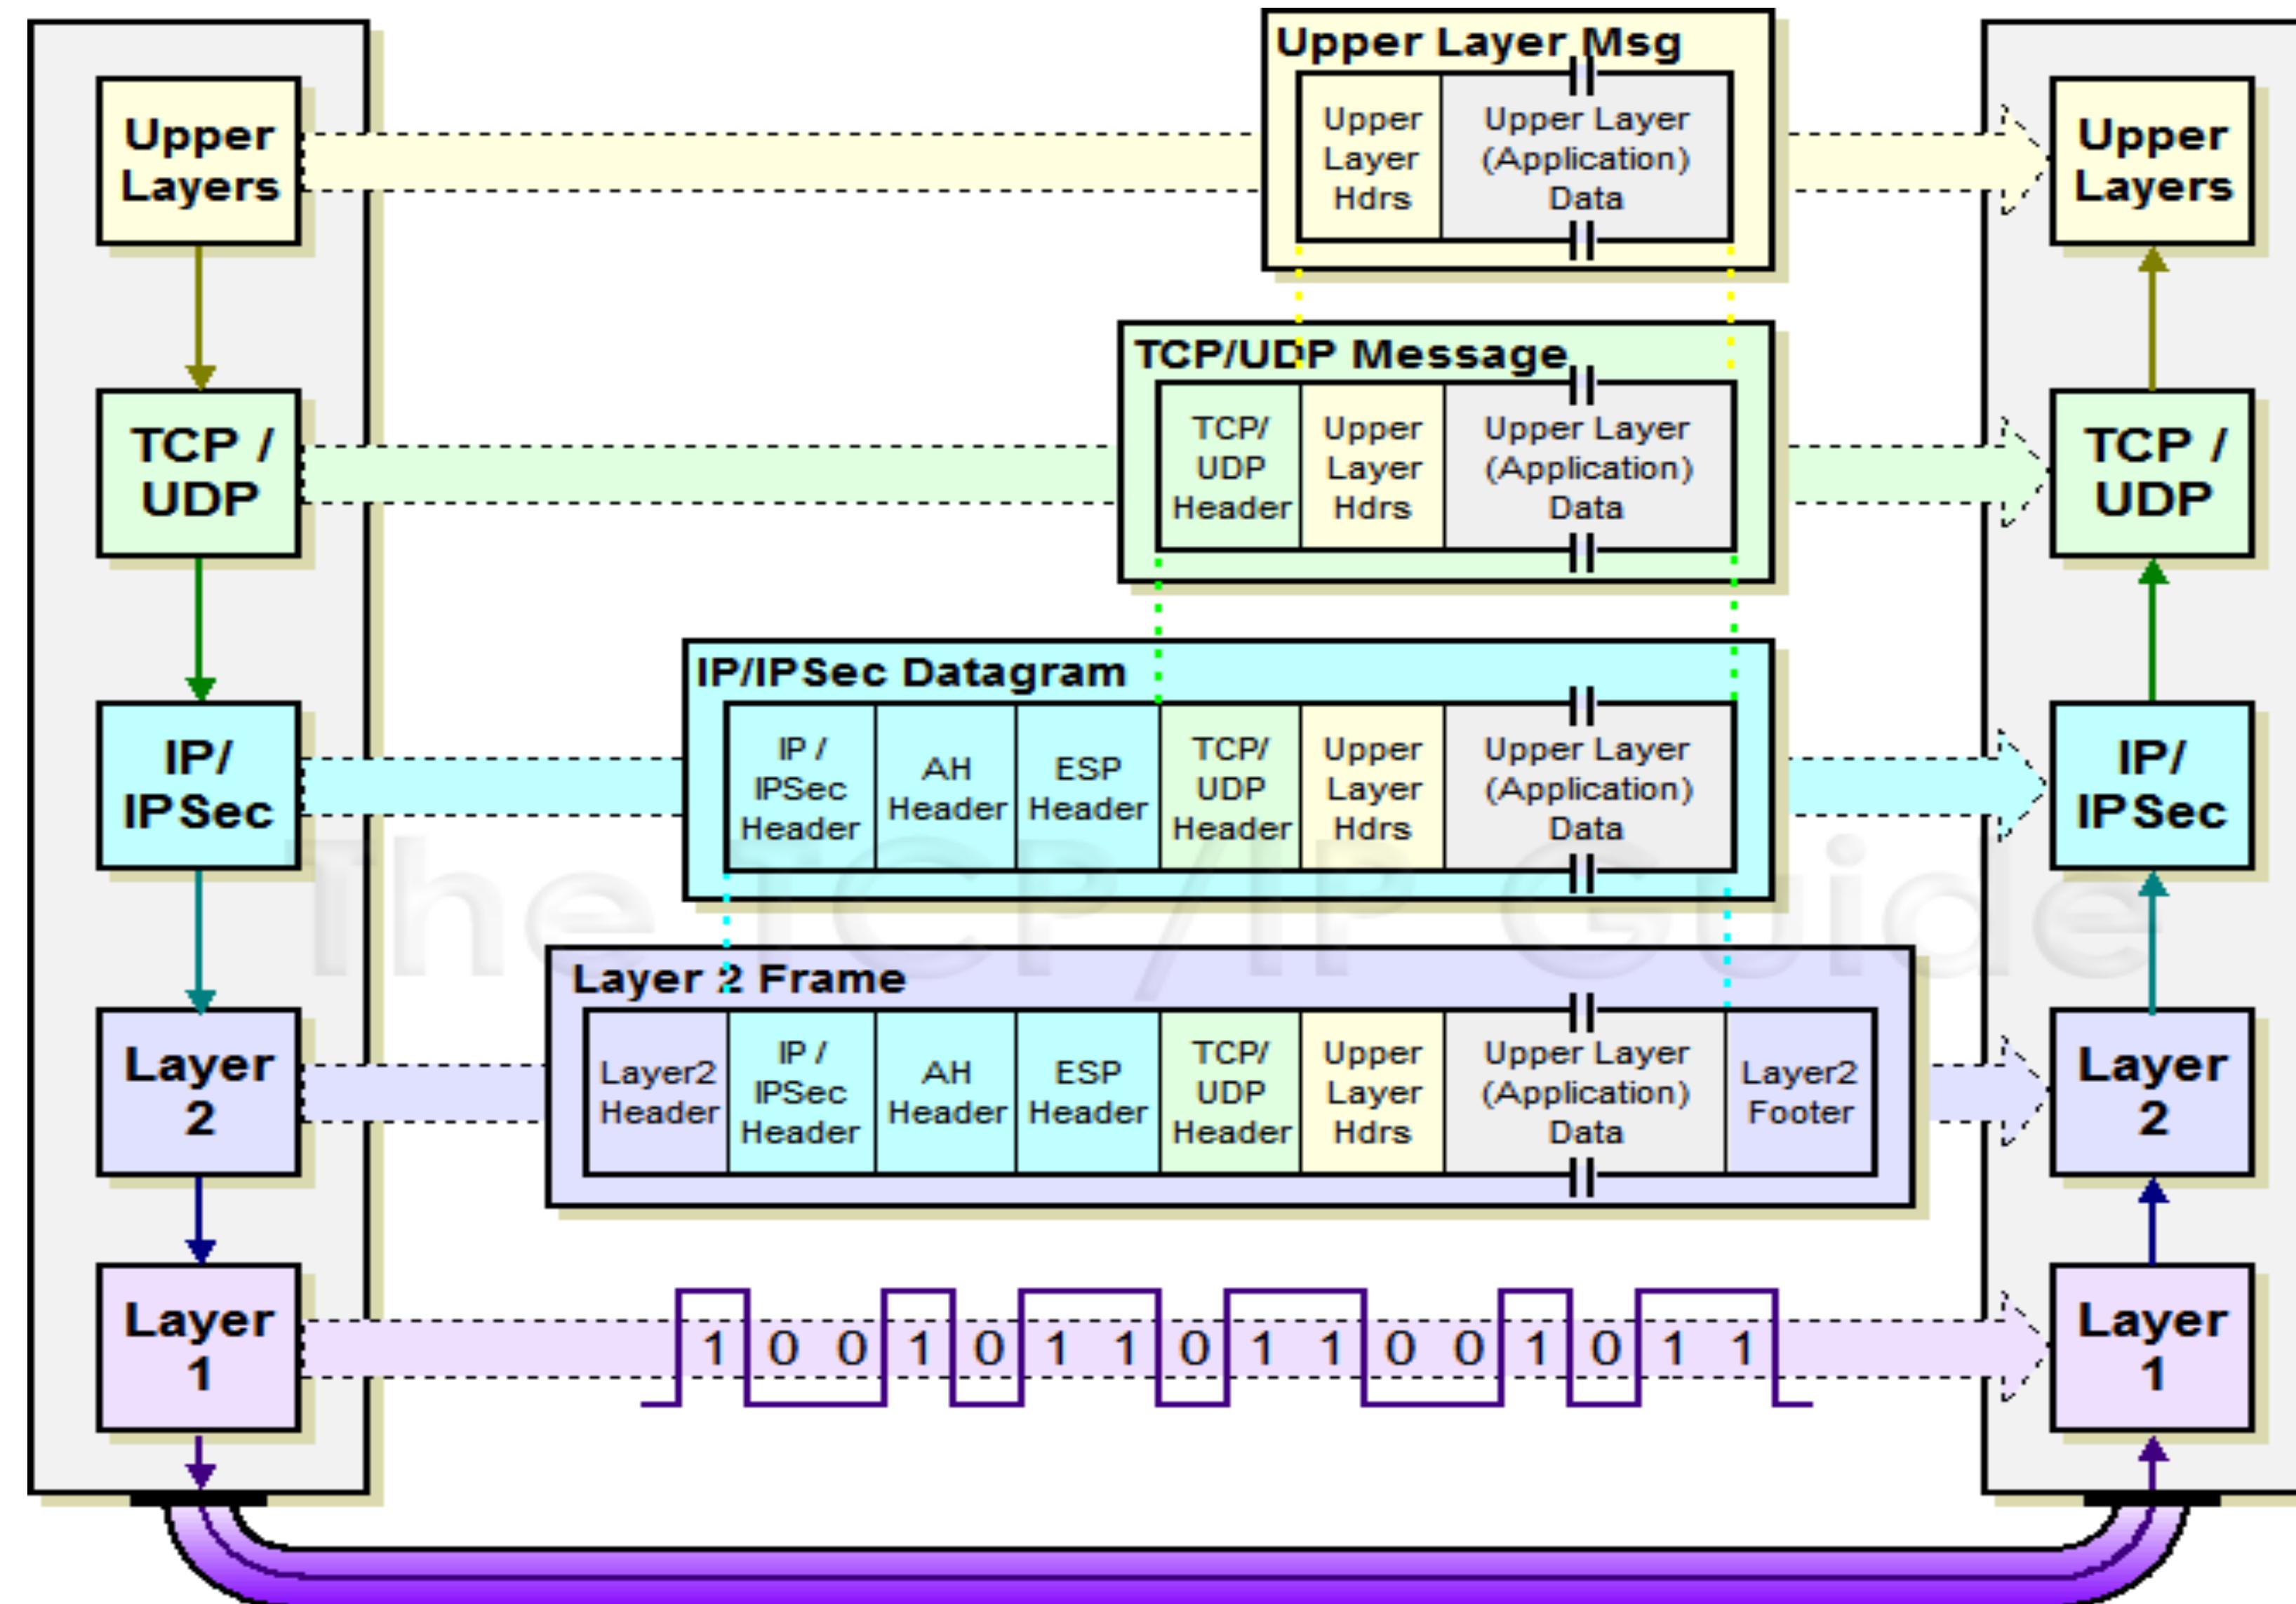
IPSec

Several VPN protocols exist (PPTP, L2TP, IPsec, OpenVPN)

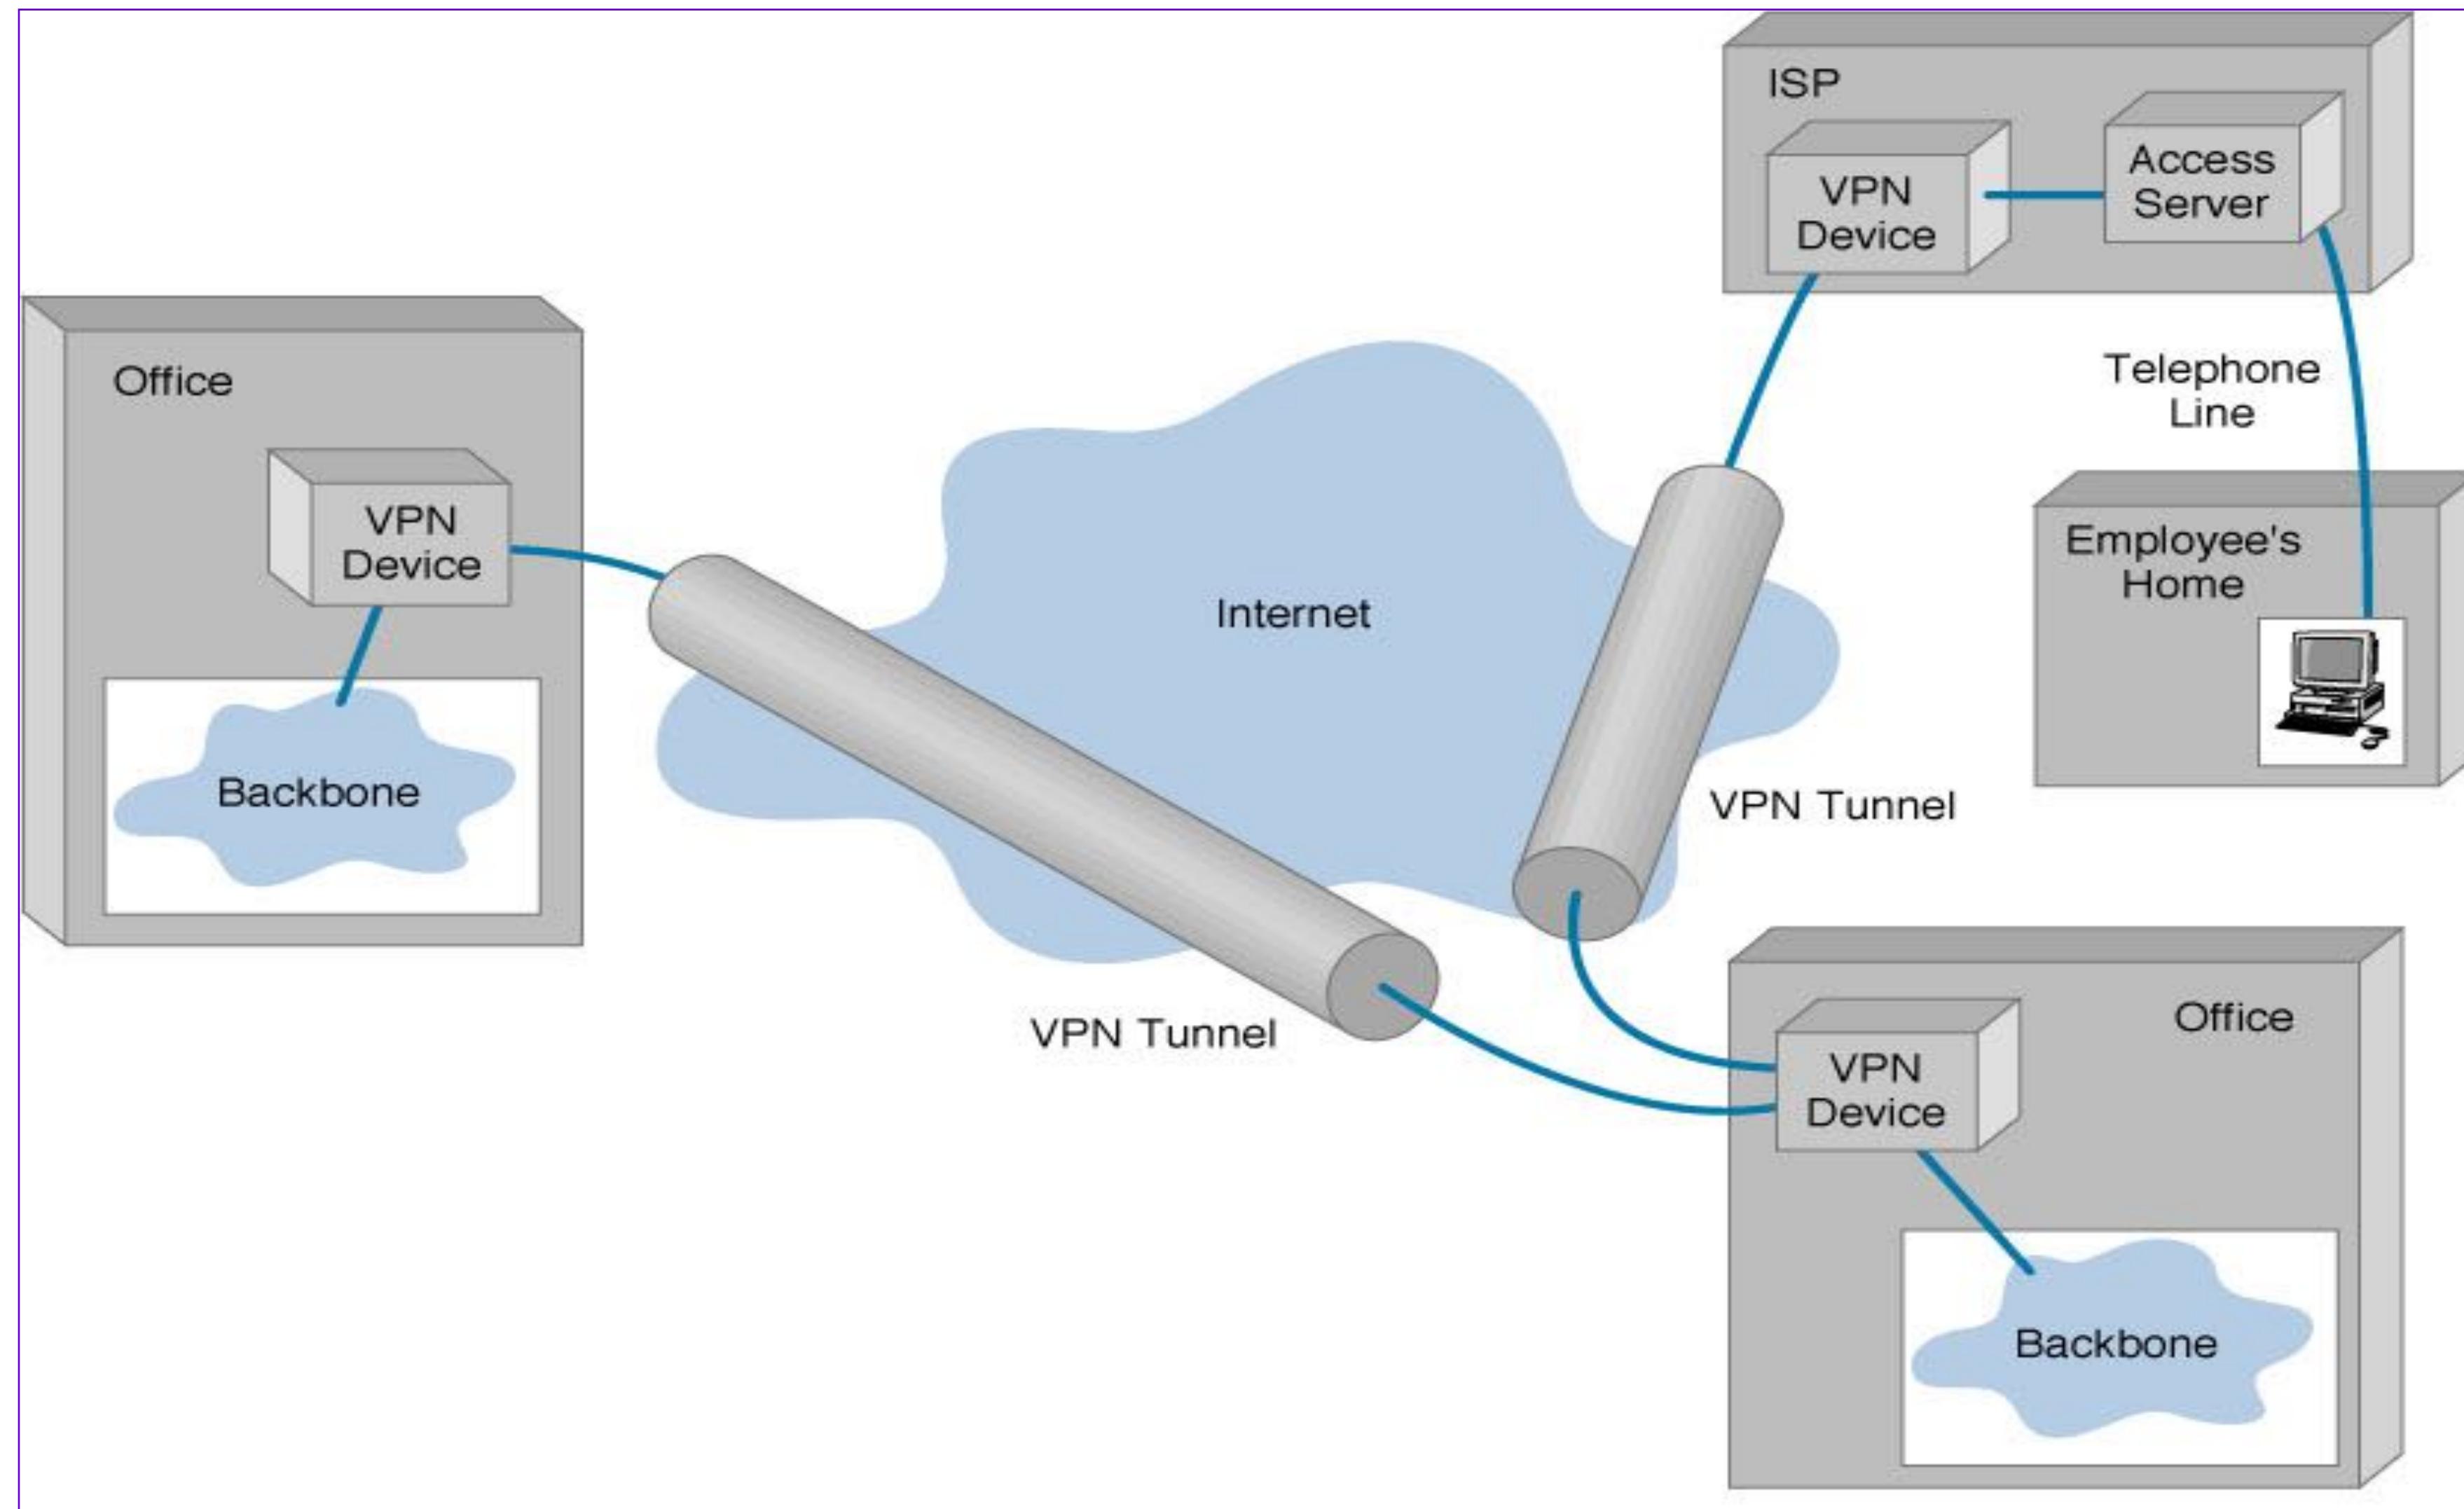
Most popular is IPsec. OpenVPN is open source.



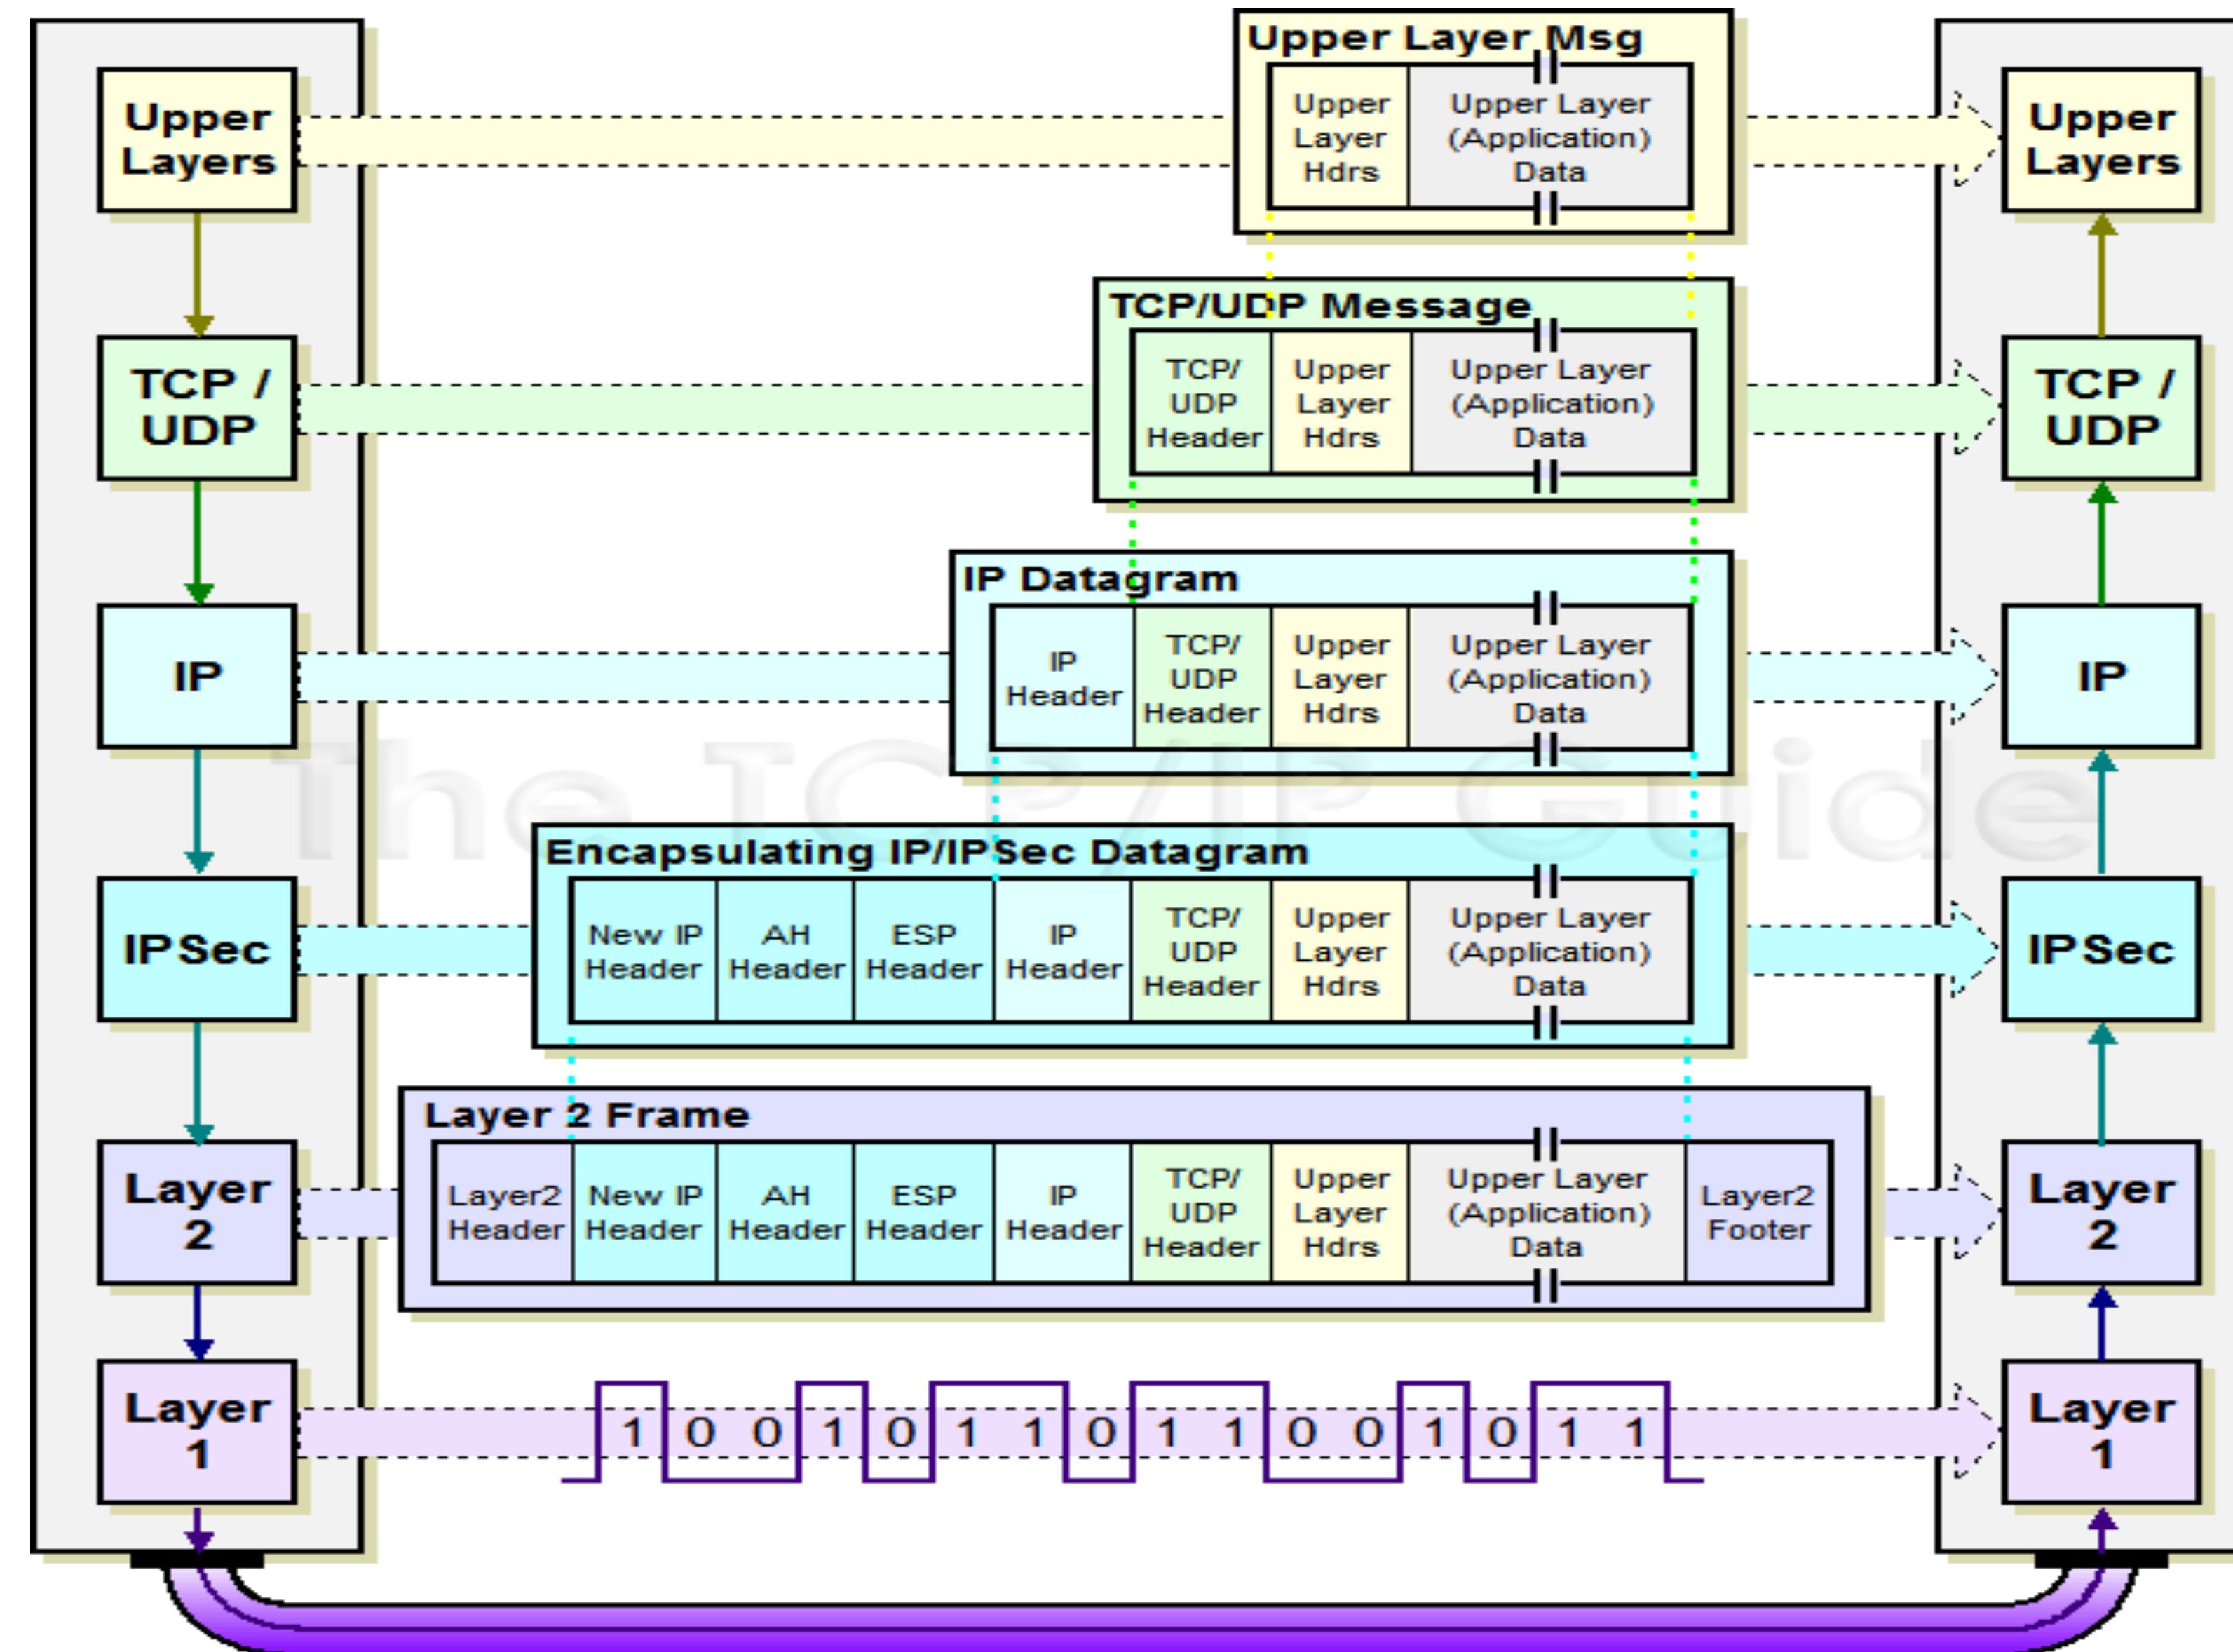
IPSec Transport Mode: IPSEC instead of IP header



IPSEC Tunnel Mode



IPSec Tunnel Mode: IPSEC header + IP header



Cisco AnyConnect

Many organizations use Cisco AnyConnect

Encapsulates traffic in TLS! Initial handshake uses normal TCP-based TLS for initial handshake and then DTLS (UDP-based TLS) to transport data

Gooey Middle

VPNs support the idea of having a secure internal network and untrusted public Internet. Unfortunately, attacker has a ton of access once the network perimeter is breached.

Unfortunately, internal networks aren't *that* secure. Computers are compromised all the time and attackers have free rein.

Zero Trust Security (BeyondCorp)

Google: assume internal network is *also* out to get you. Remove privileged intranet and put all corporate applications on the Internet.

Access depends solely on device and user credentials, regardless of a user's network location

Protect applications, not the network